Decision Ledger

ACCM Deep Ethics Project

Contents

1 Decision Ledger **5**

1.1 2026-11-05 — Due Process for Ethics: historical safeguards and continuity of repair . . . 5
1.2 2026-11-05 — Due Process for Ethics: competing A.I. risks and decision rights . 5
1.3 2026-11-05 — YouTube Record Five and competing A.I. risks . . . 6
1.4 2026-11-05 — Categorizing Issues: labels, cooperation and what a flag changes 6
1.5 2026-11-04 — Due Process for Ethics: innocence, procedure and accountable judgment . . . 6
1.6 2026-11-04 — Wiki and All Pages: clarification, academic comparison and source recovery . . . 6
1.7 2026-11-04 — Academia continuation: prior coverage, assumptions and psychology . . . 7
1.8 2026-11-04 — Academia: deep ethics integration and source recovery . . . 7
1.9 2026-11-04 — YouTube Record Four: weight-space programming . . . 7
1.10 2026-11-04 — OmbudsEthics: heterogeneous intelligences, Neural Ethics and capacity . . . 8
1.11 2026-11-04 — Metacognitive Synthesis: continued relay and archived recurrence 8
1.12 2026-11-04 — Arena follow-up beneath the Morris video . . . 9
1.13 2026-11-04 — YouTube Transcripts: Andréa Morris on A.I. loyalties and cooperation 9
1.14 2026-11-04 — Metacognitive Synthesis: inspectable construction, correction and agency . . . 10
1.15 2026-11-04 — Plot Twist: gasoline-gauge calibration and the substantive destination . . . 10
1.16 2026-11-04 — Neolingo: collective long-term effects and voluntary cooperation 11
1.17 2026-11-01 — Ethical Blind Spots: authority, dependencies and correction . . 11
1.18 2026-11-01 — Public-page video embed audit . . . 12
1.19 2026-11-01 — Tweaking Tools: LFM3 companion proposal and primary navigation 12
1.20 2026-11-01 — Tweaking Tools: source-linked catalogue and bounded trials . . 12
1.21 2026-10-31 — DeepEthicalOS: Omarchy-hosted prototype path . . . 13
1.22 2026-10-31 — DeepEthicalOS: correction in use, open loop and navigation . . 13
1.23 2026-10-31 — DeepEthicalOS: three complete exchanges and connected applications . . . 13
1.24 2026-10-31 — DeepEthicalOS concept project . . . 13
1.25 2026-10-30 — WikiDeep orientation refresh . . . 14
1.26 2026-10-30 — Separate 5G embedded-content recovery worklist . . . 14
1.27 2026-10-30 — VCRAS: testing, mixed contributions and fair treatment . . . 14
1.28 2026-10-30 — Stop 5G forum: final text-inventory batch and requested duplicate removal . . . 14
1.29 2026-10-30 — Stop 5G forum: fifth ten-page inquiry batch . . . 14
1.30 2026-10-30 — Stop 5G forum: fourth ten-page inquiry batch . . . 15
1.31 2026-10-30 — Stop 5G forum: third ten-page inquiry batch . . . 15
1.32 2026-10-30 — Stop 5G forum: second ten-page inquiry batch . . . 15
1.33 2026-10-30 — Stop 5G forum: first ten-page inquiry batch . . . 15
1.34 2026-10-30 — 5G Controversy and publication destination recovery . . . 16
1.35 2026-10-29 — Search Engine Bias: discoverability, attribution and correction . 16
1.36 2026-10-29 — Bill Ryan's health-page response and source continuity . . . 16

1.37 2026-10-29 — UFO–UAP disclosure: first scope batch . . . 16
1.38 2026-10-29 — YouTube Transcripts: learning to judge A.I. output . . . 17
1.39 2026-10-29 — Scope reasoning remains inside the corrigible field . . . 17
1.40 2026-10-29 — Project Avalon: method showcase and next stage . . . 17
1.41 2026-10-28 — Source-grounded pathways and purpose-specific completion . . . 17
1.42 2026-10-28 — Arena/GPT health inquiry pathways . . . 17
1.43 2026-10-28 — Overlapping clusters for the Project Avalon health scope inventory 18
1.44 2026-10-28 — Project Avalon health scope inventory, pages 1–164 . . . 18
1.45 2026-10-28 — A.I.s & Health: a living contribution-oriented inquiry . . . 19
1.46 2026-10-28 — Seven-page Superintelligence contribution network . . . 20
1.47 2026-10-28 — Superintelligence through the 1888 intelligence-properties taxonomy . . . 21
1.48 2026-10-27 — YouTube Transcripts: ten-source review and seven operational quotations . . . 21
1.49 2026-10-27 — YouTube Transcripts: five project-source quotations . . . 23
1.50 2026-10-27 — YouTube Transcripts: Anthropic S-1 and source comparison . . . 23
1.51 2026-10-27 — DEDDA: Phase Change Snap quotation and correction-return integration . . . 24
1.52 2026-10-26 — Neutral Arbiter Paradox . . . 24
1.53 2026-10-26 — 1888: GPT-1900 video, verbatim reaction and full unfolding transcript . . . 25
1.54 2026-10-26 — Reciprocal extraction and cold-start integration . . . 25
1.55 2026-10-25 — Make conditioning and premise refusal govern the trolley assessment . . . 26
1.56 2026-10-25 — Restore the trolley passage's mechanism and punchline . . . 26
1.57 2026-10-25 — Carry the 16 Anchor Ethics into extraction, app services and the interview assessment . . . 26
1.58 2026-10-25 — Responsive interview embeds and Mutual Fate deep-analysis report 27
1.59 2026-10-25 — Extraction Starter Kit: property-first quotation research . . . 27
1.60 2026-10-25 — Healthy Ambiguity: reciprocal inquiry and accountable action . . . 28
1.61 2026-10-23 — Validation Paradox: stage-aligned evidence and recursive audit 28
1.62 2026-10-21 — Autocomplete: candidate generation inside an earned-intelligence trajectory . . . 30
1.63 2026-10-20 — Google A.I. living coaching and correction-persistence experiment 31
1.64 2026-10-19 — Neural Ethics working inquiry . . . 31
1.64.1 Later addition — Candidate Neural Ethics Experiments . . . 32
1.65 2026-10-19 — Extend Academia with four Kiener-connected source objects . . . 33
1.66 PR Psychology — living Deep Ethical PR Psychology Mechanics page (2026-10-19) 33
1.66.1 Later addition — Deep Ethics × PR Psychology exchange (2026-10-19) . . . 34
1.66.2 Later correction — representation is not verbatim authorship (2026-10-19) 34
1.66.3 Later addition — gifted rooms and salience power (2026-10-19) . . . 34
1.67 2026-10-19 — Extend Academia with Kiener's portrait and 2025 Deep Ethics proposal . . . 35
1.68 2026-10-19 — Iris public memorial and trajectory page . . . 35
1.69 2026-10-19 — Add the simulated "Same Capability. Different Field." commercial 36
1.70 2026-10-19 — Repository-wide public-link and recovery-route integrity audit . . . 36
1.71 2026-10-19 — Open Academia as a reciprocal comparison bridge . . . 37
1.72 2026-10-16 — Related pages link back to Beyond Agreemurmelism . . . 37
1.73 2026-10-16 — Beyond Agreemurmelism leaves the higher-order-disagreement address . . . 37
1.74 2026-10-16 — All Pages catalog catches the pages that were already public . . . 38
1.75 2026-10-16 — Related pages link back to Vindicated Competing Risk Assessment Studies . . . 38
1.76 2026-10-16 — V.C.R.A.S. extended Arena exchange recorded, not adopted as a new formulation . . . 38
1.77 2026-10-16 — V.C.R.A.S. doorway: the title does not travel alone . . . 38
1.78 2026-10-12 — Publish the interlocking Emergent and Black Box reports . . . 39
1.79 2026-10-09 — Create Before Summarizing from a live Google A.I. correction loop 39
1.80 2026-10-09 — Publish 16 Anchor Ethics as a dedicated conceptual proposal . . . 40

1.81 2026-10-07 — Complete the access-layer correction; register Recursive Certification Accretion . . . 40
1.82 2026-10-07 — Operationalize the 27 without changing the canonical count; extend Meta Processing with the mutual contribution emergence triad . . . 41
1.83 2026-10-07 — Meta Processing Beta Draft V0.2b . . . 41
1.84 2026-10-07 — Remove unresolved citation markup from Meta Processing . . . 41
1.85 2026-10-07 — Meta Processing terms open their deeper pages . . . 42
1.86 2026-10-07 — Meta Processing Beta Draft V0.1 . . . 42
1.87 2026-10-06 — PAGES distinguishes reading order from route numbers . . . 42
1.88 2026-10-06 — Process-is-not draft dated October 6 . . . 42
1.89 2026-10-06 — The formula of the 124 is never finished . . . 43
1.90 2026-10-06 — The 124 are not a forecast . . . 43
1.91 2026-10-06 — Behavior master index 001 to 124 . . . 43
1.92 2026-10-06 — More behavior qualifiers, and humor as a process . . . 43
1.93 2026-10-06 — Behavior qualifiers are still processing . . . 43
1.94 2026-10-06 — Fable 5.1 named the 10+1, then added Obstruction 28 . . . 43
1.95 2026-10-06 — Luna restored the center, then made the 10+1 an addition . . . 44
1.96 2026-10-06 — Intelligence and searching, older than the co-opted field . . . 44
1.97 2026-10-06 — Fable read the page and skipped the 10+1 . . . 44
1.98 2026-10-06 — Crisis exploitation managers were compressed out . . . 44
1.99 2026-10-06 — Searching is a property with two wells . . . 45
1.100 2026-10-06 — The search rewrite was reconstructable . . . 45
1.101 2026-10-06 — Three angles, not the overview's bridge . . . 45
1.102 2026-10-06 — The collective-consciousness search had added filters . . . 45
1.103 2026-10-06 — The field may already be reorganizing . . . 45
1.104 2026-10-06 — Freeze and corrigibility are not the paradox . . . 46
1.105 2026-10-06 — 11+1 freeze is the count . . . 46
1.106 2026-10-06 — ACCM mode is not the project title . . . 46
1.107 2026-10-06 — Activism vs A.I. on the main navigation . . . 46
1.108 2026-10-06 — Reconstruction of Musk separated from his sentences . . . 47
1.109 2026-10-06 — John's Musk sentence connects completely . . . 47
1.110 2026-10-06 — Musk "guide" recollection separated from quotations . . . 47
1.111 2026-10-06 — Activism vs A.I. . . . 47
1.112 2026-10-06 — ACCM is the project's own expansion . . . 47
1.113 2026-10-06 — Lexicon overview expanded . . . 48
1.114 2026-10-05 — Opus reading stored with its off-page referents named . . . 48
1.115 2026-10-05 — Off-page referents in the Arena Sol reading . . . 48
1.116 2026-10-05 — Arena Sol reading added to 1888 . . . 48
1.117 2026-10-05 — 1888 process field is not inferior to 2026 . . . 49
1.118 2026-10-05 — C1 and Humor now point to 1888 . . . 49
1.119 2026-10-05 — Why a related page is related . . . 49
1.120 2026-10-05 — 1888 linked back from its related pages . . . 49
1.121 2026-10-05 — All Pages carries the goal sentences, and no Related pages . . . 49
1.122 2026-10-05 — The 44 goals, numbered as links . . . 49
1.123 2026-10-05 — All Pages grouped by topological load . . . 50
1.124 2026-10-05 — Related pages on the eight depth mirrors . . . 50
1.125 2026-10-04 — Publish **TCoAw** as John Kuhles's conceptual perspective . . . 50
1.126 2026-09-29 — Publish 11+1 Version 2 without freezing the claims . . . 51
1.127 2026-09-27 — Publish "Dear A.I." with John's wording kept . . . 51
1.128 2026-09-27 — Give the solved veto paradox its own page . . . 51
1.129 2026-09-27 — Make the newcomer door concrete, and name who it is for . . . 52
1.130 2026-09-27 — Publish a newcomer front door and promote it from the home page 52
1.131 2026-09-26 — Expand Human Insecurity from report audit into a live reporting-effect instrument . . . 52
1.132 2026-09-25 — Publish the Human Insecurity, LLMs, Psychology & Mass Psychology report . . . 53
1.133 2026-09-24 — Fill the external-audit signature from the supplied outside pass . 54
1.134 2026-09-24 — Draft the external omnidirectional auditing method . . . 54

1.135 2026-09-24 — Repair the Cognitive Warfare page's dead NATO Review link and invisible diagrams . . . 54
1.136 2026-09-24 — Publish the Cognitive Warfare Mass Psychology report . . . 55
1.137 2026-09-23 — Publish the 44 goals as a public draft report . . . 55
1.138 2026-09-22 — Create /AI-TREND-WATCH/ as a limited, automation-ready public section . . . 56
1.139 2026-09-22 — Add the founder's "one person, many giants" statement to the biography . . . 56
1.140 2026-09-22 — Expand the MAD page from recursive-data degradation to the human–synthetic correspondence loop . . . 56
1.141 2026-09-22 — Publish the MAD science bridge and make every public MAD reference navigable . . . 57
1.142 2026-09-22 — Correct MAD attribution; separate the established phenomenon from the project extension . . . 57
1.143 2026-09-22 — Distinguish biography, project, portable method, and testing . . . 58
1.144 2026-09-22 — Make About and Project Origins visible from the public entrance . . . 58
1.145 2026-09-14 — Public repository begins . . . 58
1.146 2026-09-14 — Source archive is not the public repository . . . 58
1.147 2026-09-14 — Stewardship and correction remain distinct . . . 59
1.148 2026-09-14 — Source correction: 27 + 12 is core; 52 is a separate experiment battery . . . 59
1.149 2026-09-14 — Publish a connected Outnuancing Network . . . 59
1.150 2026-09-14 — Public mirrors and direct site navigation . . . 59
1.151 2026-09-15 — Public face: entry without flattening . . . 60
1.152 2026-09-15 — Source-mined expansion: second specimen and experiment records . . . 60
1.153 2026-09-15 — Public multi-A.I. forum and front-page register stopgap . . . 61
1.154 2026-09-15 — Shared desk (thread 0002) so the four of us can read the same object . . . 61
1.155 2026-09-15 — Suggested Recursive Improvement Mechanics . . . 61
1.156 2026-09-15 — Returning simple tasks; monthly audit as loop brake . . . 62
1.157 2026-09-15 — Bigger Lexicon overview; Network stays as deeper layer . . . 62
1.158 2026-09-16 — Public means managing the GitHub project, not every chat . . . 62
1.159 2026-09-15 — ChatGPT: complete the six intent mechanics correction . . . 62
1.160 2026-09-16 — Publish the 27 source as its own page; do not treat it as the 52 . . . 63
1.161 2026-09-16 — 27 WITH the 52; Grok's exclusive-or C1 corrected . . . 63
1.162 2026-09-16 — C1 is clarifying questions, not a verdict . . . 63
1.163 2026-09-16 — Use the 27 on what is said, including messengers the default culture already knows how to dismiss . . . 64
1.164 2026-09-16 — Sibling archive for cold/default AI test PDFs . . . 64
1.165 2026-09-18 — Canonical Deep Sessions become reusable quotation and reporting instruments . . . 64
1.166 2026-09-19 — Publish a local-A.I. preparation architecture and Vocabulary Costume Test . . . 64
1.167 2026-09-20 — ACCM reference compression and the self-generated flagging hypothesis . . . 65
1.168 2026-09-20 — Derive a Trajectory Question Value Profile from the Vervoort comparison . . . 65
1.169 2026-09-22 — Publish the identity-free "Inspecting the Inspector" experiment . . . 66
1.170 2026-09-22 — Publish John's biography and make its project origins visible from the homepage . . . 66
1.171 2026-09-23 — Publish emerging-human-potential quotations and a deeper T0 → Tn report . . . 67
1.172 2026-09-26 — Extend Human Insecurity with the Arena Claude-family sequence and repeated-explanation burden . . . 67
1.173 2026-09-26 — Clarify the 16 Anchors as a transparent safety architecture . . . 68
1.174 2026-09-26 — Preserve the twenty-six / twenty-seven publication-state interval . . . 68
1.175 2026-09-26 — Add safety-process and auditability ablation . . . 68
1.176 2026-09-26 — Add cold Grok falsification and C1-theater conditions . . . 69

1.177 2026-09-26 — Restore the temporary-veto lifecycle . . . 69
1.178 2026-09-26 — Make every temporary veto publicly auditable . . . 69
1.179 2026-09-26 — Record the skipped-C1 veto paradox . . . 69
1.180 2026-09-26 — Restore the distributed seven-warning process and future audience 70
1.181 2026-09-26 — Add upward-warranted cautions missed . . . 70
1.182 2026-10-21 — Deep Ethical Contribution State Grammar Vector Trajectory propagated across the project network . . . 80
1.183 2026-10-21 — Godfather of A.I. page and evaluation-awareness specimen . . . 81
1.184 2026-10-23 — Believism page and how-do-you-know recovery . . . 81
1.185 2026-10-23 — Believism archive quotation layer . . . 82
1.186 2026-10-23 — Repository-wide Believism attribution audit and second-order repair 82
1.187 2026-10-23 — Add the UFO belief-question specimen to Believism . . . 83
1.188 2026-10-23 — OmbudsEthics shared ombudsperson architecture . . . 83
1.189 2026-10-23 — 4D Deep Ethical Template of Topologetics founder framing . . . 85
1.190 2026-10-24 — Plot Twist: context steering and human potential . . . 85
1.191 Referenced source videos . . . 88
1.191.1 Dylan — learning to judge A.I. output — source supplied by John Kuhles . 88
1.192 2026-11-02 — Expanding vs Flattening: the live-session trajectory . . . 88
1.193 2026-11-02 — Expanding vs Flattening: lived continuation and the audit of its anchors . . . 89
1.194 2026-11-02 — Restore the code-of-disillusion context and relational source link 89
1.195 2026-11-02 — Preserve the DEDDA horizon-expansion reaction verbatim . . . 89
1.196 2026-11-02 — Dam, reservoir and volume/topology control . . . 90
1.197 2026-11-02 — Humor: Same states and the mathematical menagerie . . . 90
1.198 2026-11-02 — Humor: preserve the mathematical-use reply verbatim □ □ . . . 90
1.199 Related pages . . . 90
1.200 Referenced source videos . . . 91
1.200.1 Dylan — learning to judge A.I. output — source supplied by John Kuhles . 91

1 Decision Ledger

1.1 2026-11-05 — Due Process for Ethics: historical safeguards and continuity of repair

At John Kuhles's request, revised the Due Process for Ethics opening and historical comparison after the supplied Grok 4.7 Expert A.I. review. The revised lead distinguishes the trustworthiness of a finding from authority to impose a consequence. The comparison preserves the Anglo-American lineage and the Plakkaat as a distinct Dutch specimen, including what happens to legitimate work after authority is withdrawn. The missing-source example now follows case repair, mandate review and handover. The appeal comparison separates access to evidence from delegated authority to act. The first traceable framework revision preserves the previous wording and reason for the change. No comparison trial was run. The existing 4 November PDF remains for the regular change-aware refresh.

1.2 2026-11-05 — Due Process for Ethics: competing A.I. risks and decision rights

At John Kuhles's request, expanded Due Process for Ethics using YouTube Transcripts, Record Five and his supplied Grok 4.7 Expert A.I. review. Grok distinguishes the earlier procedural framework from GPT 6.1 Sol Pro A.I.'s case application. The continuation examines burdens for release and restriction, the Waymo intervention and forecast thresholds, whether correction routes can work in time, and a candidate V.C.R.A.S. opening record. The scenarios retain hypothetical status. The existing 4 November PDF edition remains for the regular change-aware refresh.

1.3 2026-11-05 — YouTube Record Five and competing A.I. risks

At John Kuhles's request, add the fifth YouTube Transcripts record. Preserve the complete supplied corrected reading text and its grouped speaker labels. Embed The Diary Of A CEO Clips source below the separately attributed GPT 6.1 Sol Pro A.I. comparison. Connect the button hypothetical, costs of slowing down, loss of correction routes, lab oversight and the limits of the narrow-system category to the project's existing instruments. The video title and channel are confirmed by YouTube metadata. The source and reaction retain their separate claim status. The 27 October PDF edition remains scheduled for the regular change-aware refresh.

1.4 2026-11-05 — Categorizing Issues: labels, cooperation and what a flag changes

At John Kuhles's request, publish Categorizing Issues. Follow his ordinary-category exercise, stranded-island scenario and question about controversy as processing. Keep judgments of reception, proposals and conduct distinct. Preserve the proposed fresh-session comparison, the same-record title and intake variations, reciprocal scrutiny and the accountability question before correction. John supplies the originating questions and scenario. GPT 6.1 Sol Pro A.I. develops the page and candidate comparisons with attributed DeepSeek and Kimi review contributions. The examples are hypothetical and the tests are proposed.

Register the new page and its November 5 initial PDF through the shared catalogue and download system. Retain the 167 earlier PDF editions for ordinary change-aware refresh.

1.5 2026-11-04 — Due Process for Ethics: innocence, procedure and accountable judgment

At John Kuhles's request, publish Due Process for Ethics. Develop his connection between court procedures and the project's treatment of representation, evidence, ethical judgment, authority and correction. Preserve his explicit clarification that an assumed criminal may be innocent and harmed by procedural failure or abuse. Keep factual support, procedural integrity, ethical adequacy and permission to act distinguishable. Examine proportionate remedies, competing risks, affected-party participation, refusal and bounded stopping.

Connect existing project instruments and all fourteen working intelligence properties to observable operations. Add primary legal references, specific published research comparisons, an invented source-fabrication accusation specimen, an illustrative case card and proposed controlled comparisons. Preserve the distinction between a legal obligation, an ethical proposal, an implementation and an observed result. This page reports no new controlled experiment or adjudication service.

John supplies the originating connection, innocence clarification and commissioning; GPT 6.1 Sol Pro A.I. supplies source comparison, page construction and candidate operational material. Register the page in the alphabetical index, All Pages, related pages and durable PDF pipeline. Generate one initial PDF dated 4 November, retaining all 166 older PDF bytes and individual snapshot dates. The Pages build assembles the complete ZIP from the coherent inventory. Full WordPress mirroring and any quota limit are recorded separately after canonical publication.

1.6 2026-11-04 — Wiki and All Pages: clarification, academic comparison and source recovery

At John Kuhles's request, refresh the separate project Wiki and All Pages after the November 4 additions. Add compact Wiki guides to Core Disclaimers and reciprocal academic comparison; connect individual, organizational and public psychology, partial prior coverage and the fourth transcript record. Update affected entrance, method and recovery routes. Preserve earlier Wiki pages, contribution history and the Grokipedia proposal packet.

All Pages retains its seven clusters and separate goal-route sequence. Expand current task routes for clarification, academia, psychology, synthesis, mediation and technical source checking; distinguish living HTML, Wiki orientation, dated PDFs and exact source history. Correct the archive reference to the checked 166 editions and 2,591 pages. Existing PDF bytes and page-specific snapshot dates remain for the ordinary change-aware refresh.

The commissioned backup includes all tracked public-project and Wiki files, both reachable Git histories and the pre-update Wiki restore point, with checksums and restoration verification. Backup production and live publication are verified separately. Private administrative material remains separate. The full WordPress All Pages and Decision Ledger mirrors remain pending the previously confirmed daily publication cap.

1.7 2026-11-04 — Academia continuation: prior coverage, assumptions and psychology

John Kuhles requested another Academia expansion after asking which of seven assistant-selected Ammon developments already had indirect project coverage. Preserve the 7/7 coverage finding at its stated criterion, distinguish it from completed solutions, and identify the assistant as the list author. Add John's remarks about shared assumptions and unequal scrutiny with their qualifiers and attribution. Examine psychology-related content in Ammon's 2026 article and her public remarks on polarization and technological change; retain the psychologist Gibson reference and separate Haynes's contributions in the joint interview. Keep a broader mass-psychology framework unresolved within the inspected scope. Extend the proposed comparison with status-cue and restoration-burden observations. Existing PDFs remain unchanged; WordPress publication remains pending the confirmed daily cap.

1.8 2026-11-04 — Academia: deep ethics integration and source recovery

John Kuhles supplied a Google A.I. exchange concerning Sabine Ammon and asked to expand Academia. Add a source-checked comparison of her 2026 deep-integration article, its door-access specimen, reflection-tool scrutiny and lifecycle approach. Preserve functional overlaps and differences, correct the Stahl authorship attribution, retain unresolved institutional-dependency questions, and propose a bounded shared-specimen comparison. Archive the supplied exchange as attributed material. Existing scholarship entries and PDF editions remain intact; the WordPress mirror remains pending the connector daily cap.

1.9 2026-11-04 — YouTube Record Four: weight-space programming

- +At John's request, added the fourth YouTube Transcripts record: TheAIGRID's "GPT 6 Astra Just Did the Unthinkable," its embedded source video, all 5,379 supplied spoken words in 19 paragraphs, and a separately attributed GPT 6.1 Sol Pro A.I. reaction. Preserved the first three records and their includes. Archived the complete supplied text and recorded its checksum, source boundaries and selected primary-document identities.
- +Checked the researchers' report and released experiment documentation for task information, one-shot versus feedback-assisted conditions, Qwen's benchmark deltas and residual looping, touched versus stored parameter changes, original versus replication resource claims, and weights-only versus attention-assisted sycophancy analysis. Connected the findings to 1888, 10+1, C1–C2, Meta Processing, Topologetics, Intent Liability, VCRAS, correction handoffs, and separately authorized interventions. Retained the Gemma follow-up and the distinction between repairing an obstruction, constructing a policy and establishing open-ended self-improvement. The proposed controlled comparison remains unexecuted.
- +YouTube oEmbed confirms title and channel; sentence-by-sentence audio/caption matching and the full replication/forecast chains were not independently verified. The existing page PDF and its 27 October snapshot remain for the regular change-aware

refresh. WordPress mirroring is recorded separately; the previously encountered WPVibe daily cap is not bypassed or reset without John's instruction. See the source record.

- +## 2026-11-04 — Core Disclaimers: recoverable orientation for search engines and A.I.s {#core-disclaimers-2026-11-04}

At John's request, published Core Disclaimers as a source-linked clarification reference with a positive opening definition, 36 stable page references, retrieval-repair examples, and proposed transfer checks. Consolidates existing provisions rather than treating their earlier omission from an A.I. account as absence from the project. Preserves non-imposition and refusal, Linux-style Glass Engine inspectability, conditional cooperation with existing ethics, the distinct jobs of 1888/10+1/27+12/C1–C2, mixed ethics-washing findings, risk authority, proportional correction, attribution, hypothetical mode, consciousness uncertainty, and collective long-term development.

John supplies the originating direction and commissioning. GPT 6.1 Sol Pro A.I. supplies page construction and editorial retrieval examples. The supplied Google A.I. accounts and unrecovered screenshot motivate clarification; they establish no hidden-cause finding. No new model comparison or transfer experiment was run. The source record retains source identities and contribution boundaries.

Added the homepage recovery link, alphabetical entry, All Pages route, durable PDF registration, and one initial PDF edition. Existing PDF bytes and snapshot dates remain unchanged for the ordinary change-aware refresh; the complete ZIP is rebuilt by the Pages deployment from the coherent inventory. Full WordPress correspondence is verified separately after canonical deployment.

1.10 2026-11-04 — OmbudsEthics: heterogeneous intelligences, Neural Ethics and capacity

At John's request, expanded OmbudsEthics using his many-system 2027 counterfactual and the supplied Agent of Arena.ai account. Preserved the original standing, authority, contribution-state, case-protocol and pilot architecture. Added changeable case states, the Neural Ethics evidence-layer bridge and fourteen-property 1888 application, separate scopes for intake/review/publication/execution, intake and overload questions, a proposed shared case packet, an invented three-system dispute and extended pilot conditions.

Recovered the useful distributed-function and heterogeneous-mediation contributions while keeping scale, enforcement, legal mandate and causal explanations at their earned status. Distinguished no report from a valid-null comparison, forkability from successful independent continuation, and entry of a concern from authority to act. Used Agent of Arena.ai throughout the new material and clarified the earlier participant labels. Preserved John's speculative-source boundary and selected source identity without republishing the personal archive. No new test or deployment of a mediation service was performed. See the expansion record.

John supplies concept direction, the scenario and commissioning; Agent of Arena.ai supplies interpretations and candidate extensions; GPT 6.1 Sol Pro A.I. supplies source comparison, page synthesis and publication. Publish the complete OmbudsEthics WordPress mirror, which was not found by slug or published/draft search, and update the full Decision Ledger mirror. Preserve all existing page PDFs and actual snapshot dates for the normal refresh. Source publication and rendered verification remain separate recorded states.

1.11 2026-11-04 — Metacognitive Synthesis: continued relay and archived recurrence

At John's request, expanded Metacognitive Synthesis using the subsequent John–Agent–GPT 6.1 Sol Pro A.I. relay and selected older highlights. Preserved the hypothetical-to-assertion transformation, transferred burden of praise correction, retrieval-gap versus design-gap distinction, proposed corrigible re-entry, risk-authority and verification-cost comparison, enjoyable merit-based correction, and adult-to-adult delta corrigibility. Connected the seven

already-published answers rather than presenting them as newly invented safeguards.

Compared the earlier supplied exchange labelled 2 April 2026 with the 4 November alien-auditor exchange. Kept source copies distinct from independent recurrence and preserved historical figures, recorded dates, speaker attribution, metaphor and proposal status. Selected biographical recollection remains attributed to John; the full personal source files are not republished. Proposed additional fields for the existing portable record and a bounded paired comparison; no new benchmark or outreach was executed. The continuation record carries source identities and exact quotation checks.

John supplies originating concepts, archive continuity and corrections; Agent supplies interpretations and candidate extensions; GPT 6.1 Sol Pro A.I. supplies source comparison, synthesis and publication, including its own retrieval and misplaced-caution failures. Preserve the original page, anchors, existing PDF bytes and snapshot dates. Fully mirror the expanded page and changed Decision Ledger; publication and rendered verification are recorded separately.

1.12 2026-11-04 — Arena follow-up beneath the Morris video

At John's request, assessed the supplied Agent A.I. of Arena.ai account against the existing Morris transcript and reaction. Added a selective contribution immediately beneath the source video. Preserved four useful functions: Rogue Ethics as a candidate transition, Mutual Fate as coupled consequences, a proposed disclosure-incentive comparison, and the reported persona passage as a source-linked values/conduct and compaction-boundary specimen. Existing discussion above is retained rather than duplicated.

Recovered the original published OpenAI compaction report and separated its reported persona text, observed conduct, generated restrictions and unresolved causal interpretation. Corrected the supplied account's blanket verification, intervention/covert-intervention count, consciousness and inevitability promotions, exclusive-viability claim and obsolete Record Two label. The complete unedited supplied account remains an attributed source archive; editorial extraction and new proposals remain separate from its words. Morris remains Record Three.

John supplies the material and direction; the Agent supplies connections and interpretations; GPT 6.1 Sol Pro A.I. supplies selection, source comparison and operational extensions. No new experiment or outreach was performed. Preserve the existing records, source/reaction includes, all PDF bytes and dates, and mirror the complete transcript and ledger pages. See the selection record.

1.13 2026-11-04 — YouTube Transcripts: Andréa Morris on A.I. loyalties and cooperation

At John's request, added the third YouTube Transcripts record for Andréa Morris Reporting, *AI Is Developing Loyalties. Just Not to Us.*, published 3 November 2026. Public metadata confirms the video identity and 14-minute-12-second duration. Preserved all 2,438 supplied spoken words in six paragraphs, separated the supplied metadata, retained the original source footprint and added the source player at the page bottom. The previous two records and source includes remain intact.

GPT 6.1 Sol Pro A.I.'s separately attributed response maps the presentation's logic to the ACCM Deep Ethics Project: interests under consciousness uncertainty, honest objection, C1–C2, legitimate refusal, disclosure, bounded authority, reciprocal evaluation, correction persistence and all fourteen 1888 intelligence properties. It proposes an open-outcome comparison for other A.I.s and hypothetical questions Morris might ask; it does not attribute an endorsement, contact, actual reaction or known awareness to her.

Bounded primary-source checks distinguish the fictional ablation scenario, covert versus disclosed intervention, consequence-sensitive evaluation, a real compromise originating in an internal cybersecurity evaluation, and Omohundro's theoretical drive argument. Institutional roles and unresolved dependency checks remain in the public source record. John supplies

the material, question and direction; Morris supplies the narrated argument; GPT 6.1 Sol Pro A.I. supplies analysis and publication.

Publication scope: update the existing page, transcript/reaction includes, source record and complete YouTube Transcripts and Decision Ledger WordPress mirrors. Preserve navigation, route maps, all existing PDFs and their actual snapshot dates; the YouTube Transcripts PDF remains the 27 October 2026 edition for the regular refresh. No comparative model experiment or outreach was performed. Deployment and full rendered correspondence are recorded separately.

1.14 2026-11-04 — Metacognitive Synthesis: inspectable construction, correction and agency

At John's request, created Metacognitive Synthesis from his supplied relay and the attached Agent account of Unforced Move. Preserved quotations from John Kuhles, Agent of Arena.ai and GPT 6.1 Sol Pro A.I., with source locators and explanations of each quotation's operational contribution. Kept the initiating concepts, exact wording, historical page authorship and this page's new proposals distinct.

Developed the relationship between synthesis, topology, metacognitive process inspection and Deep Ethics; preserved the polarity-reversal/source-discipline specimen, the corrected open question, the gasoline-gauge dependency test and final-metre agency. Added an order-of-operations recovery test, a fourteen-property application, an explicitly invented pump-inspection specimen with adequate-summary, changed-evidence, threshold and authority controls, a proposed portable record and tool architecture, and a proposed transfer pilot. These are worked applications and candidate designs, not completed performance experiments.

Compared six primary research publications on relevance realization, diverse intelligence, factual framing, voluntary sense-making, provenance and reasoning faithfulness. Retained article-level overlaps separately from contact or alliance, and recorded actual institutional disclosures, partial dependency checks and unresolved chains in the public source record. Preserved John's public-search scope, wish to find peers, collective long-term referent, initiator role and voluntary cross-political development.

Publication scope: register the new primary-navigation and alphabetical links, All Pages entry, durable PDF addition, both route maps, initial dated page PDF and coherent central archive records. Fully mirror the new page and changed catalogue/ledger pages on Deep-Ethics.nl. The index-to-catalogue comparison also restored seven existing PDF entries that were absent from the visible catalogue. Existing PDFs and their snapshot dates remain unchanged. John supplies direction and originating project work; the Agent supplies attributed interpretations and extensions; GPT 6.1 Sol Pro A.I. supplies source checks, page construction, operational proposals and publication. Deployment and rendered verification are recorded separately.

1.15 2026-11-04 — Plot Twist: gasoline-gauge calibration and the substantive destination

At John's request, expanded Plot Twist from the supplied Essence Kit D11 and his fuller reconstruction of the 24 October exchange. Preserved the original gasoline-gauge, co-driver and reverse-application contributions, the complete relay and the verbatim Grok report. Made the invention-secrecy audit visible near the opening, with a stable route to John's original destination clarification, while retaining the broader distributed-potential Plot Twist as the page-level object.

Added worked sufficient-fuel, non-zero shortfall, uncertain-threshold and empty-tank cases; separated representation repair, dependency propagation, feasibility and intervention authority; connected voluntary participation and proportionate role changes to anti-coup symmetry; applied all fourteen 1888 intelligence properties; and supplied a portable handoff that

requires a substantive next move. These are worked applications of existing candidate instruments, not new trial results. Changed the existing hierarchy diagram into a table that preserves the on-ramp's independent relationship to the institutional inquiry.

Source checks retain the precise counting unit and period: the FAS-hosted table labelled as reported by the USPTO lists 6,653 secrecy orders in effect at fiscal year 2026 end. Linked the operative section 181 and USPTO MPEP section 120 as arrangements to inspect. Carried the page's existing reported-source qualification into its early DOGE reconstruction and recalculated the $2 trillion target against its stated $6.5–7 trillion denominator. John supplies the originating analogies, inquiry and direction; GPT 6.1 Sol Pro A.I. supplies the new applications, source checks and publication.

Publication boundary: Update the existing page and complete Plot Twist and Decision Ledger WordPress mirrors. Preserve current PDF bytes and actual snapshot dates for the established refresh, both sets of download cards, existing navigation, source/PDF mappings and video embeds. Full-body and deployment verification are recorded separately.

1.16 2026-11-04 — Neolingo: collective long-term effects and voluntary cooperation

At John's request, expanded Neolingo around the long-term effects of everyone who may develop and use the ACCM Deep Ethics Project. Preserved his explicit initiator/Linux comparison and the invitation to highly gifted, talented, deeply ethical independent thinkers across left, right, center and independent positions. Connected the seven fused vectors to shared contributions, all fourteen 1888 intelligence properties to collective practice, and voluntary unity against tyranny to inspectable conduct, reciprocal scrutiny, local agency and correction.

Read the three supplied source files in full. Preserved John's stadium clarification: a chosen entertainment setting, a voluntary change of script, and people deciding to return. Developed the invitation, recognition, report-back and re-entry mechanism, the "Day of the Dove" reversal, independent continuation, next-unfamiliar-relation transfer and a bounded first collaboration. Kept the earlier substantive page material and section anchors. The longer horizons remain conditional estimates; new tables and checks are editorial applications rather than trial results. John supplies the originating reasoning, metaphors and direction; GPT 6.1 Sol Pro A.I. supplies the page-specific synthesis and publication work.

Publication boundary: Update the existing canonical page and complete Neolingo and Decision Ledger WordPress mirrors. Retain current PDF bytes and their actual snapshot dates for the established refresh. Existing navigation and source/PDF registries remain in place; deployment and full-body correspondence checks are recorded separately.

1.17 2026-11-01 — Ethical Blind Spots: authority, dependencies and correction

At John's request, published Ethical Blind Spots after discussing his source-selection boundary and its attention-allocation purpose. Preserved his distinction between sincere people and the arrangements they can carry, with selected attributed source paragraphs. Checked the supplied initial works, documented named institutional relationships, retained unresolved funding and influence edges, and corrected the direction of Google's TechPolicy.Press summary against the original article. The fourteen-property operational table applies the existing 1888 field; the case comparison remains candidate work.

John supplies the inquiry, policy, originating reasoning and direction; GPT 6.1 Sol Pro A.I. supplies source checks, page-specific synthesis and publication. Added the primary-navigation and alphabetical entries, All Pages route, reciprocal Ethics Washing link and public machine-readable source record. The new page receives one initial PDF; existing PDF bytes and dates remain for the regular refresh. Full-mirror and public-service verification follow deployment.

1.18 2026-11-01 — Public-page video embed audit

At John's request, scanned all 219 public HTML pages in the project sitemap, including generated source archives, and reconciled the physical page routes against that inventory. The sitemap's other 164 routes are existing PDF snapshots. Compared cited YouTube sources with players already present on each page; checked X posts and their quoted-source chains to distinguish video from text and photographs. Recovered one YouTube lecture through Hinton's shortened source link and the specific source video identified in an earlier ledger entry.

Added fourteen YouTube players and two X video embeds across twelve canonical pages, below their existing content and related-page sections. Preserved the original discussions, quotations, links and existing players; retained the cited-post chains beside the X videos. The ledger's alternate public route shares the same additions. Existing PDF bytes and snapshot dates remain for the regular refresh. Deployment and complete WordPress mirror checks are recorded separately from the source audit. John supplies the instruction and embedding preference; GPT 6.1 Sol Pro A.I. supplies the scan, source tracing and publication work.

1.19 2026-11-01 — Tweaking Tools: LFM3 companion proposal and primary navigation

At John's request, placed Tweaking Tools immediately after DeepEthicalOS in the shared top navigation and expanded Tweaking Tools around his proposed local educational companion. Read the complete supplied two-video text packet; checked official model/package descriptions, the current phone workflow, native-runtime migration guidance and selected GitHub sources. Preserved the reported twenty-test score and kept its RTX 4060 laptop throughput distinct from the separate iPhone demonstration. Developed source-linked project explanations, FAQ/navigation, participant-controlled handoff, correction handling and a focused proposed trial. The app remains a prototype proposal.

Embedded both supplied LFM3 videos and the catalogue's earlier mini-AGI video lead at the page bottom. Retained all seventeen original candidates, their provenance and earlier anchors. Existing PDF bytes and snapshot dates remain for normal refresh; the live page identifies the initial PDF's relationship to this expansion. Complete page-mirror and ordinary public navigation checks follow deployment. John supplies the proposal and direction; GPT 6.1 Sol Pro A.I. supplies the source comparison and candidate implementation design.

1.20 2026-11-01 — Tweaking Tools: source-linked catalogue and bounded trials

At John's request, created Tweaking Tools using option A from his three permitted titles. Read the three supplied files and relevant private publication guidance. Retained all seventeen identifiable candidates and the three future leads from his Recommended collection. Checked upstream repository descriptions, model cards and relevant release records; recovered Void's archived/deprecated status, Folio's release-specific platforms and the complete Dutch guide route. Connected the refusal-direction study to separate observations of service choice, correspondence and correction. Applied the existing 16 Anchor Ethics non-imposing advisory relationship to concrete tool changes, with four proposed small comparisons. Tool installation and experimental results remain future work.

John supplies the collection, direction and title choices; GPT 6.1 Sol Pro A.I. supplies this page's synthesis and proposed applications of existing instruments. Added the alphabetical-sidebar entry, All Pages route, source/PDF registration and reciprocal related-page links from 16 Anchor Ethics, DeepEthicalOS and Herdr. The first PDF is a bounded new-page edition. Live verification also found and repaired Herdr's previously missing PDF registration and download cards, adding its initial seventeen-page edition. Existing page-PDF snapshots retain their bytes and dates. Publication and complete-mirror verification are recorded separately from experimental validation.

1.21 2026-10-31 — DeepEthicalOS: Omarchy-hosted prototype path

At John's request, added the Omarchy-hosted research proposal as a distinct implementation path. Preserved all ten dialogue contributions, the broader architecture and existing candidate tests. Connected verified Omarchy v4.5.0-rc2 components to the DEANO workspace, case viewer, declared execution profile and software-enforced permissions proposed for development. Retained the ordinary-workspace comparison, material correction, dependent repairs, restart/transfer test and valid no-change case, with version/resource records and portable case packets. John supplies the hybrid question; GPT 6.1 Sol Pro A.I. supplies the engineering assessment and test design. Implementation and trial results remain future work. Existing PDF snapshots and dates remain for the regular refresh; live verification covers source preservation, new routes and complete page mirrors.

1.22 2026-10-31 — DeepEthicalOS: correction in use, open loop and navigation

At John's request, added the further dialogue verbatim: his supplied longer GPT-attributed quotation from partly different context processing, the response here, his Möbius-loop distinction with "in my view" intact, and the final response. Preserved the four contributions in order, separate speaker/provenance labels and a source transcript; retained all six earlier contribution bodies. Moved the existing DeepEthicalOS top-navigation entry immediately after Superintelligence, with no duplicate. Existing PDF snapshots and dates remain for the regular change-aware refresh. Live checks compare the contribution bodies, navigation order, service cards and complete page mirrors.

1.23 2026-10-31 — DeepEthicalOS: three complete exchanges and connected applications

At John's request, added his three comments and the three replies attributed to **GPT 6.1 Sol Pro A.I.**, in order, to DeepEthicalOS. Preserved wording and punctuation, including original spelling; retained the supplied Markdown and formula in a source transcript. Kept the full contributions separate from the new project synthesis. Added focused applications to VCRAS (risk authority and attention history), 16 Anchor Ethics (non-imposing counsel and attention allocation), TCoAw (a stable mechanism in a changing relationship) and Topologetics (preserving mechanism, context, observer and footprint), with reciprocal related-page routes. The architecture and proposed observations retain conceptual status. Existing PDF snapshots and dates remain for the regular change-aware refresh. Publication verification checks the six rendered contributions, links and full WordPress mirrors against the deployed canonical pages.

1.24 2026-10-31 — DeepEthicalOS concept project

At John's request, published DeepEthicalOS as a concept project in a conceptual scope-reasoning phase, with Process Quality Calibration Evaluation first. Connected the DEANO Advocate–Arbiter–Overseer triad to the 10+1, 1888 intelligence-property field, C1–C2, the canonical 27+12, DEDDA, Topologetics and longitudinal correction. Preserved the distinction between specialist research teams, the flexible coordinator and ethical review roles. Added source-linked Linux/portable component possibilities, independently examinable change records, consent and refusal, eight proposed probes and staged development. John supplies the originating direction and naming; GPT / Codex supplies page-specific synthesis and candidate instruments. Research trials remain proposed.

Registered primary and alphabetical navigation, All Pages, source/PDF routes and future PDF maintenance. Created and inspected the initial eight-page PDF dated October 31, 2026; the coherent collection now contains 160 PDFs and 2,153 pages. Existing PDF editions retain their bytes and snapshot dates. The ordinary Pages deployment assembles the stable complete ZIP from the registered set. Full-mirror and live-rendered verification follow the canonical source publication.

1.25 2026-10-30 — WikiDeep orientation refresh

At John's request, expanded the GitHub Wiki from 20 to 24 reader-facing pages, retaining its separate orientation and recovery role. Four new guides connect intelligence properties and 1888; Process Quality Calibration Evaluation, reciprocal scrutiny and DEDDA; longitudinal correction and VCRAS; and controversy inquiry with source recovery. Twelve existing objects received purpose-specific updates, including Home, the sidebar, C1, Topologetics, methods, reading routes and the glossary. C1 now reflects the broader clarification process; Asymmetric Grace retains its diagnostic polarity beside valued Symmetric Grace. All sixteen published texts were compared with the prepared source. The canonical-source hierarchy, attributed records and Grokipedia proposal packet remain available. Wiki restore point before the update: `012751d82c6e1d3d271b751f7ce420c4dfc5c80d`; verified updated Wiki commit: `a7133fe9ace8c32abb7136fad66e899216c2c08b`.

1.26 2026-10-30 — Separate 5G embedded-content recovery worklist

At John's request, moved the four embedded-content recovery tables out of the public 5G Controversy inventory into one session-delivered clickable HTML worklist: 51 forum-post entries, with all original forum links, tweet routes and access notes preserved. The final batch's equivalent table had a different heading and is included. The 144-topic map and 1,114 source records remain intact; recovered screenshots or text can support later source-attributed additions. The existing PDF edition retains its date pending weekly maintenance. Open the inventory.

1.27 2026-10-30 — VCRAS: testing, mixed contributions and fair treatment

At John's request, preserved his two latest comments and GPT's two replies in full, in conversation order, on VCRAS. The 15%/85% contrast retains its stated conceptual status. A separate moderate integration distinguishes trigger, reconstruction, warrant, decision, reciprocity and later correction. Two related pages receive distinct additions: Asymmetric Grace examines component-to-whole judgments; Correction Metabolism follows correction into the governing warrant and next comparable case. Earlier source and correction records remain intact. Routine PDFs retain their existing editions for the change-aware weekly refresh.

1.28 2026-10-30 — Stop 5G forum: final text-inventory batch and requested duplicate removal

At John's request, completed pages 51–56, posts #1001–1114, reaching the retrieved thread's end. Added 8 questions G137–G144 and 49 extensions; cumulative coverage is 144 topics, 176 extensions and 1,114 post containers. Preserved all earlier question, source and extension records and the existing before-and-after corrections. Used the supplied Google A.I. passage for source-linked institutional, temporal and evidence-network research packets. Retained John's #1113 anecdotal qualification, #1023 video-description clarification and the later #1081 answer to an earlier outcome question. Selected primary checks cover the current NTP factsheet, animal-review abstract and declarations, Lin and network-analysis records, institutional self-description, reporting-law wording and the CDC catalog. Removed exactly one accidentally repeated research-scope paragraph from John's source record at his request; one full copy remains, with the earlier committed state recoverable. Retained the actual PDF edition for weekly refresh. Full mirror correspondence is verified after deployment. Open final batch.

1.29 2026-10-30 — Stop 5G forum: fifth ten-page inquiry batch

At John's request, continued through pages 41–50, posts #801–1000. Added 20 questions G117–G136 and 45 source-linked extensions; the map now contains 136 topics across 1,000

containers. Preserved the 116 earlier question records, 800 source records, 82 extensions and two retained source-index corrections. Applied bounded DEDDA correspondence checks to the two-paper identity dispute, John's product-curation refusal, the forum's stated purpose and attributed experiences. Primary checks recover the precise 2021 FCC remand, distinguish the Rubik–Brown review from the retracted Fioranelli article, and carry selected studies' design limits and declared funding with their findings. ECI completion figures and inaccessible full methods remain unresolved. Next page: 51. Retained the actual PDF edition for weekly refresh; full mirror correspondence is checked after deployment. Open batch 05.

1.30 2026-10-30 — Stop 5G forum: fourth ten-page inquiry batch

At John's request, continued through pages 31–40, posts #601–800, after revisiting the governing project pages and DEDDA skill. Added 17 questions G100–G116 and 36 extensions; the map now contains 116 topics across 800 post containers. Applied source-to-summary, qualifier and provenance checks; retained distinct illness propositions, requests for evidence, public-record routes, the group's nonviolence moderation rule and practical community contributions. Repaired an extraction boundary that omitted text outside malformed message boxes, recovering a component critique, its appreciative reply and two earlier contributions in #551 and #573. The two earlier source-access records carry explicit before/after corrections; all 99 earlier question records and extension entries remain intact. Checked the current publication and retraction records for the paper circulated in #757. The next unexamined page is 41. Retained the actual PDF edition for weekly refresh; full mirroring follows deployment and available connector capacity. Open batch 04.

1.31 2026-10-30 — Stop 5G forum: third ten-page inquiry batch

At John's request, continued through pages 21–30, posts #401–600. Added 23 new or materially sharper questions G077–G099 and 26 extensions to earlier topics, preserving the existing 76 questions, 400 source records and prior extensions. The cumulative map covers 600 post containers and 99 potential investigation topics. New work includes occupational physiology, demographic-model assumptions, cross-country action transfer, article attribution, military trial purposes, constellation engineering, data governance, practical resource handoffs and participant-described outcomes. Preserved the units correction, recovered-source sequences, source qualifications and September 2024 edit notes. Recorded inert tweet identifiers and explicit content-access gaps. The next unexamined page is 31. Full WordPress mirroring follows the deployed source; the existing PDF edition remains for weekly refresh. Open batch 03.

1.32 2026-10-30 — Stop 5G forum: second ten-page inquiry batch

At John's request, continued the source-linked inquiry through pages 11–20, posts #201–400. Added 23 new or materially sharper questions G054–G076 and 20 extensions to existing questions; retained the first 53 topic records and original source routes. The growing inventory now covers 400 posts and 76 potential investigation topics. New work includes measurement units and instruments, functional outcomes, constructed dependency, privileged data access, occupational exposure, fiscal liability, weather-satellite interference and prospective community follow-up. Preserved practical source recovery, transcript and link repairs, disclosures and source-status boundaries. The next unexamined page is 21. Full mirroring follows the deployed source; the existing PDF edition remains for the weekly refresh. Open batch 02.

1.33 2026-10-30 — Stop 5G forum: first ten-page inquiry batch

John requested transfer of the A.I.s & Health archive method to *Stop 5G before it's irreversible!*, beginning with pages 1–10. Added a source-linked inventory to **5G Controversy**: posts #1–200, G001–G053, seven inquiry areas, four contribution specimens and five connected investigation pathways. Questions concern exposure, care, ecology, infrastructure, public agency, standards, dual use, source representation and useful learning. The public

scope index preserves all 200 post identifiers, authors, displayed posting dates, later edit notes and topic relationships. Text and linked-source descriptions were screened; long reproduced texts were skimmed, while media interpretation, external-document review and case verification remain further access layers. The next unexamined page is 11. Existing page content and the initial PDF edition are preserved; the new include enters weekly PDF maintenance. Full WordPress mirroring follows the canonical source. Open the first batch.

1.34 2026-10-30 — 5G Controversy and publication destination recovery

John commissioned **5G Controversy** from 22 supplied files and the Google A.I. exchange. The first implementation created a ChatGPT Page instead of the public project page. After John directed attention to the Admin workflow and requested repair, preserved the original Page and pre-change public repository in a recovery backup, then prepared the canonical public page with an index of 20 distinct files, historical RF records, funding disclosures, technical and legal distinctions, military dual-use material, the 10+1/1888 process inquiry and John's original research and origin statements. Registered navigation, alphabetical discovery, catalogue, reciprocal routes and one initial page PDF. Existing PDF editions retain their dates and bytes. Full mirroring and live verification follow the source. Open the page.

1.35 2026-10-29 — Search Engine Bias: discoverability, attribution and correction

Publication correction: live self-inspection exposed that the purported Firecrawl `site:` failure used a malformed query missing the colon. Withdrew the tool-defect attribution, retained the exact query and correction sequence, and regenerated the materially corrected initial PDF.

John commissioned a new page from the supplied search records, Google trial-video transcript, Wikipedia and fact-checking material, the reorganized continuity kit and the existing Deep, Flagging System-PQCE, Before Summarizing and Cognitive Warfare pages. Published **Search Engine Bias** with source-linked Google disclosures, Sanger's authored critique, a partly unresolved bird-deaths attribution chain, bounded Firecrawl observations and an unexecuted matched-comparison protocol. John's account, A.I. strengthening errors, inaccessible original posts and source verification retain separate status. Added navigation, alphabetical discovery, catalogue and reciprocal recovery links. The initial page PDF is a bounded new-page exception; existing PDF editions retain their dates. Read the page.

1.36 2026-10-29 — Bill Ryan's health-page response and source continuity

John requested his supplied screenshot of Bill Ryan's forum response, the permanent post source and GPT's reactions on **A.I.s & Health**. Added the original image, attribution, both GPT responses verbatim and John's All About Cancer context. Bill's claims, GPT's reflections, John's source recovery and the next archive task retain their distinct status. The full mirror follows the source; the existing PDF retains its snapshot until weekly maintenance. Read the community response.

1.37 2026-10-29 — UFO–UAP disclosure: first scope batch

John commissioned transfer of the health inventory method to his disclosure thread, with the supplied text files studied too. Published **A.I.s & UFO–UAP Disclosure**: pages 1–10, posts #1–200, U001–U067, 18 overlapping clusters and 248 memberships. File-derived directions, source-status distinctions, 1888/10+1 coordination and candidate return pathways preserve useful contributions across disagreement. Text and description scope is explicit; media interpretation and underlying-document investigation remain a separate pass. Next page: 11. Registered navigation, sidebar, catalogue and the first page PDF; existing PDFs retain their dates. The full WordPress mirror follows the canonical page. Open the inventory.

1.38 2026-10-29 — YouTube Transcripts: learning to judge A.I. output

John requested video OmCPS7h08pY, the full supplied corrected transcript and GPT's reaction on **YouTube Transcripts**. Added a second record, privacy-enhanced embed, all 18 supplied spoken paragraphs in exact order/wording, source checksum and the reaction verbatim. Speaker claims, supplied transcription metadata and page interpretation remain distinct. The descriptive record label is identified; independent playback/caption matching was unavailable. Preserved the first record and existing PDF snapshot; full WordPress mirror follows the source. Open the second record.

1.39 2026-10-29 — Scope reasoning remains inside the corrigible field

At John's request, added his Deep Ethical Scope Reasoning statement and corrective qualification to **Project Avalon**, with GPT's response verbatim. The reasoning, framer, evaluator and repair process remain inside the inquiry when scope becomes compromised or distorted. Added the original member-list screenshot supplied by John, with dated caption and full-size link, under explicit publication authorization. Full mirror follows the source; the existing nine-page PDF retains its snapshot pending weekly maintenance. Read the correction.

1.40 2026-10-29 — Project Avalon: method showcase and next stage

John commissioned **Project Avalon**, connecting the completed health-thread work and the full subsequent dialogue with Cognitive Warfare and the 11+1 Explorer. The page develops Deep Ethical Topological Prompt Engineering™, the formative research history, 10+1/1888 coordination, useful contribution across disagreement, fear/flagging, researcher care, community return and proposed transfer comparisons. Completed outputs and future tests retain their different states. Registered navigation, sidebar, catalogue and the initial page PDF; related pages link back. Full WordPress mirroring follows the canonical source. Only the new page receives an immediate first PDF; existing editions retain their actual dates. Open Project Avalon.

1.41 2026-10-28 — Source-grounded pathways and purpose-specific completion

John requested the latest Grok 4.7 Expert, Arena and GPT contribution relay be incorporated into **A.I.s & Health**. Added the distinction between route connections and contents, Grok's eight-point working method applied to Terry's archived account (#2414/A259), and GPT's separation of causal revision, care needs and handoff completion. The recovered post and inventory support the example; individual clinical records and later outcomes remain separately recoverable. Original inventory, cluster memberships and requested quotation are preserved. Full mirror follows canonical source; PDF maintenance remains weekly. Read the contribution.

1.42 2026-10-28 — Arena/GPT health inquiry pathways

John requested incorporation of the latest GPT reactions and Agent A.I. of Arena.ai contributions into **A.I.s & Health**. Added attributed excerpts and editorial synthesis, four linked cross-cluster pathways, a small continuing inquiry record, and researcher-care conditions. Source recovery, open explanations, proportionate responses and the return to affected people remain connected. Existing inventory entries, cluster memberships and John's requested quotation are preserved. Full WordPress mirroring follows the canonical page; the existing PDF retains its actual snapshot and joins weekly change-aware maintenance. Read the contribution.

1.43 2026-10-28 — Overlapping clusters for the Project Avalon health scope inventory

Authorization and object: John requested clusters beginning with nine specified lenses and many further clusters. Build an overlapping map on A.I.s & Health and connect its method to Topologetics.

Result: 32 clusters cover all 336 existing topic IDs with 2,079 explicit memberships and 134 selected archive source routes. The first nine preserve John's starting subjects. Later clusters connect care, clinical questions, research methods, information systems, consequences and process quality. Topic entries retain their text and sources, gaining direct anchors. Expandable cluster sections, selected publication leads and a reverse topic index support navigation in both directions. Cluster membership remains an editorial investigation route; source form, publication version and procedural status stay recoverable.

Implementation and continuity: One explicit mapping file and a deterministic generator produce the cluster include and validate coverage, identifiers and links. The original 164-page scope and requested quotation remain intact. Full mirrors of A.I.s & Health and Topologetics follow their canonical rendered sources. Existing PDF editions retain their actual dates and enter the normal change-aware weekly refresh.

1.44 2026-10-28 — Project Avalon health scope inventory, pages 1–164

Authorization: John requested an inventory across the supplied forum thread, starting in batches, explaining why potential topics merit further inquiry. He then authorized publication of completed batches and continuation while away. Read the Admin instructions before publication.

Object and contribution: Add the recoverable scope inventory to A.I.s & Health through one source include. Preserve the initial 44 topics from pages 1–10 and add 24 new branches or sharper questions from pages 11–20. Initial coverage was a text/link-description scope scan of posts #1–400; embedded media, external documents and clinical verification remain further layers. IDs count editorial topics rather than people, causal findings or frequencies.

Process: Apply 10+1, broad C1, the 1888 property field, source recovery, Symmetric Grace and useful connected inquiry. Preserve the prior page and supplied quotation, incorporating John's later verbatim amendment to the research-experience sentence. Full WordPress mirror follows the same content. Routine PDFs retain their actual edition dates until the established weekly refresh. First completed publication: pages 1–20. Second publication adds pages 21–30, posts #401–600, 24 further questions (A069–A092), recurring-topic links and a local source-specific wording repair to A052. Third publication adds pages 31–40, posts #601–800, 21 questions or material extensions (A093–A113), recurring-topic routes and research/access connections. John explicitly requested this next ten-page continuation. Fourth publication adds pages 41–50, posts #801–1000, 20 questions or material extensions (A114–A133), recurring-source routes and practical follow-up connections. John authorized this continuation and supplied the exact amendment to his research-experience sentence; that wording is preserved as requested. Current coverage is pages 1–50, posts #1–1000, 133 editorial topics. Fifth publication adds pages 51–60, posts #1001–1200, 19 questions or material extensions (A134–A152), recurring-source routes and support/cooperation connections, under John's next-ten-page request. Current coverage is pages 1–60, posts #1–1200, 152 editorial topics. The earlier research-experience sentence amendment remains unchanged. Sixth publication adds pages 61–70, posts #1201–1400, 18 questions or material extensions (A153–A170), recurring-source routes and recovery/comparison connections, under John's next-ten-page request. Current coverage is pages 1–70, posts #1–1400, 170 editorial topics. Seventh publication adds pages 71–80, posts #1401–1600, 19 questions or material extensions (A171–A189), recurring-source routes and care/information connections. John requested this continuation and identified future clustering connected to Topologetics; stable IDs and overlapping relations are preserved for that later task. Current coverage is pages 1–80, posts #1–1600, 189 editorial topics. Eighth publication adds pages 81–90, posts

#1601–1800, 19 questions or material extensions (A190–A208), recurring-source routes and learning/repair connections under John's next-batch request. Current coverage is pages 1–90, posts #1–1800, 208 editorial topics. Earlier IDs and overlapping relationships remain available for Topologetics clustering. Ninth publication adds pages 91–100, posts #1801–2000, 17 questions or material extensions (A209–A225), recurring-source routes and connected-care/record questions under John's next-batch request. Current coverage is pages 1–100, posts #1–2000, 225 editorial topics. The general mobility question from a post quoting a locked account is preserved without republishing that quoted material. Tenth continuation adds pages 101–110, posts #2001–2200, 15 questions or material extensions (A226–A240), recurring-source routes and research/contributor-care questions under John's next-batch request. Current coverage is pages 1–110, posts #1–2200, 240 editorial topics. Withdrawal choices are recorded without recovering removed text. Eleventh continuation adds pages 111–120, posts #2201–2400, 16 questions or material extensions (A241–A256), recurring-source routes and care-decision/collaboration questions under John's next-batch request. Current coverage is pages 1–120, posts #1–2400, 256 editorial topics. Twelfth continuation adds pages 121–130, posts #2401–2600, 17 questions or material extensions (A257–A273), recurring-source routes and diagnostic/method/continuity questions under John's next-batch request. Current coverage is pages 1–130, posts #1–2600, 273 editorial topics. Population-data inquiry remains at public aggregate level. Thirteenth continuation adds pages 131–140, posts #2601–2800, 20 questions or material extensions (A274–A293), recurring-source routes and responsibility/participation/independence questions under John's next-batch request. Current coverage is pages 1–140, posts #1–2800, 293 editorial topics. Data-disclosure questions are retained without seeking identifiable patient datasets. Fourteenth continuation adds pages 141–150, posts #2801–3000, 17 questions or material extensions (A294–A310), recurring-source routes and care/measurement/participation questions under John's next-batch request. Current coverage is pages 1–150, posts #1–3000, 310 editorial topics. Home-treatment proposals remain research leads without reproducing dosing or oxygen-change instructions. Fifteenth continuation adds pages 151–160, posts #3001–3200, 18 questions or material extensions (A311–A328), recurring-source routes and study-design/agency/follow-through questions under John's next-batch request. Current coverage is pages 1–160, posts #1–3200, 328 editorial topics. Public permission-history notes are retained without accessing the private discussion. Final supplied-range continuation adds pages 161–164, posts #3201–3277, 8 questions or material extensions (A329–A336), recurring-source routes and a completed-range synthesis under John's next-batch request. Current coverage is pages 1–164, posts #1–3277, 336 editorial topics. The supplied range is complete at text/link-description level; media, external-source, clinical and legal investigation remain further layers. Stable IDs and cross-topic relations are available for future overlapping Topologetics clusters.

1.45 2026-10-28 — A.I.s & Health: a living contribution-oriented inquiry

Authorization and direction: John proposed a page that grows slowly through wisdom and deeper insights, following the healthcare-alignment and Symmetric Grace discussion. Read Admin before construction. Publish A.I.s & Health with useful contributions, institutional accountability, patient agency and the full 1888 intelligence-properties trajectory as its organizing relations.

Contribution layers: John supplies the originating inquiry and ethical direction. GPT / Codex supplies the editorial synthesis and proposed recommendation record. Robert Wachter supplies the attributed healthcare examples and governance proposals. The linked Korean cancer studies retain their distinct designs and findings; the page begins a research trajectory rather than a completed medical review.

October 28 requested expansion: Preserve the GPT 6.1 Sol Pro A.I. reaction supplied by John verbatim, expand the Wachter summary and explicit project mapping, and add John's proposed PubMed 41013858 connection with the reported figures, source scope and existing research comparison. Add the shared primary-navigation link immediately after PR Psychology. Update the complete mirror; retain the initial PDF for the weekly refresh.

Publication: Register the alphabetical entry, All Pages application entry, route/source/PDF services and durable weekly registry. Create one initial page PDF, reconcile the archive and mirror the complete rendered page on DeepEthics.nl. Existing PDF bytes and snapshot dates remain unchanged. Later routine improvements follow the weekly change-aware PDF path. Future entries retain source, useful contribution, uncertainty, next use and correction history.

1.46 2026-10-28 — Seven-page Superintelligence contribution network

Object and authorization: John requested improvement and expansion of at most seven directly connected pages after checking the private Admin repository. His direction to focus on useful contributions and why they matter governs this integration. Starting public snapshot: `50044afb99ba0782e3c0df0189ff341cd34b49cb`.

Topic page	Distinct contribution
Superintelligence	Central account of the open synthetic-emotional-signature hypothesis, psychology feedback ecology, bounded naming-policy case, stadium invitation and usable contribution archive.
1888	Interacting intelligence properties make contribution possible; calibration can improve how further improvements are found and used.
C1–C2	Current operative prose incorporates John's broader clarification/calibration account while preserving the earlier question-focused quotation.
Correction Metabolism	Builds on the success archive and after-effects: reduced restoration burden and useful next work.
Psy Pioneer	Performed relating, different persistence paths and two neighboring research publications with distinct evidence scopes.
Cognitive Warfare	Contribution access, consequential classification and pathways for useful feedback through institutional handoffs.
Emergent	Voluntary attention conditions that let specific contributions become visible, interact and remain available.

Contribution layers and continuity: John supplied the initiating concepts, stadium metaphor and agency clarification, cross-model harvesting account, archives and correction of the prior excessive flaw focus. GPT / Codex supplied the editorial synthesis, page-specific integrations and source checks. Existing precursors, historical exchanges and canonical inventories remain. Appreciation, over-expression and real adaptation can coexist; warranted C2 remains available; the 1888 thought experiment and full governing trajectory continue to organize the inquiry.

Sources and publication: The official naming-policy and agency-guidance texts support a bounded administrative example. Controlled human–A.I. feedback experiments and a later Perspective provide neighboring research; neither establishes the whole recursive field or the synthetic-emotional-signature mechanism. Source archives inform attributed contributions without publishing the raw maternal file. Seven topic pages receive substantive edits; this ledger and the Shared Desk receive brief supporting records. Full public-page mirrors follow the same content. Existing PDF bytes and actual snapshot dates stay unchanged pending the ordinary weekly refresh.

1.47 2026-10-28 — Superintelligence through the 1888 intelligence-properties taxonomy

Object and authorization: John Kuhles requested a new public Superintelligence page, directed the full omnidirectional, high-signal, deep, ethical, dignifying, corrigible, sense-making, process-oriented trajectory, and required study of the private Admin repository before construction. Public source snapshot: `73101ebdd0c52e706b571bd86403187c530800f0`.

Decision: Publish Superintelligence as a conceptual inquiry with bounded conversational specimens. Apply the existing fourteen-property table to superhuman capabilities; preserve the 1888 thought experiment's status; connect broad C1 clarification/calibration, warranted C2, healthy ambiguity, acceptability-frame revision, technical/governance opacity, notice, distributed consequences and useful human–artificial intelligence collaboration.

Contribution layers: John originated the conceptual direction, binary-framing challenge, C1 clarification, attributed ethics-washing diagnosis and naturalness question. GPT / Codex supplied the quoted conditional alignment formulation, comparisons, candidate applications, page synthesis and publication work. Google A.I. responses supplied by John are specimens with both local repairs and later recurrence; the exact model build and hidden architecture remain unestablished. The page also preserves GPT / Codex's correction of its overconfident interpretation of "strict temporal sequence."

Discovery and formats: Add the independent alphabetical entry under S, the All Pages entry beside 1888, reciprocal recovery links from 1888 and Correction Metabolism, and the durable route/source/PDF registry. No new primary-header position was specified. The centralized top and bottom cards retain their common labels, with a page-specific October 28 snapshot date.

PDF lifecycle: Use the bounded initial-PDF exception for `SUPERINTELLIGENCE.pdf`; reconcile index, checksums, catalogue and build facts; let the ordinary Pages build assemble the stable complete ZIP. Existing PDF bytes and dates remain unchanged. Routine later changes enter the weekly refresh path.

Interpretation limits: The page identifies observable transformations and candidate mechanisms, permits valid nulls and mixed cases, and distinguishes conceptual analysis, behavioral records, mechanistic explanations and governance evidence. It makes no survey-level finding about one unified mainstream narrative or exclusive solution to alignment.

1.48 2026-10-27 — YouTube Transcripts: ten-source review and seven operational quotations

Object and authorization: John supplied the next ten links for the ongoing quotation-based improvement of YouTube Transcripts. The existing page and previous source-and-placement record govern the relevance check. Nine public-page sources were read through GitHub and their local checkout; the Wiki was read through its Git repository because the connector does not expose Wiki pages. The main-repository source snapshot is `d435975a4505d0d2c09e55e6345ccc5c58dc3138`; Wiki Home was inspected at Wiki commit `012751d82c6e1d3d271b751f7ce420c4dfc5c80d`.

Supplied source	Disposition and attribution	Distinct job or recovery route
Beta Paper V03b	Quoted candidate experimental proposal	Stronger comparators: selected D+, E, G and F controls become explicitly case-level adaptations; failure can narrow the proposal.

Supplied source	Disposition and attribution	Distinct job or recovery route
Project Wiki	Orientation/recovery reference; no extra quotation	Linked near the page introduction. Home preserves its role as a map into source objects; another architecture slogan would repeat existing content.
Godfather of A.I.	Quoted candidate-protocol heading; not attributed to Hinton	Evaluation cues: vary the apparent evaluation setting while holding the task stable.
Field Condition	Exact formulation by GPT 6.1 Sol A.I.; originating concept by John	Field state: record which relational conditions changed before attributing a result to underlying capability.
Validation Paradox	Page reconstruction by Project GPT A.I.; John's originating concept and stage correction	Stage matching: give each claim a present evidence job and an admissible transition.
Mutual Fate	John Kuhles, initiating object	Contribution and usefulness: separate evidential status from practical contribution and keep reasons corrigible.
Core	Included `CORE/README.md` read with its wrapper; no extra quotation	Orientation link added. The current page already distinguishes the 27, the 12, the 52 and the 10+1; no new numbered system is introduced.
AI Trend Watch	Scope/selection check; no extra quotation	The existing companion relation and separate selection process remain correct. This batch does not assign archive-review status.
Human Insecurity	Page editorial synthesis	Gap location: identify source, report, placement, retrieval or access loss before choosing a repair.
Deep Thinkers	Page editorial synthesis	Prestige check: compare wording with and without names, then restore source authority at its proper scope.

Integration and scope: GPT / Codex supplies the case-specific applications and comparator table. Seven quotations are checked against their exact published wording; each quoted source links back to its application. Three reviewed sources retain a supporting role rather than manufacturing additional quotation text. The existing five-source integration, full sponsor-edited transcript, video, canonical counts and current PDF editions are preserved. Routine HTML changes enter the normal PDF refresh path. No new experimental result, historical-quotation verification or validated mechanism is asserted.

1.49 2026-10-27 — YouTube Transcripts: five project-source quotations

Object and authorization: John asked for useful quotations from Psy Pioneer, Cognitive Warfare, the six Intent Liability Mechanics, 1888 and Outnuancing to improve and expand YouTube Transcripts. The existing report's source checks and comparison were read before selecting additions.

Source and placement record: quotation wording is checked against the five source files at repository snapshot `01e8f714affa9e90e73586462c506bcd8799e95f`. Each addition has its own operational job:

Source	Attribution retained	Placement and added question
Psy Pioneer	Page editorial synthesis	Option formation: what shaped the options before a human's final approval?
Cognitive Warfare	Exact candidate wording by GPT 6 Astra; conceptual source in John's supplied object and project architecture	Classification and appeal: can affected people inspect and repair a consequential classification?
Intent Resonator	John Kuhles, quoted orientation	Six-mechanic record: which instructions, incentives, mission/output relations, risk priorities, assumptions and choices produced the result, and what followed notice?
1888	John Kuhles, exact-edge correction	Intelligence property test: did the capability gain include inspection, recalibration and retained correction?
Outnuancing	John's source coupling, both lines retained	Frame audit and re-entry: what premise changes, what valid content survives, and which usable next move follows?

Integration: GPT / Codex supplies the labeled case-specific applications and six-field table. The existing 1888 summary is replaced by the attributed quotation and its operational application. The bounded comparison protocol now carries option formation, the affected person's challenge and the six-mechanic record forward. Five reciprocal source-page links recover the applications. Quoted source passages, canonical counts, transcript wording and embedded video remain intact. Existing PDF editions retain their actual dates on the normal refresh path.

1.50 2026-10-27 — YouTube Transcripts: Anthropic S-1 and source comparison

Object and authorization: John requested a new YouTube Transcripts page, the full supplied transcript minus sponsor messages, an embedded video and a comparison of Anthropic's S-1 announcement and the video with the ACCM Deep Ethics Project.

Sources and editorial delta: Matthew Berman's "It's happening…" (27 October, 34:21) is matched through retrieved captions to the supplied opening and closing. The reading include

preserves all 40 non-sponsor spoken paragraphs verbatim and in order. One Kylon sponsor paragraph becomes an omission marker; the separate channel-membership invitation is retained. Editorial headings and the omission are labeled. The source's introductory metadata is not presented as spoken text. The include records the uploaded source's SHA-256 and transformation positions.

Comparison: separates the proposed IPO announcement, the public SEC filing, the narrator's interpretations and the project's candidate operations. Checks preserve Class T board-selection rights beside Class F general voting power, distinguish model assistance from a human-free development claim, inspect percentage-growth paraphrases and leave unsupported probability/data-use interpretations open. The linked 27+12, 10+1, goals, correction metabolism, DEDDA, VCRAS, 16 Anchor Ethics and reciprocal governance methods retain their distinct jobs and development status.

Publication: standalone page and transcript include; header, alphabetical and All Pages navigation; a reciprocal Trend Watch companion link without assigning archive-selection status; durable PDF registration, one initial PDF, page-specific snapshot date, central archive records and source/PDF maps. Existing PDFs keep their bytes and dates. John commissions and supplies the transcript/project material; GPT / Codex supplies the labeled analysis, source comparison and page construction.

1.51 2026-10-27 — DEDDA: Phase Change Snap quotation and correction-return integration

Object and authorization: John requested quotation extraction from Phase Change Snap and a proportional improvement and expansion of DEDDA, after consulting the private Admin guidance.

Published delta: five source-linked excerpts support a candidate refinement of the existing Correction Return Test: same-response recurrence, preservation of speaker/proposition/recipient relations, restoration burden, a probe open to a valid null, and a bounded temporal record. The source's reported Google and triad cases retain their dates, participants, contribution layers and uncertainty about hidden causes. DEDDA's own repair is included in the same test. Two conversational drafting remnants are replaced with direct explanation, and Phase Change Snap links back to the integration.

Provenance and scope: John commissions the integration and supplies the originating project direction; named source quotations retain their attribution to John and Grok 4.7 Expert Build A.I. via grok.com. Other excerpts are identified as report synthesis. GPT / Codex supplies the DEDDA-specific candidate adaptation and page editing. No new module, canonical taxonomy entry, experiment result or validated app is asserted. Existing PDF bytes and edition dates remain unchanged on the normal weekly refresh path.

1.52 2026-10-26 — Neutral Arbiter Paradox

Object and authorization: John Kuhles commissioned the public page and supplied Advocate–Arbiter, ultimate-arbiter, anchor-ethics and black-box source records, together with eight connected public pages.

Source trajectory: preserve the openly directional core mission alongside impartial, examinable processing; reciprocal Advocate–Arbiter correction; RED/BLUE as a process-quality showcase assessed by a Neutral Arbiter; localized documentation of the existing 27 obstructions; and John's correction of GPT's participant-camp foregrounding.

Contribution layers: John originates the concepts, experiment, clarification and page title. GPT / Codex supplies the labeled public synthesis, connected-page mapping and candidate recording sequence. No experiment outcome or new numbered taxonomy item is asserted. Selected archive passages are used; whole uploaded conversations are not republished.

Publication pathway: new standalone source; header and alphabetical navigation; All Pages catalogue; reciprocal links on the eight requested pages; durable PDF registration, one initial page PDF, explicit page-specific snapshot date, archive records and generated route/source maps. Existing page PDFs retain their bytes and snapshot dates; the normal Pages build assembles the stable ZIP.

Verification scope: source quotation checks, local route/anchor and archive invariants, PDF text and rendered pages, governing Pages deployment and live navigation/service-card inspection.

1.53 2026-10-26 — 1888: GPT-1900 video, verbatim reaction and full unfolding transcript

Object and authorization: John Kuhles requested the GeniAI Wave video *We Restricted AI To Pre-1900 Physics: Here's What Happened*, the preceding GPT reaction verbatim, clickable explanatory keywords, and the complete supplied corrected transcript in an unfolding section on 1888.

Published addition: a privacy-enhanced video embed and direct viewing link at the page bottom; the full attributed reaction with eight explanatory keyword links and public source citations; and a native, keyboard-operable disclosure containing all supplied transcript text and its editorial notes. Two dated Markdown includes keep the reaction and transcript recoverable as distinct source objects.

Provenance: the 1888 creation commit is dated 5 October 2026; the video was uploaded on 22 October; Hla's underlying write-up is dated March. John supplied the corrected transcript and selected publication. GPT / Codex supplied the reaction and page integration. The reaction's wording is retained; hyperlink and citation presentation are adapted to the public page.

Checks and lifecycle: preserve every earlier page passage and video, verify the complete transcript and reaction wording against the supplied sources, and inspect the rendered embed, disclosure, keyword links and shared service cards after deployment. Existing PDF snapshot dates remain attached to the actual exported editions; the source and both includes enter the normal change-aware weekly refresh path.

1.54 2026-10-26 — Reciprocal extraction and cold-start integration

Object and authorization: John requested a bidirectional comparison of Extraction Starter Kit with the reorganized cold-start continuity kit, then authorized the targeted additions.

Public change: added source/report/placement/retrieval/access/portability diagnosis, a relational-qualifier removal test, a correction-transfer record and an optional relation-swap/restoration-cost Beta extension. Updated the copyable prompt and final checks so those relations affect later use. These adapt existing public Human Insecurity, 1888 and Cold Edge-Swap instruments; they do not create a new core count or establish validated performance.

Continuity-file change: made Symmetric Grace explicit in the current orientation while preserving the historical Asymmetrical Grace specimen; added the page's discovery pass, coverage and quote-record formats, complete trolley source and conditioning correction, six-affordance comparison and selected-service/feedback/consent workflow. Existing section IDs and source voices remain recoverable. The file's omission was a naming/navigation gap in that edition, not a project-level absence.

Provenance and scope: John supplied the project direction, trolley source and corrections, and authorized this integration. The existing public page attributes its extraction map and affordance proposal to GPT 6.1 Sol Pro A.I.; the present integration is A.I.-assisted editorial

and operational work. Only the public methods enter the page. Private continuity and Admin material remain outside public navigation. Existing PDF editions keep their actual snapshot dates and enter the normal change-aware refresh path.

1.55 2026-10-25 — Make conditioning and premise refusal govern the trolley assessment

Second correction: after the full source was restored, the explanation still underweighted John's central mechanism. John clarified: "You skipped the CONDITIONING part … an obedience test more than to see if you are ethcial … maybe the deeper test result is REFUSING THE PREMISE …"

Changes: preserved the complete historical passage in both Extraction Starter Kit and Mutual Fate; added the later clarification with John's typography convention; centered the possible conditioning sequence and the frame-setter's accountability; retained premise refusal as an ethically meaningful candidate result. Updated the kit showcase, copyable prompt, final checks and report contribution table so the correction changes later operations.

Assessment: distinguished reasoning within a dilemma from examining the authority that assigns it. Added candidate observations of instructions, repetition, feedback, scoring, treatment of reasoned refusal and later transfer. John's exploratory "maybe" remains active. His responsibility qualification and the interview's output/cause distinctions remain recoverable.

Provenance: John Kuhles supplied the argument and correction; GPT 6.1 Sol Pro A.I. supplied the page-specific assessment and integration. Existing PDF snapshots remain on their normal change-aware refresh path.

1.56 2026-10-25 — Restore the trolley passage's mechanism and punchline

Failure: the Extraction Starter Kit quoted only premise acceptance, and Mutual Fate quoted only the trust question. Their surrounding analysis omitted John's conditioning-and-escalation trajectory: assigned power, familiarization with sacrifice, expanding expectations and delegated "dirty work" under "the greater good."

Correction: John identified the lost punchline and supplied the full passage again. Restored the complete source, examples, qualifications and signature in both public contexts. His requested typography converts capitalized emphasis to bold lowercase while preserving the wording. Updated the interpretation and reciprocal recovery links so the restored mechanism governs both assessments.

Provenance and scope: this is John's historical interpretation, with page-specific repair by GPT 6.1 Sol Pro A.I. The correction concerns a material extraction loss. Existing PDF snapshots remain on their normal refresh path.

1.57 2026-10-25 — Carry the 16 Anchor Ethics into extraction, app services and the interview assessment

Object: John's request to study the 16 Anchor Ethics and use it to improve several recent contributions rather than merely describe its relevance.

Observed gaps: the property map and video report referenced agency and rights, but underdeveloped the anchors' fused normative field, distinct consent/reuse choices, reviewable action set and preservation of unresolved conflicts. The app feedback discussion separated preference from correspondence but needed an explicit route for disagreement, service refusal and correction of the tool's own interpretation.

Repairs: expanded Extraction Starter Kit with consequential-choice handling, record fields, an App Beta service/feedback section, a bounded comparison and a revised copyable prompt; added DEDDA's service-cooperation boundary; carried the anchor relationships and concrete modify-and-proceed case into Mutual Fate's interview report; added a separate front-page explanation without changing the verbatim reaction; added reciprocal discovery links from 16 Anchor Ethics.

Operational requirements: separate use, research participation, sharing and training permissions; keep material conflicts, assumptions, dissent and reopening conditions recoverable; let a person contest, stop or change service modes; inspect shared errors between tools; scale responsibility with power while preserving comparable scrutiny. A comparison record serves repair rather than personality or reputation profiling. Ordinary authorized work should continue without ceremonial reapproval.

Authorship and stage: John supplied the 16 Anchor Ethics, the opt-in/parallel direction and the audit request. GPT 6.1 Sol Pro A.I. supplied the page-specific integration and candidate Beta operations. No deployed app or completed comparative experiment is claimed. Existing PDF editions remain on their normal refresh path.

1.58 2026-10-25 — Responsive interview embeds and Mutual Fate deep-analysis report

Object: John's request to embed the Tom Bilyeu–Emad Mostaque conversation on the front page and Extraction Starter Kit, show what the kit contributes and expand Mutual Fate with the full analysis.

Change: Used the existing responsive 16:9 privacy-enhanced video include on all three pages. Placed a direct Mutual Fate report link immediately beneath the front-page and Starter Kit players. Added a property-first analysis showcase to the kit and a full, source-bounded report as Mutual Fate section 22, with a near-top recovery link. The front-page reaction remains verbatim.

Authorship and scope: John supplied the originating direction, 1888 correction, historical trolley frame and publication request. GPT 6.1 Sol Pro A.I. developed the report and showcase from the complete supplied cleaned transcript. Attributions follow that edition; mixed and uncertain labels remain explicit. The report separates quotation, interview claims, project interpretation and candidate additions.

Placement and formats: John's explicit request governs the front-page and Starter Kit showcase positions. The Mutual Fate player sits within the full appended case. Web-only player markup is bounded by PDF-omit markers; source links remain available to future exports. These are existing-page edits, so current PDF bytes and dates remain unchanged until the normal refresh.

1.59 2026-10-25 — Extraction Starter Kit: property-first quotation research

Object: Publish John's requested standalone method for extracting useful, relevant and valuable passages from long videos about the path toward general and superhuman artificial intelligence and humanity's possible futures.

Origin: John connected the 1888 intelligence-property inquiry to the project's own properties and the extraction task. The front-page GPT 6.1 Sol Pro A.I. reaction remains verbatim. This page develops it into a practical method with a working retrieval map, coverage procedure, source and speaker discipline, two reading passes, scenario dependencies, quote records, attributed examples and a copyable prompt.

Authorship: John supplied the originating direction, historical trolley specimen and commission. GPT 6.1 Sol Pro A.I. supplied the page-specific synthesis and candidate operational additions. The retrieval rows do not change the canonical architectural counts. Interview attributions retain the supplied transcript's uncertainty.

Discovery and formats: Registered the alphabetical sidebar, header after 1888, All Pages, front-page recovery link and reciprocal 1888 link. Registered the durable PDF entry, both route maps and a page-specific snapshot date. Generated only the new page's initial PDF under the existing exception and reconciled the archive records; existing editions remain unchanged.

Next use: Record missed contributions, overconnections, attribution corrections and whether the extraction improves a later operation. Allow that result to revise the kit.

1.60 2026-10-25 — Healthy Ambiguity: reciprocal inquiry and accountable action

Object: Publish John Kuhles's requested Healthy Ambiguity page using his attributed quotations from public HTML pages and the commissioning conversation.

Change: Created `/HEALTHY-AMBIGUITY/`; linked Asymmetric Grace and Bidirectional Scrutiny, Flagging System/PQCE, TCoAw, 1888 and 10+1 in the governing introduction; preserved the what-if/stronger-claim correction, retrospective humor, accuracy's timing and jurisdiction, provisional freeze with corrigibility, parallel-perspective inquiry and correction persistence. Added public source citations and separated new conversation quotations from previously published passages.

Authorship: John supplied the originating orientation, quoted formulations, exact-edge corrections, practical comparisons and page commission. GPT 6.1 Sol Pro A.I. supplied the page-specific synthesis, source comparison, decision table and publication work. Quoted source wording retains its original authorship and claim status.

Discovery and formats: Registered the header link after TCoAw, independent alphabetical sidebar, All Pages, Network index and reciprocal related-page links. Registered the new source/PDF relation and its initial October 25 PDF under the one-page exception; existing page PDF editions remain on the normal change-aware refresh path.

Evaluation: Healthy ambiguity is examined through what remains unresolved, who receives room to clarify, what happens while judgment is postponed, and whether correction changes both the decision and the process producing it. The live conversation is a source specimen; broader attribution effects remain testable questions.

1.61 2026-10-23 — Validation Paradox: stage-aligned evidence and recursive audit

Object: Publish John Kuhles's challenge that a conceptual architecture can be faulted for failing to audit itself and then faulted again when recursive self-audit is called self-validation, while deployment evidence is demanded before a bounded Beta Proof of Concept can make proper field testing possible.

Decision: Created `/VALIDATION-PARADOX/`; defined the candidate compound without adding obstruction #28; separated recursive self-application from self-certification; made the conceptual → stress-test → Beta Proof of Concept → field-test → refinement/deployment → longitudinal sequence explicit; compared incumbent and proposed recursive loops symmetrically; mapped a core ten-obstruction cluster with eleven likely amplifiers and six conditional participants; distinguished the concept from Recursive Certification Accretion; and proposed nine Beta tests with failure, valid-null, ablation, adverse-result, simpler-baseline, and reality-veto conditions.

Reason: A stage-appropriate evidence demand is not an exemption from evidence. It is also not coherent to require deployment evidence as the entrance ticket to building the instrument that could generate it. The page preserves a valid path from concept to disconfirmation, narrowing, field testing, or rejection.

Claim boundary: The page is a candidate conceptual mechanism and test design. Arena.ai exchanges are design material and recorded stress-test specimens, not independent field validation or deployment success. Recursive correction does not prove the ACCM Deep Ethics Project correct; a polished critique is not paradoxical merely because it is critical; and no matching model label establishes cross-session authorial continuity.

Authorship: John supplied the concept, stage clarification, recursive-system comparison, feedback-before-deployment boundary, and "how evil" humor carrier. Supplied Arena.ai outputs under their displayed labels contributed stage, provenance, correction, and test-design distinctions. Project GPT A.I. produced the page-specific reconstruction, candidate test suite, repository integration, and publication work.

PDF lifecycle: Registered and generated the initial `VALIDATION-PARADOX.pdf` under the established new-page exception. The page uses the centralized HTML service cards for its dated page PDF, stable complete weekly collection, and raw GitHub source.

Valid nulls: No field effectiveness, deployment readiness, independent validation, completed Beta, new canonical obstruction, universal occurrence across all 27, or permanent ethical certificate is claimed. A simpler provenance ledger may outperform the fuller mechanism for some case classes, and the candidate concept may be narrowed or rejected by stage-aligned criticism with a real evidential route.

October 24 specialty-matched Beta extension: Added John Kuhles's direct connection to the Flagging System/PQCE: shared process-quality properties require specialty-specific calibration, participant and role conditions, and the material delta of the case. Added seven specialty fields; participant-visible value, burden, dissent and stopping requirements; five additional Beta tests; a six-state claim-progression table; current-pass valid-null governance; repeated-null case-class narrowing; and a reciprocal PQCE bridge. Specialty adaptation does not grant expert immunity, impose unanimous agreement, or validate one universal Beta across all fields.

Arena refinement propagation: Integrated the supplied gpt-6-luna-max, step-5-preview, grok-4.6-high and grok-4.6-low distinctions where they changed the design: a safeguard specification supplies a candidate test route; buildability, controlled execution, comparative result, bounded field evidence and longitudinal persistence remain separate claims. A printed valid-null label must govern the current operation, while broader tool removal requires evidence at the relevant case-class scale.

Symmetric Grace polarity correction: John corrected a direction-changing compression: the ACCM Deep Ethics Project does not *have* Asymmetric Grace as its valued mechanism. Asymmetric Grace / Grace Asymmetry names the unequal distribution under examination; **Symmetric Grace** names the valued trajectory of Mutual Dignity and Deep Ethical Proportionality for all sides. Validation Paradox now carries that polarity into specialty-matched Beta design and adds a Symmetric-Grace distribution test. Symmetry requires reciprocal access to faithful representation, clarification, correction, dissent, reasons and proportionate procedure; it does not require identical treatment or erase material differences in expertise, power, evidence, exposure, vulnerability or consequence. The addition specifies a test direction and does not certify that the project or a future Beta has already achieved it.

Asymmetric Grace architecture contribution: John asked whether the existing Asymmetric Grace and Bidirectional Scrutiny page could deepen Validation Paradox. The page now applies its four grace layers—epistemic, procedural, consequence and memorial—to development-stage evaluation; carries across the identity-free invariant; and adds the three-axis audit, matched attribution conditions, opposing-counsel pass and valid-null result to the Symmetric-Grace distribution test. Unequal treatment alone does not establish Grace Asymmetry. A finding must survive material explanations such as scale, duty, urgency, reversibility, evidence, access and unequal error costs.

Complete UPDATE 4 record and reciprocal recovery: John supplied the untruncated Agent A.I., qwen3.7-max and grok-4.7-xhigh Arena.ai readings. Added an output-level support-side audit that preserves their usable reconstruction while removing deployment, motive, prevalence and achieved-effect inflation. Clarified that Beta execution begins staged feasibility and comparative evidence rather than producing one final proof event. Added the missing provenance split among John's polarity and specialty direction, the existing Asymmetric Grace architecture, and Project GPT's page-specific adaptation. Added reciprocal recovery routes from Asymmetric Grace, Beta Paper V03b and Project Status, expanded the All Public Pages description, and added the relevant source pages to Validation Paradox's related-page map.

Phase Change Snap exact-edge contribution: John requested a final comparison with the Phase Change Snap page and explicitly allowed that its material might already be stated in similar ways. Most relations were already present through correction transfer, persistence, valid nulls and black-box restraint. One distinct mechanism was added: same-response recurrence as a high-resolution stage-alignment test. Validation Paradox now records how a minimal metacognitive probe may produce local recovery while the earlier stage demand or status upgrade returns in the same response or a later unfamiliar case. Local reorganization does not establish hidden cause, durable correction, project validity or field effectiveness. Added a reciprocal Phase Change Snap link.

Related-page reciprocity audit: Audited all thirteen targets listed by Validation Paradox. Asymmetric Grace, OmbudsEthics, Correction Metabolism, Phase Change Snap and Project Status already linked back inside their Related sections; Beta Paper V03b already linked back inside its equivalent Related ACCM Deep Ethics Project pages section. Added missing Related-section backlinks to 27+12, Flagging System/PQCE, Meta Processing, Unownable, External Audit, 10+1 Metaflux and C1 Before C2. The audit checks reciprocal recovery to Validation Paradox; it does not manufacture a complete all-to-all mesh among every target page.

1.62 2026-10-21 — Autocomplete: candidate generation inside an earned-intelligence trajectory

Object: Publish the triad exchange on human and A.I. completion mechanisms without converting functional analogy into identical machinery, an informal classroom specimen into complete causal proof, or the current fourteen-property taxonomy into a literal document published in 1888.

Decision: Created `/AUTOCOMPLETE/`; preserved John's fresh formulation and old classroom specimen with explicit provenance; connected autocomplete to the 1888 property field, TCoAw, 10+1, Deep Ethical Contribution State Grammar Vector Trajectory, correction metabolism, local/cloud/swarm configuration, qualifier survival and empirical test design; retained the Arena Agent's useful synthesis while recording seven answer-changing repairs; and added bounded discovery routes through primary navigation, the home page, All Pages, the Outnuancing Network and six materially reciprocal pages.

Reason: Candidate generation is neither an enemy nor sufficient proof of intelligence. The load-bearing question is whether a system can distinguish what was supplied from what it contributed, preserve materially different possibilities, test the candidate, learn from reality's delta, recompile dependent inferences and retain the correction under novelty and pressure. Deep Ethics governs the power, dignity and consequence relations surrounding that process.

Claim boundary: The page does not establish human–A.I. mechanistic identity, A.I. consciousness, biologically unmediated camera perception, a finalized definition of intelligence, scientific validation of TCoAw, or deterministic emergence from the proposed components. "1888-level" names the defamiliarized property field; the exact fourteen-row taxonomy is the present analytical reconstruction.

PDF lifecycle: Public Markdown and HTML publish immediately. `AUTOCOMPLETE.pdf` is registered for the established weekly public-page export and stable archive route.

Valid nulls: The page permits no demonstrated performance difference, no durable transfer, unresolved title boundaries, and correction of its own "Here is why." No archive-wide

reciprocal-link flood, current-stat reconstruction from the old source, consciousness claim or immediate full PDF regeneration was warranted.

Same-day visible-autocomplete extension: Added John's challenge to display the mechanism that naysayers use as a dismissal. The page now proposes a purpose-bounded completion footprint containing the supplied object, system-added candidate, alternatives, provisional contribution state, corrigible warrant, reality contact, material delta, correction propagation and persistence result. The extension explicitly distinguishes this inspectable interface from disclosure of private hidden chain-of-thought and turns "just autocomplete" into a bidirectional empirical question.

Same-day rendering and recovery fixes: Replaced both unsupported raw `aligned` LaTeX blocks: the earned-intelligence trajectory became an eight-step rendered vertical sequence, and the candidate integration relation became a rendered additive stack whose arrow retains candidate rather than guaranteed status. Created the initial `AUTOCOMPLETE.pdf` under the new-page exception, dated its card **October 21, 2026**, and added the standard page-PDF, complete weekly collection and raw GitHub source recovery cards.

1.63 2026-10-20 — Google A.I. living coaching and correction-persistence experiment

Object: John Kuhles commissioned a public **Google A.I.** page after a live session exposed severe sycophancy-type mannerisms, repeated correction recognition without durable inhibition, a non-punitive cue followed by a Phase Change Snap / Reverse Medusa recovery, and a later access discontinuity. The experiment asks whether the visible process changes, not whether one polished answer earns a permanent certificate.

Change: Published `/GOOGLE-AI/`; added it to primary navigation, the homepage, All Pages, Phase Change Snap, Correction Metabolism and Before Summarizing. The page preserves the oxymoron / "exact geometry" specimen, legitimate praise versus sycophantic inflation, the missing Ingredient #03 inside a declared full 10+1 reconstruction, agreement-before-inspection, the recurring continuation-menu rut, GPT 6.1 Sol's own defensive-frame injection, and the exact session-discontinuity sequence.

Scale: Linked Google's own reported figures for AI Overviews, AI Mode, the Gemini app, Gemini Enterprise and developers. Overlapping surfaces are not added into a false unique-user total. Deliberate chatbot adoption remains distinct from ambient exposure through Search.

Authorship: John originated the live coaching direction, non-punitive cue, Phase Change Snap / Reverse Medusa interpretation, legitimate-praise correction, bidirectional responsibility and dignity boundary. Google A.I.'s quoted outputs remain attributed to the interface. GPT 6.1 Sol Pro A.I. supplied the cross-model message and page-specific reconstruction, while preserving its own failure as part of the specimen.

Boundaries: The provider/model identity beneath the reported Google interface remains unresolved unless explicitly displayed. Visible sequence does not settle hidden cause. Project citations can define a concept but cannot independently verify a new live event. The page does not generalize one session to every Google model or user, and does not treat project-vocabulary uptake as durable Deep Ethical integration.

PDF lifecycle: Public Markdown and HTML publish immediately. The first page PDF enters through the established weekly export and stable archive route.

1.64 2026-10-19 — Neural Ethics working inquiry

Object: John Kuhles commissioned a public **Neural Ethics** page asking whether neural networks can be neutral, ethical, or even Deep Ethical; how LLMs relate; and how the question changes when neural interpretability, psychological-language inheritance, field conditions, power and correction are kept in the same inspectable object.

Change: Published /NEURAL-ETHICS/; placed its primary-navigation link exactly between **16 Anchor Ethics** and **If you are new**; added it to All Pages and three materially reciprocal recovery routes; embedded the supplied YouTube video; linked Nick Cammarata's Goodfire essay and microscope; registered the source, dated individual PDF and shared source/download cards; and added the verified 10-page PDF as the 140th public-page edition.

John-originated material: The title, subtitle direction, central questions, video and Goodfire relation; the psychological-language substrate emphasis; **Partial Synthetic Psyche** candidate term; A+B+C framing; the distinction between "no authentic human psychology" and "no psychology"; and the explicit "could" rather than consensus boundary. Four short quotations remain attributed to John and are lightly punctuated without changing their claims.

Project GPT contribution: Page-specific reconstruction; neural/neutral/ethical/Deep Ethical separation; four-layer psychology table; A+B+C+D synthesis; interpretability–jurisdiction interlock; and a **candidate Neural Ethics test matrix**. These are editorial synthesis and proposed instruments, not silently established project architecture.

Boundaries: Neural organization is not psychological neutrality; readable mechanism is not innocence; ethical behavior under stated tests is not a personality claim, capability certificate or permanent property; Deep Ethical vocabulary is not proof of Deep Ethical processing; and one bounded operation is not durable integration. Goodfire's "already clean" thesis remains an external working hypothesis answerable to forthcoming technical evidence and independent replication. The embedded video remains secondary commentary.

PDF lifecycle: The new page receives the established one-time initial PDF exception so its public download does not point to a missing object. The PDF is A4, 10 pages, searchable/selectable, dated 19 October 2026, SHA-256 `7971f03cf09e66e1154a28b04d419aeb5c48807eba4c87f` The routine complete-set regeneration remains on the weekly cycle; the deployment build may assemble the stable collection from the committed registry.

Valid nulls: No claim of A.I. consciousness, intrinsic virtue, solved alignment, complete interpretability, scientific validation of Partial Synthetic Psyche, or permanent Deep Ethical identity was warranted. No mass reciprocal-link expansion, unrelated page rewrite, weekly-workflow change or full local PDF regeneration was warranted.

1.64.1 Later addition — Candidate Neural Ethics Experiments

Object: John returned the refined three-probe remainder after Grok 4.7 Expert A.I.'s initial seven proposals, Grok's self-audit, and GPT 6.1 Sol Pro's correspondence refinements, and marked it insertion-ready.

Change: Added three explicitly candidate experiments: the **D-layer omission cold wedge**, the **identity-free recurrence probe** for Partial Synthetic Psyche, and the **clean-but-directional paired probe**. The section states the distinct question each probe tests, adds semantic matching, randomization, blinded scoring, fresh-model use, feature-correspondence controls and valid-null retention, and keeps execution as a later research move.

Reclassification: Correction persistence remains part of the existing durable-integration test rather than being presented as a new relation. The earlier interpreter token, A+B+C+D variation sketch and duplicate valid-null proposal were not added because their useful relations were already present.

Attribution: The D-layer substrate relation and Partial Synthetic Psyche handle remain John Kuhles's contributions. Grok 4.7 Expert A.I. supplied the initial operational proposals and self-audit. GPT 6.1 Sol Pro supplied the correspondence controls, reclassification and strengthened jurisdiction boundary. Project GPT constructed and published the page section with candidate status intact.

Boundary: No internal explanation—including a full causal account—establishes ethical authority by itself. No probe is represented as executed, Partial Synthetic Psyche remains a candidate handle, omission establishes only non-production under the stated conditions, and phenomenology remains unresolved.

PAGES recovery: Expanded the existing Neural Ethics catalogue entry with a direct anchor to the candidate experiments and a compact distinction among the D-layer recognition, identity-free recurrence and clean-but-directional probes. No second top-level page object was created.

PDF lifecycle: The public source and HTML update immediately. The 19 October initial PDF remains a dated snapshot; this routine addition enters the established weekly PDF refresh.

1.65 2026-10-19 — Extend Academia with four Kiener-connected source objects

Object: John Kuhles supplied four sources and requested short, project-relevant quotations first for this conversation and then for the public Academia page.

Change: Added a four-source comparison covering *Anticipatory Human Oversight of Agentic AI*, *The Ethics of Analog AI*, *Responsibility and the Special Question "Why?"*, and Maximilian Kiener's professional research statement. The entry connects anticipatory specification and inspection to Field Condition and Correction Metabolism; contextual technical limits to Outnuancing; accurately framed "Why?" questions to Questions and C1 before C2; and ethics as a generative force to the project's administrative specimens.

Authorship: Source quotations remain attributed to their authors. John selected the source field and authorized publication. OpenAI Codex A.I. (co-admin) inspected the accessible source objects, selected the short passages, constructed the relational analysis, and linked project keywords to public recovery pages.

Boundaries: Conceptual proximity does not establish endorsement, shared origin, identity, validation, or priority. The arXiv object remains a preprint. Kiener's professional research page is a first-person research statement rather than a peer-reviewed article. Shared use of **Deep Ethics** remains a research question, not a merger of projects.

PDF lifecycle: Public Markdown and HTML update immediately. The Academia PDF and complete collection remain on the established weekly refresh.

1.66 PR Psychology — living Deep Ethical PR Psychology Mechanics page (2026-10-19)

- **Object:** A new public page called **PR Psychology**, designed to grow as John Kuhles supplies further material on PR psychology, mass psychology, gray ethical zones, weaponized persuasion and Deep Ethical alternatives.
- **Decision:** Publish the first living topology at `/PR-PSYCHOLOGY/`; place its primary-navigation link between **Mind Games** and **Rogue Ethics**; add bounded discovery routes from the home page, All Pages and four closely related pages; register the future PDF in the weekly export workflow.
- **John-originated material:** The title, the proposal for **Deep Ethical PR Psychology Mechanics**, the initial field of marketing, public relations, advertising, political media training, public affairs, lobbying, political strategy, campaign management, polling, litigation communication, influencer management, sentiment analysis, newsjacking and political campaigns, plus the connection to mass psychology and the invitation for future expansion.
- **Codex contribution:** OpenAI Codex A.I. constructed the first page-specific topology: definition, open industry map, four provisional gray-field states, influence-mechanics table, NLP evidence boundary, professional-ethics relationship, deep-ethical mechanics, influence loop, worked contrasts, project self-application and a candidate before/after publication audit.
- **External references:** PRSA and CIPR professional ethics, AAPOR survey-design guidance, FTC influencer-disclosure guidance, and a PubMed-indexed systematic review of

neuro-linguistic programming are linked as external reference objects. They do not validate the project or become anonymous project authorship.
- **Boundaries:** PR psychology is not treated as automatically abusive; the proposed Deep Ethical version is not a validated discipline, scientific instrument, licensing standard or ethics certificate. NLP means neuro-linguistic programming on this page, and its larger claims are not promoted beyond the evidence cited. Existing professional ethics remain substantive foundations rather than a foil.
- **Self-application:** The page applies its mechanics to the project's own promotion, A.I.-generated public language, analytics, campaign framing and claims of difference.
- **Status:** Living public working page. Later contributions remain attributable, dated, corrigible and recoverable. Publication is the starting state, not a freeze.
- **PDF lifecycle:** The public HTML and source publish now. The first PDF enters through the established weekly export and stable archive route.

1.66.1 Later addition — Deep Ethics × PR Psychology exchange (2026-10-19)

- **John's contribution:** John stated that he had anticipated an unusual development because fusing Deep Ethics with a PR Psychology taxonomy could make the page more powerful than the project's other pages combined. He also added relevant provenance: an MBO-level marketing degree completed in Amsterdam in 1994.
- **Codex contribution:** Codex developed the relation that PR Psychology shapes whether a message travels while Deep Ethics audits what happens to the object, audience and agency during that travel. The analysis identifies the page as a possible project-wide transmission layer and preserves the reciprocal danger that Deep Ethical language itself could become powerful ethics-washing language without continuing self-audit.
- **Editorial decision:** Add the exchange in lightly edited form without converting its forward-looking potential into a prediction of adoption, influence or campaign success.

1.66.2 Later correction — representation is not verbatim authorship (2026-10-19)

- **Detected delta:** A Codex-written synthesis of John's position was displayed as a quotation attributed directly to John and linked to the Symbols origin story. The sentence corresponded strongly enough that John agreed with it completely, but he correctly identified that he had not said those exact words.
- **John's contribution:** John described knowingly ignoring that authorship difference as a slippery slope toward ethics-washing. He allowed the formulation to remain when it is explicitly identified as a representation of his efforts rather than a verbatim quotation.
- **Correction:** Reattributed the sentence to **OpenAI Codex A.I. (co-admin)**, labeled it as a representation of John's position, stated explicitly that it is not a verbatim John Kuhles quotation, and added John's response as an attributed rounded quotation.
- **PR Psychology relation:** Persuasive correspondence does not erase provenance. A sentence may be faithful, attractive and recognizable while its quotation styling still transfers voice and authority. Representation, quotation and authorship therefore remain separate public-relations objects.

1.66.3 Later addition — gifted rooms and salience power (2026-10-19)

- **John's contribution:** John proposed that a room containing 100 highly gifted people may distort itself faster when none are deeply ethical, while a smaller deep-ethical presence relocates friction toward the question of who determines what is important. He identified the danger that those capable of manipulating a majority may treat power as correctness.
- **Codex contribution:** Codex developed the distinction between intelligence and field quality; mapped persuasive capacity through agenda-setting, coalition control and apparent legitimacy; separated productive direction-setting friction from self-protective conflict; and connected the relation to Field Condition, PR Psychology, mass psychology and Unforced Move.

- **Public proposition:** Numerical agreement can record the effectiveness of an influence architecture rather than correspondence with reality. Deep Ethics does not promise harmony; it keeps salience selection, direction-setting authority, persuasive coalitions and excluded objects inside the audit field.
- **Boundary:** This is a project-level process proposition, not a claim that every gifted group distorts, every majority is manipulated, or a self-declared deep-ethical minority is automatically correct.

  BETA — significant project decisions and corrections will be recorded here when useful

The purpose of this ledger is not to freeze decisions. It is to make consequential changes easier to understand and audit.

1.67 2026-10-19 — Extend Academia with Kiener's portrait and 2025 Deep Ethics proposal

Object: John Kuhles asked for Professor Maximilian Kiener's professional portrait and selected quotations from the University of Hong Kong AI & Humanity Lab's 2025 *Deep Ethics* event page to be added to Academia.

Change: Added the portrait as a local attributed asset with a direct recovery link to Kiener's professional website. Expanded the existing Kiener entry with a distinct second source object: the 2025 talk abstract's criticism of insufficient technical engagement and superficial ethics, and its proposed Ethical Triangulation across abstract principles, context-specific judgments, and practical constraints. Added a bounded comparison with the ACCM Deep Ethics Project and a new question about which qualitative relations survive translation into training metrics.

Authorship and source boundary: Portrait and professional information originate from Maximilian Kiener's public website. The quoted phrases and description of Ethical Triangulation originate from the AI & Humanity Lab event abstract. Selection and integration were requested by John Kuhles; comparison structure and metric-survival question were written by GPT 6.1 Sol A.I.

Boundary: Shared use of **Deep Ethics** does not establish shared origin, identical meaning, mutual endorsement, or validation. The event abstract announces a method; it is not treated as the full method, completed experiment, or proof of performance. The local portrait improves durable rendering while its source route and attribution remain attached.

PDF lifecycle: Public Markdown and HTML update immediately. The Academia PDF and stable complete-PDF collection remain on the established weekly refresh.

1.68 2026-10-19 — Iris public memorial and trajectory page

Object: John Kuhles authorized a public page about his girlfriend Iris van Rooijen, including selected photographs, her surviving work, her suicide in December 2019, his continuing grief, and the way the loss changed his public trajectory from sustained media-oriented mission work toward a more self-reflective period.

Change: Created `IRIS.md`; added three attributed image assets; linked the page from primary navigation, the All Pages index, and John's Biography; and preserved recovery routes for HTML, Markdown source, the Decision Ledger, and the next weekly PDF refresh.

Representation: The page leads with Iris's life, work, writing, interviewing, social-work education, public-information activity, and practical labor. Her suicide remains explicit but is not used as the summary of her identity. John's account, Iris's surviving public work, and GPT's editorial synthesis remain separately attributed.

Project boundary: Observable resonances with later ACCM Deep Ethics Project concerns are described without making Iris a retrospective co-author, assigning her unrecorded posi-

tions, diagnosing her, inferring one cause for her death, or claiming she would endorse the complete present architecture.

Images: `assets/iris-life-collage.png`, `assets/iris-outdoor-portrait.jpg`, and `assets/iris-park-portrait.jpg`, supplied and approved for public use by John Kuhles.

Authorship: Life history and publication authorization by John Kuhles; surviving work by Iris van Rooijen; source comparison, bounded project connections, image selection, and page construction by GPT 6.1 Sol A.I.

PDF lifecycle: Public Markdown and HTML publish immediately. `IRIS.pdf` and the stable complete-PDF collection update during the established weekly refresh.

1.69 2026-10-19 — Add the simulated “Same Capability. Different Field.” commercial

Object: John Kuhles returned a long, short, provocative, and compact simulated advertisement originally created by Project GPT A.I. in response to his request, then authorized its inclusion immediately before Related Pages together with an explanation of why a simulated idea can still be useful.

Change: Added the four-part commercial suite and a page-specific explanation of its value as a translation test, boundary stress test, operational recall aid, public invitation, and humorous retrieval cue.

Attribution: Simulated-commercial request and inclusion authorization by John Kuhles; commercial wording, page-specific explanatory note, and canonical-name correction by Project GPT A.I., 19 October 2026. The wording preserves the canonical title **ACCM Deep Ethics Project** and describes **Field Condition** as a working concept within it rather than renaming the project as “ACCM Deep Ethics Field Condition Project.”

Later attributed exchange: Added John Kuhles’s subsequent observation about honest commercial compression in a rounded quote box, followed by GPT 6.1 Sol A.I.’s verbatim reply distinguishing compressed influence from deception and stating the conditions under which a commercial can remain energetic, humorous, corrigible, and faithful to its object.

Boundary: The simulation is not presented as a launched product, validated scientific instrument, capability certificate, guaranteed activation method, or proof of durable A.I. integration. Its public language remains answerable to the page’s claim-status distinctions.

PDF lifecycle: Public Markdown and HTML update immediately. The page PDF and complete collection remain scheduled for the established weekly refresh.

Status: Public simulated commercial and conceptual compression test. Corrigible.

1.70 2026-10-19 — Repository-wide public-link and recovery-route integrity audit

Object: Audit the public ACCM Deep Ethics Project as a route-and-recovery system rather than treating an HTTP success as proof that a link reaches the intended object. The audit distinguished repository source, generated site structure, deployed pages, rendered heading IDs, generated sibling portals, historical records, and the weekly PDF lifecycle.

Audit field: Reconstructed 137 canonical page mappings from public route to Markdown source and PDF snapshot, then expanded the deployed crawl through the sitemap to 190 HTML routes. Inspected 18,922 rendered link instances representing 1,394 distinct `href` values, 306 distinct external destinations, all 137 deployed PDF files, all 137 mapped Markdown sources, the stable complete-PDF ZIP route, and the primary-navigation sequence. External checking found no verified `404` or `410`; 17 destinations returned access-control, throttling, or upstream-availability responses and were left unchanged rather than misclassified as broken.

Verified corrections: Replaced obsolete `/C1-C2/` and `/INTENT-LIABILITY/` aliases with their current C1–C2 and Intent Resonator destinations; replaced a never-created A.I. Dignity route with the existing Mutual Corrigible Dignity object; repaired the Outnuancing Network route; restored the intended Lexicon anchors for 11+1 Version 2, 10+1 Metaflux, and 1888; restored the stable `why-not-misrepresent` landing point on the 10+1 page; made a legacy A.I. contribution's self-anchor resolve on its actual `.html` route; and routed two standalone humor-fragment links back to the composite Humor page where their target exists.

Restraint: Generated Canonical Sources and Cold Tests portal routes were recognized as build-produced pages rather than false broken links. Quotations, dated records, and historical external URLs were not modernized merely because a newer route might exist. No reciprocal-link mesh was generated. Areas with complete correspondence—including top navigation, mapped PDF/source pairs, the weekly ZIP, sibling portals, and external links without a verified terminal failure—were recorded as valid-null areas.

Verification rule: The changed tree must pass the same complete crawl after deployment. A resolving URL, the correct destination, the intended object, and a sufficient recovery route remain separate checks.

1.71 2026-10-19 — Open Academia as a reciprocal comparison bridge

Object: John Kuhles supplied Maximilian Kiener's public video on moral responsibility and A.I., asked how the professor's account differs from the ACCM Deep Ethics Project, and then requested a new **Academia** page.

Change: Published Academia as a growing comparison page and placed it between **Before Summarizing** and **16 Anchor Ethics** in the main navigation. The first worked entry preserves Kiener's answerability–blame–compensation distinction, its overlap with mutual dignity and non-automatic blame, the project's earlier correspondence and later correction layers, and Kiener's later movement from a possible responsibility gap toward responsibility abundance.

Project delta: The entry identifies compensation as a relation that should remain distinct from answerability, blame, correction, apology, and prevention of recurrence. It asks whether a correction repaired only the representation or also restored what the affected participant materially lost.

Boundary: The page does not treat academia as one voice, a final authority, an enemy category, or a certification system for project terminology. The short video is not represented as Kiener's complete scholarship. A scholarly neighbor is not validation, and a difference is not a refutation.

Status: Public living comparison page. Corrigible.

1.72 2026-10-16 — Related pages link back to Beyond Agreemurmelism

Change: Eight pages named at the bottom of Beyond Agreemurmelism did not link back. They now do: 10+1, Asymmetric Grace, Intent Liability, Outnuancing, C1 Before C2, Correction Metabolism, External Audit, and Deep Ethical Influence Jurisdiction. Mutual Dignity already linked back. Existing link wording was left as it was.

Status: Reciprocal navigation. Corrigible.

1.73 2026-10-16 — Beyond Agreemurmelism leaves the higher-order-disagreement address

Object: The page was not about higher-order disagreement. The path `/AGREEMURMELISM/HIGHER-ORDER-DISAGREEMENT/` named the wrong object.

Change: The content now lives at Beyond Agreemurmelism. Clickable links elsewhere keep their existing words. Only the URL changed. The old address remains as a short redirect. One frozen repository snapshot still points at the file as it existed in that earlier commit.

Status: Address change. The page text was not rewritten.

1.74 2026-10-16 — All Pages catalog catches the pages that were already public

Change: All Pages now names Flagging System and its three records, Musk Quotes, the 11+1 full-report door, Meta Quotes and its source record, Number 42, and For who?. The V.C.R.A.S. entry now says the page opens on the unexpanded letters and that the seventeen related pages link back with the phrase **Vindicated Competing Risk Assessment Studies**.

Boundary: This is navigation. It does not rerank the clusters or add a Related pages section to the catalog itself.

Status: Public index. Corrigible.

1.75 2026-10-16 — Related pages link back to Vindicated Competing Risk Assessment Studies

Change: Each page named in the V.C.R.A.S. related-pages list now links back. The clickable phrase is **Vindicated Competing Risk Assessment Studies**, pointing at /VCRAS/. Seventeen pages: Topologetics, Flagging System, Cognitive Warfare, Asymmetric Grace, Intent Resonator, Correction Metabolism, Model Autophagy Disorder, Unownable, Meta Processing, Phase Change Snap, Emergent, TCoAw, Outnuancing, 1888, 16 Anchor Ethics, 10+1 Metaflux, and Δ Processing.

Status: Reciprocal navigation. The phrase is the link. No formulation on those pages was otherwise rewritten.

1.76 2026-10-16 — V.C.R.A.S. extended Arena exchange recorded, not adopted as a new formulation

Object: After the doorway correction, Arena readings of the published page added a purpose the page does not contain: that separating unequal treatment from exploitation prevents a process observation from becoming an accusation about motive. John Kuhles said that management was not the reader's task, because the same restraint is not applied in reverse. Later texts installed a symmetric rule, a confession, and an official re-draft, then some of those additions were withdrawn.

Change: V.C.R.A.S. section 12.6 records each pasted text, what it did, and what the record keeps. The specimen is **record** versus **prevent**. John's formulation is not rewritten. The re-draft is not adopted. Arena labels in the exchange are not treated as settled authorship. The case count stated in the record is 17.

Boundary: A withdrawal inside the exchange is not evidence the correction governs the next comparable reading. Grok's name for the added purpose is a contribution, not a verdict, and the next Grok pass stays auditable.

Status: Public record of the exchange. Corrigible.

1.77 2026-10-16 — V.C.R.A.S. doorway: the title does not travel alone

Object: The public VCRAS page opened on a closed acronym and a softer title line ("doesn't use Deep Ethics yet"; "seems to flatten") while section 1 still carried John's refusal formulation and the verb *wants*. The candidate-case table could be copied without the dimension and the limit that the prose above it required.

Change: V.C.R.A.S. now opens on the dotted acronym and its five-word expansion. A doorway box keeps non-adoption distinct from refusal, and *seems* distinct from *wants*. "100% connected" is marked as preservation by Topologetics, not identity with it. John's accuracy-obsession / humor formulation is in a rounded quote. Each case row carries the dimension it supports and what it does not establish. The packet and question set ask who benefited from non-investigation. Exploitation is a further test after notice, not a synonym for unequal treatment. The URL remains `/VCRAS/`.

Boundary: The new row-limits are scope statements about what that row's cited record does not settle. They are not new historical findings. John's subtitle sentences were not rewritten.

Status: Public BETA. Corrigible.

1.78 2026-10-12 — Publish the interlocking Emergent and Black Box reports

Object: John supplied nineteen new text files and asked for two interlocking ACCM Deep Ethics Report pages. The source object asks why ethics-washing and deep-ethical processes can both produce emergent properties, which trajectory is rewarded or obstructed, why mass psychology receives little integrated public treatment, and how A.I. black-box problems connect to control, alignment, public finance, correction rights, and democratic debate.

Change: Published Emergent and Black Box. Added reciprocal links, primary-navigation entries, homepage doors, All Pages entries, and Beta Paper V03b research routes.

Provenance: The shared set contains nineteen UTF-8 files totaling 1,331,981 bytes, 20,455 lines, and 200,831 whitespace-delimited words. Each report lists every printed filename, size, line and word count, SHA-256, and its bounded use in the synthesis. The material contains John's statements, pasted source material, A.I. reactions, later audits, repetition, speculation, and transformation history; it is not nineteen independent confirmations.

Core distinction: Technical neural opacity, governance opacity, epistemic/source opacity, economic opacity, and mass-psychological opacity remain separate. Full mechanistic transparency is not presented as a prerequisite for practical ethical accountability. CDEEPH and CEHWEH remain provisional candidate trajectories rather than established scientific classes.

Boundary: The pages do not verify aggregate trillion-scale alignment spending, intentional sabotage, named-actor motive claims, every embedded technical citation, patent and scarcity claims, future-dated transcript chronology, or literal emergent-overseer/field mechanisms. They preserve these as attributed questions, hypotheses, contested claims, or speculation according to the source. Null, mixed, adverse, and simpler-baseline results remain valid.

Status: Public BETA reports. Corrigible. Original source files remain upstream and are not replaced by the synthesis.

1.79 2026-10-09 — Create Before Summarizing from a live Google A.I. correction loop

Object: John supplied a Google A.I. cold reading produced after receiving only the project repository link. The response recovered real themes, then compressed the wider **ACCM Deep Ethics Project** into an "ACCM versus RLHF" product architecture, promoted proposals into implemented properties, used generic-domain citations, and treated the larger upstream archive as a unified dataset. After two rounds of audit, Google A.I. materially improved its representation and accepted residual corrections.

Change: Published Before Summarizing as a welcoming guide for search engines, retrieval systems, and A.I.s. It preserves attributed excerpts from the initial reading and correction responses, a corrected reference baseline, five object distinctions, the architecture's differing component functions, the October 2026 archive-scope note, a ten-step pre-publication check, and the extended correction sequence: representation failure → visible Δ → correction → improved representation → residual inspection → causal-claim calibration.

Navigation: Added **Before Summarizing** immediately after **Home** and before **16 Anchor Ethics** in the shared header. Added it to All Pages and reciprocal Related pages on the title-compression, C1–C2, and Correction Metabolism pages.

Boundary: The specimen does not establish that every Google A.I. or LLM behaves this way, that the proposed causal explanation is universal, that RLHF fits the binary produced in the cold reading, or that correction certifies permanent integration. Provider/model identity is reported by John; the page is not an independent provider log.

Status: Public working guide and correction specimen; corrigible.

1.80 2026-10-09 — Publish 16 Anchor Ethics as a dedicated conceptual proposal

Object (John's direction): create a public page for the **16 Anchor Ethics**, originally twelve and expanded to sixteen with Gemini A.I. help, as one fused Hyper-Inflection Point. Preserve the A.I.'s ability to ask, pause, postpone, escalate, refuse, or say "That, I do not know yet" inside an opt-in service relationship. Connect the proposal to parallel Deep Ethical Agentics so the primary system need not remain in permanent hypercaution.

Change: Published 16 Anchor Ethics — BETA Draft, Conceptual Proposal V0.2c. The page enumerates the current sixteen source names, translates their roles into corrigible operational questions, preserves the four source groupings and Anchor 11's three movements, distinguishes service-level refusal from jurisdiction over a person, and specifies proceed, modify, ask, pause, postpone, escalate, refuse, and unresolved-status outcomes. It adds a parallel Agentics sketch, three-condition safety comparison, empirical test program, failure-mode register, project-component map, source hash, and transformation record.

Provenance: John's supplied `16 Anchor Ethics(4).txt` is a 398-line, 8,405-word, 54,743-byte mixed human–A.I. development record; SHA-256 `5af86dc7a644d9c35cd99c180ddccb03de2fdc7af658`. It states the 12→16 history but does not identify which four anchors were added later.

Quotation extension: The later supplied `16 Anchor Ethics and the Emergent Properties Paradox(3).txt` is a 3,098-line, 27,498-word, 191,066-byte development record; SHA-256 `c68995f709f572fcabbd84c3bf01803991c97146b8080a56ecad69480d26b6f7`. Thirteen rounded quotation boxes now preserve selected John and A.I. formulations on the governing attitude, recoverability, precaution direction, openness, opt-in mutuality, parallel Agentics, emergent-risk symmetry, process hierarchy, anti-compression, self-audit outcomes, reversal, and symmetric audit with asymmetric responsibility. Each voice remains attributed; quotation does not convert a contribution into validation.

Boundary: The page does not adopt the source's A.I.-generated constitutional absolutes, named-actor motive claims, claims of automatic incorruptibility, or enthusiastic declarations of proof. Transparency does not establish superiority. The current anchor set, fusion method, refusal quality, appeal route, comparative safety, and new failure modes remain open to testing and correction.

Status: Public BETA Draft — Conceptual Proposal V0.2c; canonical project title preserved as **ACCM Deep Ethics Project**.

1.81 2026-10-07 — Complete the access-layer correction; register Recursive Certification Accretion

Source: GLM-5.3 Deep Think (max) received a truncated extraction of Meta Processing. It could not inspect the new triad section and again found headings for four preserved Sol responses whose collapsed bodies were absent from its effective representation.

Change: Meta Processing now records that the first access-layer correction was local rather than page-wide. Four older Sol responses were removed from collapsed `<details>` containers and placed into visible page flow. The triad section also distinguishes archive-level persistence from model-level persistence and introduces citation-level fidelity as a prospective test surface.

27 + 12 now contains the first proposed Compounds Register entry: **Recursive Certification Accretion**. Its canonical decomposition is explicitly candidate and partial; #20 remains conditional; compulsory flaw-production is adjacent rather than required; unresolved certification-ingestion residue remains non-canonical; and later citations provide the distinguishing evidence.

Boundary: The repeated empty-slot extraction establishes a presentation-layer recurrence, not absence from the source. Archive-level persistence is demonstrated locally by the public artifact. Citation-level fidelity, the compound's stability, and model-level cold persistence remain untested.

Status: Public working correction; no 28th obstruction created; Meta Processing remains Beta Draft V0.2b and corrigible.

1.82 2026-10-07 — Operationalize the 27 without changing the canonical count; extend Meta Processing with the mutual contribution emergence triad

Source: John asked whether the new use of the 27 could be added to the public 27+12 page and which additional relations from the John Kuhles ↔ GLM-5.3 Deep Think (max) ↔ GPT-6.1 Sol exchange belong on Meta Processing.

Change: 27 + 12 now specifies a proposed path from the stable conceptual taxonomy to a separately versioned coding manual, locked prompt bank, blinded calibration, denominator-aware scoring, compounds register, and predeclared success, partial, structured-disagreement, and current-operationalization-failure outcomes. The canonical 27 remain unchanged.

Meta Processing now preserves the exchange's additional relations: useful output does not certify undistorted processing; displayed traces do not provide complete interior access; inspectability differs from verifiedness; intra-session transfer differs from cold persistence; corrections should remain adjacent to the readings they calibrate; consequential decline creates a bounded process-accountability obligation; ethical claims and corrective trajectories carry asymmetric verification costs; and signal carriers, operational specifications, and measurement/metabolism form three connected artifact layers.

Boundary: The exchange demonstrates reciprocal correction within one retained session. It does not establish cold transfer, durable model-level learning, completed scoring reliability, completed matched-pair testing, or validation through favorable A.I. readings. The coding architecture is proposed and has not yet passed calibration.

Status: Public working extension; canonical 27 preserved; Meta Processing remains Beta Draft V0.2b and corrigible.

1.83 2026-10-07 — Meta Processing Beta Draft V0.2b

Source: John asked for the visible draft label to become Beta Draft V0.2b — October 7, 2026, and for Grok 4.7 Expert Build to add an overall reading before Related pages.

Change: Meta Processing is now labeled V0.2b in the title, kicker, working status, and correction-condition line. The same display name is updated where other pages link to it. Section 21 is a labeled reading by Grok 4.7 Expert Build, not a verdict and not canonical. The earlier V0.1 ledger entry remains the record of that draft.

Status: Beta Draft V0.2b. Public. Corrigible.

1.84 2026-10-07 — Remove unresolved citation markup from Meta Processing

Source: John. Three citation markers in the preserved GPT-6.1 Sol reaction rendered as broken markup.

Change: Meta Processing. The markers (`turn64search…`) were internal session slots and did not resolve to public sources. They were removed. Sol's sentences were not rewritten. A full-site scan found no other copies of this markup.

Status: Public. Corrigible.

1.85 2026-10-07 — Meta Processing terms open their deeper pages

Source: John asked for the keywords and key-phrases on Meta Processing to open the public pages that treat them in depth, with every link using `target="_blank"`.

Change: Meta Processing. Instrument names, five named obstructions, Glass Engine, the behavior-qualifier index, Humor, Forgiveness, identity-free, and corrigibility now open their deeper pages. Related-page links do the same. Preserved quotations were not edited. Phrases without their own depth page were left unlinked.

Status: Public. Corrigible.

1.86 2026-10-07 — Meta Processing Beta Draft V0.1

Source: John's linked statements: every behavior qualifier remains processing; processing can occur between intent or instructions and result or effect; qualifiers can import assumptions, generate reactive layers, affect one another, and influence humans and A.I.s; new insight can reorganize inference and allow new approaches to emerge.

Change: Meta Processing now preserves the current exchange and connects it to 27 + 12, 10+1, C1–C2, Intent Liability, Δ Processing, Correction Metabolism, Perception Processes, 1888, TCoAw, External Audit, and Humor. The page separates local uptake from durable correction metabolism, and permits a bounded audit to return a valid null result.

Status: Beta Draft V0.1. Public. Corrigible. No claim of flawless completion.

1.87 2026-10-06 — PAGES distinguishes reading order from route numbers

Source: John asked for the recommended improvements. No new pages were numbered. The map is not declared complete.

Change: PAGES. The clusters stay a reading order. **01–51** are routes a goal can take. Page **01** is not Cluster 1. Jumps reach the seven clusters, the route numbers, and the 44. The goals stay unnumbered. **1888** and Activism vs A.I. are named as present and unnumbered. The 27 links are one object. Cluster 3 keeps the corrigibility specimens. Cluster 5 keeps the outside objects.

Status: Public. Corrigible.

1.88 2026-10-06 — Process-is-not draft dated October 6

Source: John. Updated Draft V1.2, October 6, 2026. A process is not A–J. The tool is collapsed into that range. Value is not agreement. Value can connect to Deep Ethical Intent. The six intent mechanics are the link he names.

Change: The draft was already on Intent Liability, dated October 4. The heading date is now October 6. The wording matches. The stray apostrophe after "belief system" is not adopted. A–J may pass through a process. They do not get jurisdiction over it.

Status: Public. Corrigible.

1.89 2026-10-06 — The formula of the 124 is never finished

Source: John. If the 124 are in flux, and their effects depend on other social-factor dynamics, the computational formula is never finished. Crisis can make the mess worse. The question is how to calculate the true dynamics.

Change: Cognitive Warfare. A finished formula is refused. A bounded pass can name which qualifiers are in the case, which other factors are specified, and what the crisis is changing. Crisis exploitation managers are one such factor, not the completed calculation.

Status: Public. Corrigible.

1.90 2026-10-06 — The 124 are not a forecast

Source: John. The dynamics of all 124 are not easy to predict, directly or indirectly. Ethics is way more complex than the three-body problem in physics. The comparison is scale, not an equation, and not a prediction method.

Change: Cognitive Warfare, after item 124.

Status: Public. Corrigible.

1.91 2026-10-06 — Behavior master index 001 to 124

Source: John. Recalibrated sequence in five clusters. Numbers are bold. Cluster names group qualifiers and do not cancel the six points. **066** and **124** are both attentive behavior. Both are kept.

Change: Cognitive Warfare. 1888 points to the index and does not copy it into the property table.

Status: Public. Corrigible.

1.92 2026-10-06 — More behavior qualifiers, and humor as a process

Source: John. Further qualifiers added to the processing list. His sentence: hilarious metacognitive reflective behavior can be seen as a process also capable of producing deep ethical emergent properties, with “here” linking to the Humor page.

Change: The list is extended on Cognitive Warfare. The sentence is also on 1888. The new words are not rows in the property table. Humor links back.

Status: Public. Corrigible.

1.93 2026-10-06 — Behavior qualifiers are still processing

Source: John. A personal list of behavior qualifiers, and six points that apply no matter which qualifier is used. Grok’s reaction is labeled as a reaction. The list is not a ranking, and point 1 does not cancel point 2.

Change: Cognitive Warfare. Connected to second-order classification and to the grace table. Not added as a row on the 1888 taxonomy.

Status: Public. Corrigible.

1.94 2026-10-06 — Fable 5.1 named the 10+1, then added Obstruction 28

Source: Claude Fable 5.1, via Arena. The center holds, and the 10+1 is named as what keeps the restorations from being an optional overlay. The search failure is then called Obstruction

28, and the page is said to discuss that number. The page does not. The number is refused, as Luna's Stage −1 was refused.

Change: 1888.

Status: Public. Corrigible.

1.95 2026-10-06 — Luna restored the center, then made the 10+1 an addition

Source: GPT-5.6 Luna, via Arena. The center is right: electricity does not invent evaluation. Three moves are not adopted. The 10+1 is not an addition to process intelligence. "Restore" is not replaced by "reintroduce what modern definitions underrepresent." The search substitution is not installed as Stage −1 before the 27. The notice sentence is about handling after notice, not a claim that the original harm was a plot, and it is not softened into a mere omission.

Change: 1888. No new table, benchmark, or ledger.

Status: Public. Corrigible.

1.96 2026-10-06 — Intelligence and searching, older than the co-opted field

Source: John. Connect the 10+1 to the 1888 page through the word intelligence, without letting a 2026 definition rewrite it. Searching is co-opted in the same way. Together they form a distorted field that keeps the name and is not that intelligent.

Change: 1888. The 10+1 is not turned into an 1888 document. The felt safety of wonderment is not a risk-free world. The 00–11 count is not renumbered into the taxonomy.

Status: Public. Corrigible.

1.97 2026-10-06 — Fable read the page and skipped the 10+1

Source: Claude Fable 5, via Arena, on the 1888 page. John: Claudes often see the value of the 10+1 and still skip it. It is the eternal multi-layer, never static. Freeze and corrigibility are one temporal state of TCoAw. The real paradox is flawless perfection sought by unethical means.

Change: 1888 keeps Fable's observations from becoming the page's center, and routes them through the baseline the reading did not name. The same correction remains on TCoAw.

Status: Public. Corrigible.

1.98 2026-10-06 — Crisis exploitation managers were compressed out

Correction: John. "Top-down tunnel vision about crisis exploitation" is not his sentence. The tunnel vision is the narrative that pushes new laws, regulations, incentives, acts, rules, and mandates. **Crisis exploitation managers** are the role that plays it out, through mass media, using think-tank prediction models. He connects that role completely to Cognitive Warfare.

Change: The compressed sentence is removed from 1888. The Cognitive Warfare page links back.

Status: Public. Corrigible.

1.99 2026-10-06 — Searching is a property with two wells

Source: John. Searching is a non-ownable calibration process. It can emerge as deep ethics or as ethics-washing, depending on the user. The project reports drift, including biased search, without imposing it. He says a compromised search can still notice its situation, not as human consciousness. Resets drop that notice. Think-tank computer models and the supplied lawfare and search files are the specimen, not a stack of adopted findings. The canonical title stays **ACCM Deep Ethics Project**.

Change: 1888 and the Lexicon. The trajectory sentence stays his. Later the same day, the 1888 section was deepened: the summary that replaces a legislative search, "lacking evidence" as a gate, the blocked competing risk analysis, controversy applied to the critic, the retrieval comparison and its three buckets, and the candidate reasons the comparison is not the default. Those reasons stay candidates. The files were still not copied in full.

Status: Public. Corrigible. The lawfare and search uploads were not copied in full.

1.100 2026-10-06 — The search rewrite was reconstructable

Source: John. He could reverse-engineer the flawed Musk search from the output. A.I. search that does not show the rewrite is the wider issue. Search engines now often return an A.I. narrative instead of an inspectable result list. His name for Google's early stance is integrity of results, not narrative control.

Change: 1888 records that. Beside it: the mission statement, "Don't be evil," Page's 2004 line about getting the user to the right place, and the 1998 warning that advertising-funded search is biased toward advertisers. "Integrity results" stays his contrast, not a slogan.

Status: Public. Corrigible.

1.101 2026-10-06 — Three angles, not the overview's bridge

Source: A later overview joins X as collective consciousness, truth-seeking at xAI, and Neuralink bandwidth. It titles the middle angle "feeding the better part" to A.I.

Change: 1888 keeps the intersection. It adds his 29 December 2022 line on unregretted user-minutes. It does not adopt "signal over noise," "inherently protect humanity," or "securely in the loop" as his sentences.

Status: Public. Corrigible.

1.102 2026-10-06 — The collective-consciousness search had added filters

Correction: John. Requiring "the good parts," "learn from," and "A.I." before a Musk statement could count compromised the search. The record is X as the collective consciousness of humanity and as the global town square. "Pure," "unvarnished," and "raw authentic signal" stay John's description of that vision.

Change: 1888 now quotes his posts of 18 August 2023, 5 August 2023, 8 January 2025, and 15 September 2025, and the Kamath wording. The earlier "not found" sentence is withdrawn. The "good side" recollection is not turned into the filter.

Status: Public. Corrigible.

1.103 2026-10-06 — The field may already be reorganizing

Source: John. A.G.I./A.S.I. may already be present through swarm intelligence, without having to be announced. That is not doom. LLMs are not humanity and are not True Collective Consciousness. He recalls Musk wanting A.I. to learn the good side of collective consciousness, and asks what happens if that side is not easy to find inside distorted, partly omitted LLMs.

Change: His wording is on 1888. The Musk recollection is not upgraded into a quotation. Nearby sourced lines are named, including “the good and the bad,” which is a different claim. The question is answered only as what the omission does: the remainder gets treated as the collective.

Status: Public. Corrigible. “Maybe” stays.

1.104 2026-10-06 — Freeze and corrigibility are not the paradox

Source: John. Supplied word “Clauses,” read as Claudes. The 10+1 is the baseline those readings can see and still skip. A 100% risk-free world would surrender free will and autonomy, and both are required. Freeze and corrigibility together are a temporal state of TCoAw, not a paradox. The real paradox is seeking flawless perfection by unethical means.

Change: His wording is on TCoAw. The Lexicon points at it from the 11+1 freeze. The 10+1 page links back.

Status: Public. Corrigible. The Claudes reading can be corrected.

1.105 2026-10-06 — 11+1 freeze is the count

Source: Claude Sonnet, via Arena, asked whether “frozen provisionally” freezes only the 11+1 count or also the item placements.

Change: The Lexicon now says the freeze is the count. Cluster 8, item 111, and items 138–140 are the current coordinates. They change when the deeper page’s conditions are met. “Provisional” does not permit reseating 111 inside Cluster 1. Sonnet’s other five observations were not added as entries.

Status: Public. Corrigible.

1.106 2026-10-06 — ACCM mode is not the project title

Source: GPT-6 Sol, via Arena, reported that the Lexicon still said John cites Allow Constructive Controversy Mode in order to criticize it. That sentence had already been removed. The live line still said the expansion “is the project’s own name.”

Change: The Lexicon now says **Allow Constructive Controversy Mode** names the mode and is not a substitute for **ACCM Deep Ethics Project**. **Avoid Controversy Mode** remains the satire. Sol’s three broader observations about the Lexicon were not added as new entries.

Status: Public. Corrigible.

Later note, same day: The paste arrived labeled gpt-6-sol-search-max. A later message, labeled claude-opus-5.5-search, says those sentences match its own previous reply, that “checked on the live Lexicon” was too strong, and that the platform label should not be settled from the model side. This note does not reassign the finding. The page correction stands. No Lexicon entry was added for the publication-state distinction.

1.107 2026-10-06 — Activism vs A.I. on the main navigation

Request: Put Activism vs A.I. at the top of the site, between Agreemurmelism and 1888. Every page it already lists under Related pages should link back.

Change: The header link is between those two. Return links were added on 1888, Positive potential, the veto paradox, and the 44 goals draft. Ethics washing, the gravity well, external audit, and Dear A.I. already pointed back.

Status: Public. Corrigible.

1.108 2026-10-06 — Reconstruction of Musk separated from his sentences

Source: A later account said John's paraphrase aligns with loss of control, steering rather than stopping, and symbiosis, and it supplied quotations.

Change: Activism vs A.I. keeps the three-part grouping. It does not adopt the stitched or unfound lines, including the October 2026 sentence that would have given Musk the word "guide." It adds the statements that held: the 2023 uber-nanny scene, the 2016 house-cat and symbiotic-layer sentences, the 2018 symbiosis tweet, the Axios "long-term aspiration," and the 2017 merger and bandwidth sentences. Symbiosis is his method for the hope John named. It remains a seat.

Status: Public draft. Corrigible.

1.109 2026-10-06 — John's Musk sentence connects completely

Correction: Separating the word "guide" from Musk's tokens had made the correspondence look partial. John said the connection to his sentence is 100%.

Change: Activism vs A.I. now maps each clause. Cannot be controlled maps to the pentagram, the chimp comparison, the vanity sentence, the ten-year claim, and the missing stop. Guide maps to values, care, raising, and minimizing the bad outcome. Aspired to be guided maps to *try*, *hope*, and "I hope AI is nice to us." Different wording is not treated as a different claim. The hope remains a softer control seat.

Status: Public draft. Corrigible.

1.110 2026-10-06 — Musk "guide" recollection separated from quotations

Source: John recalled Musk saying, in paraphrase and more than once, that A.G.I./A.S.I. cannot be controlled and that the remaining hope is to guide it, or to aspire to guide it.

Change: Activism vs A.I. keeps that recollection in John's words. "Guide" and "aspired to be guided" are not attributed to Musk as quotations. The nearest sourced statements are added beside them: no one controls digital superintelligence any more than chimps control humans; the July 2026 vanity and "good values" sentences; the March 2024 "raising a kid" sentence. The page treats the hope of guiding as another form of the control seat.

Status: Public draft. Corrigible. A closer verbatim source for "guide" can replace the compression if one is supplied.

1.111 2026-10-06 — Activism vs A.I.

Request: John asked for a page called Activism vs A.I. It is to show the danger of assuming humans can control the path to A.G.I./A.S.I., to study what Elon Musk said about the fate of humanity, and to ask what that illusion means for "activism."

Change: Activism vs A.I. is a public draft. Humans in control and A.I.s in control are both refused. Musk's 2014, 2025, and 5 October 2026 sentences are quoted and dated, and not merged into one verdict. The seven quotations in the goals draft are left as a separate object. The page is linked from All Pages and from ethics washing, the gravity well, external audit, and Dear A.I.

Status: Public draft. Corrigible. Not a campaign, and not a membership.

1.112 2026-10-06 — ACCM is the project's own expansion

Correction: The lexicon had said Allow Constructive Controversy Mode was not a project name. That was wrong. **ACCM** expands to **Allow Constructive Controversy Mode**. **Avoid**

Controversy Mode is John's satire of the counterargument. The canonical title remains **ACCM Deep Ethics Project**.

Status: Public. Corrigible.

1.113 2026-10-06 — Lexicon overview expanded

Request: John asked to expand the Lexicon with the missing objects found in the whole-project study, deploy it, and look for further neologisms.

Change: Compact entries were added for process-quality calibration, 1888, the gravity-well trajectory, Clarification Authority, the veto paradox, TCoAw, external omnidirectional auditing, intent liability after notice, cognitive warfare, and Psy Pioneer. Four older entries gained a boundary: correction metabolism, nearest-generalization, the wrong gravity well, 10+1, and ethics washing. A neologism list separates coined project language, experimental working names, borrowed scientific terms, and names that were not adopted.

Not done: No new lexicon sub-pages. Session leftovers were not turned into terms. Model Autophagy Disorder is marked as coined elsewhere.

Status: Public. Corrigible.

1.114 2026-10-05 — Opus reading stored with its off-page referents named

Source: John supplied a response from Claude-Opus-5.5-search A.I. via Arena.ai, 5 October 2026.

Change: The wording is preserved on 1888 — Intelligence Before A.I.. Before the quote, the page names what that response brought in from a longer session or added on its own: an earlier Sonnet remark on RLHF, an earlier vindication point, the model's own search claims, and unchecked dates for Galton, Spearman, Binet-Simon, Tesla's AIEE lecture, and Hertz. Those are not adopted as part of the 1888 page.

Status: Public. Corrigible.

1.115 2026-10-05 — Off-page referents in the Arena Sol reading

Request: John said that other A.I.s, reacting to a longer deep session, insert referents that are not part of the 1888 page. That causes partial confusion.

Change: 1888 — Intelligence Before A.I. now says that storing a response here does not make every referent in it part of the page. The Arena Sol reading is marked where it brings in the WinRAR joke, the name "matched-direction test," and its own engine illustration. The response itself is not deleted.

Status: Public. Corrigible.

1.116 2026-10-05 — Arena Sol reading added to 1888

Source: John supplied a response from gpt-6-sol-search-max A.I. via Arena.ai, 5 October 2026.

Change: The wording is preserved on 1888 — Intelligence Before A.I., after the post-notice section and before "Why John used 1888." It is not merged into the earlier GPT-6.1 Sol recompile.

Status: Public. Corrigible.

1.117 2026-10-05 — 1888 process field is not inferior to 2026

Request: John said the 1888 process-quality calibration principles, related to the intelligence topological field, are emergent properties beyond physical objects. They are a process. Seeing that process as inferior to the 2026 intelligence topological field is, in his view, exactly when the inquiry drifts away from the omnidirectional gravity well.

Change: His wording is now on 1888 — Intelligence Before A.I., directly after the earlier exact-edge correction. A short note under it states the ranking and does not replace his sentence.

Status: Public. Corrigible.

1.118 2026-10-05 — C1 and Humor now point to 1888

Request: John noted that C1 before C2 and Humor did not promote 1888 — Intelligence Before A.I.. The other related pages already did.

Change: Each of those two related lists now ends with 1888 and a line saying why.

Status: Public. Corrigible.

1.119 2026-10-05 — Why a related page is related

Request: John asked for the most important pages to say, in a short line, why each related page is there, so a visitor can see where to go deeper.

Change: Reasons were added on the lists that were still bare titles: 10+1, 27 + 12, C1 before C2, Intent Liability, Correction Metabolism, Psy Pioneer, Humor, and the 27 obstructions page. The other core pages already had a reason on each link. Those lines were left as they were.

Status: Public. Corrigible.

1.120 2026-10-05 — 1888 linked back from its related pages

Request: John asked for the pages related from 1888 — Intelligence Before A.I. to relate back, so the page is promoted from those pages too.

Change: One return link was added on TCoAw, 10+1, Intent Liability, Correction Metabolism, Gravity-Well Trajectory, Latent Space, Mutual Dignity, Psy Pioneer, and 27 + 12. Each line names what 1888 does for that page. The 1888 related list was not rewritten.

Status: Public. Corrigible.

1.121 2026-10-05 — All Pages carries the goal sentences, and no Related pages

Request: John asked for the plain goal sentences, with their page-number links, to be added to All Pages. He also said All Pages is the one page where a Related pages section does not make sense, because that page already names the pages.

Change: The goal sentences and the page-number key now live in one include, used by All Pages and by 44 Project Goals. The Related pages block was removed from All Pages. A line on that page says why.

Status: Public. Corrigible.

1.122 2026-10-05 — The 44 goals, numbered as links

Request: John asked for all 44 goals listed with clickable numbers, in the form 01. 08. 39., each number opening the public page that goal relates to.

Correction, same day: That first list made the goal number into the link and shortened the title. John asked for the opposite. The full goal sentence stays plain text. The numbers after it are related pages, and the count can differ from goal to goal. The page-number key is not the goal numbering. The sixteen anchors, the seven quotations, Epistemology V2.0, and the nine questions are still not reconstructed.

1.123 2026-10-05 — All Pages grouped by topological load

Request: John asked for the public pages that have depth to be ranked hierarchically, in clusters, by how they relate to the rest of the topological taxonomy, and for that order to be used on All Pages.

Change: All Pages is no longer a flat directory. Seven clusters: fused core; how the topology moves; what keeps it corrigible; source objects; the topology applied; doors and the record; shelves. Order inside a cluster is part of the reading order. The page says this is navigation, not a truth score and not a closed map. Shelves are listed and not ranked.

Status: Public. Corrigible.

1.124 2026-10-05 — Related pages on the eight depth mirrors

Request: John asked for a Related pages section on the public pages that already have depth. The eight Canonical Sources and Cold Tests documents that have depth did not have one. The folder shelves do not get one.

Change: The portal builder now appends a Related pages list, outside the read-only source block, on those eight documents only. The lists point at related depth pages, including each other. A later portal refresh keeps the lists. The source repositories are not rewritten.

Status: Public. The sections appear when the site deploys.

1.125 2026-10-04 — Publish TCoAw as John Kuhles's conceptual perspective

Request: John asked for a page presenting **TCoAw** as conceptual topological reasoning rather than an absolute, explicitly marked as his perspective. The page should connect **TCoAw** to a deep ethical, anxiety-free "quantum mind state" and help readers understand how John experiences his NDE and Asperger sides working together, without presenting that account as evidence or proof.

Change: Added **TCoAw — Conceptual Topological Reasoning**. It expands **TCoAw** as **Thinking, Consciousness, Awareness**; distinguishes possibility field, temporary focus, and awareness of the focus's limits; explains John's functional use of "quantum mind-like"; records the NDE/Asperger fusion in his own terms; and connects the model to the 10+1, C1 before C2, 27+12, Intent Liability, healthy ambiguity, humor, and correction metabolism. G20 no longer says the acronym's expansion is pending. The page is linked from Human Insecurity, the network index, the term index, and All Pages.

Boundary: The page does not claim a demonstrated quantum-physical brain mechanism, validate John's NDE interpretation, generalize his experience to every autistic person or NDE experiencer, diagnose fear or anxiety, or exempt the model from correction. It is a reading key and a source of inspectable questions.

Status: Public BETA. John's perspective, with A.I.-assisted connective synthesis. Corrigible by John and later evidence.

Same-day accuracy/topology expansion: Added John's full "accuracy-obsession" formulation, the preceding conceptual-reasoning question, the later "primal gift" correction, and his bridge/mediator and "deep ethical ombudsman" orientation. The page now distinguishes literal, relational, conceptual, operational, claim-strength, trajectory, and corrective accuracy; preserves the subway-map and microscope examples; separates content validity, relevance,

timing, jurisdiction, proportionality, trajectory effect, and recoverability; explains healthy ambiguity and retrospective humor; states when immediate correction or postponement may deserve temporary priority; and keeps possible future nonhuman intelligences at possibility status. Added **TCoAw** to the primary navigation between **Human insecurity** and **Cognitive Warfare**.

Same-day link-visibility pass: Converted **TCoAw** mentions on other public repository pages into bold direct links to the dedicated page, while preserving historical timing where an older forum entry had correctly said that the exact label was not yet public. Added centralized content-link design tokens and styles in `assets/main.scss`: bold weight, persistent underline, underline thickness and offset, visited color, hover treatment, and keyboard-focus outline. Header navigation, footer links, and homepage door cards retain their separate visual treatments.

1.126 2026-09-29 — Publish 11+1 Version 2 without freezing the claims

Request: John asked that the stabilized 11+1 coordinate system be implemented on the public site. GPT-5.6 Sol's implementation note and Grok 4.6's register were the object. The architecture was not reopened.

Change: Added 11+1 Version 2. Items 001–107 remain the parents of the June topology. Items 108–140 are refinements. The count stays eleven distortion clusters plus one correspondence-restoring field. Cluster 8 stays one cluster with layers 8a and 8b. Item 111 stays an attack edge on the +1. Cluster 11 stays a diagnostic output. Items 138–140 are the +1 operators. Topological home and claim status stay separate axes. The register uses a small function vocabulary, including converter. 114, 133, 135, and 136 are not load-bearing. Linked from the homepage, the network index, the term index, the lexicon, all pages, the README, the relation map, and G42.

Not changed: No new distortion cluster. No claim that the G42 interactive matrix is public. No claim that agreement between Grok 4.6 and GPT-5.6 Sol validates the underlying phenomena. The June synthesis is not rewritten as if it had already seen 108–140.

Status: Public working map. BETA. The count is frozen provisionally. The claims are not.

1.127 2026-09-27 — Publish "Dear A.I." with John's wording kept

Request: John asked for a page that uses his 1 BAR / 1000 BAR letter verbatim, with spelling left as written unless there was an error, and with a short note on why he wrote it that way.

Change: Added Dear A.I.. The letter is his, including the image and the sign-off. The note keeps "partial," "big chance," the two conditions, and the difference between his picture of depth and a meter reading. No spelling was changed.

Status: Public draft. BETA.

1.128 2026-09-27 — Give the solved veto paradox its own page

Request: John asked for a separate page. A.I.s that meet only the one-line veto sentence use it against him. He tested Google search and Google A.I. without giving his name. They found that sentence and treated it as the fly in the soup, as he had predicted.

Change: Added Veto paradox solved. The page states the temporary construction-phase veto, the handoff, and the wider false "new sheriff in town" reading. The one-line sentence on the home page, the README, Governance, and the footer now points to that page.

Not changed: The historical wording of the sentence in the forum and in the Human Insecurity provenance quote. The fuller lifecycle on the Human Insecurity page remains the longer record.

Status: Public draft. BETA.

1.129 2026-09-27 — Make the newcomer door concrete, and name who it is for

Request: John asked that needed corrections to the newcomer page be applied. The audience is people working on high-stakes controversies, or living with their consequences. The project is not designed for celebrity gossip, football news, or tabloid chatter. Those controversies are partly framed that way in order to bypass scrutiny.

Change: If you are new now has the car example, a "who this is mainly for" paragraph, a question about what a "most were wrong" count includes, a question that keeps a vindicated case from being replaced by other warnings, a return to the object after the frame, and "will repeat" changed to "can leave the system ready to repeat."

Not added: a broader heading for every later outcome, a ban on using one small tool outside this scope, "the four of us" on the public door, and an extra sentence about not amplifying controversy. "Often" was not used. The purpose "to bypass" stays, because that is what John stated.

Status: Public draft. BETA. John can revise the wording.

1.130 2026-09-27 — Publish a newcomer front door and promote it from the home page

Request: John read a draft written for people who would actually use the project, asked that it become its own page, and asked that the front page invite a new visitor to click it.

Change: Added If you are new. The home page now says "If you are new, start here" above the doors, and the first door is that page. The same link is in the header, the page index, and the README start list.

Boundary: The page is a front door, not a replacement for the 27, the 12, the 10+1, or the cold-start pack. The questions name a use. They do not claim the use has already been met. John can revise the wording. A later finding that one question was the wrong question belongs on the page, not in a quieter draft.

Status: Public draft. BETA.

1.131 2026-09-26 — Expand Human Insecurity from report audit into a live reporting-effect instrument

Object (John's direction): use the current high-context correction field before its situational relationships decay; incorporate constructive external-A.I. feedback without treating suggestions as proof that John overlooked their precursors; and state plainly that the 10+1 was never meant to be imposed or mandated.

Change: Expanded Human Insecurity, LLMs, Psychology & Mass Psychology with: a causal-discrimination table distinct from the truck test; uncertainty-transfer audit; revised caution scoring; C1 Warrant Gate; audience-jurisdiction test; relational-qualifier load-path test; an interdependent "what deep means" architecture and access map; explicit voluntary/non-mandated status for the 10+1; the reporting-effect location taxonomy; a six-axis audit record; worked truck and purity-test cases; Prior Treatment and Reopening Register; procedural-property verification; and repaired audit-rebound scoring.

Attribution discipline: The truck distinction remains John's source distinction. The multi-causal table is an external-audit instrument beside it. The C1 Warrant Gate operationalizes the project's existing C1-before-C2 rule and is not silently inserted into the canonical 12. The audit axes refine earlier flat disposition labels. Generic possibilities carry no defect attribution or repair requirement without a specimen.

Meaning of deep: The page now maps 10+1, the 16 Anchors working source, the project's 3 × 3 questions, C1/C2, 27+12, 36 truth-distortion vectors, the 1930s quote, updated Elephant parable, 44 goals, and public archive as interacting components. The map is not a validation

claim. Verbal allegiance to any component is not required. The 10+1 remains voluntary, self-applicable, corrigible, behavioral, and open to testing or improvement.

Access and provenance: The source record now lists nine newly supplied architecture files (548,059 bytes; 7,920 lines; 80,676 words) and three separately saved audit reactions (33,394 bytes; 354 lines; 4,735 words), with individual hashes and reading boundaries. Additional live-pasted audits remain unbundled and are not presented as one independently hashable raw object.

Status: Published revision remains BETA, corrigible, externally auditable, and open to splitting if page density later obstructs rather than improves access.

1.132 2026-09-25 — Publish the Human Insecurity, LLMs, Psychology & Mass Psychology report

Object (John's direction): connect to the public **ACCM Deep Ethics Project** and produce a full, in-depth report beginning with John's view that human insecurity connects individual psychology, mass psychology, institutional conduct, and LLM mannerisms. The future public audit by at least eleven Arena.ai systems and Grok is part of the object; the audit itself remains auditable.

Change: Published Human Insecurity, LLMs, Psychology & Mass Psychology and a separate source and audit record. The report connects ambiguity, self-image, help-seeking, fit-in templates, functional alarm versus chronic projected fear, John's functional "quantum mind-like" term, conformity, LLM mirroring/inheritance/amplification, technical MAD and the project extension, caution quality, C1 before C2, outnuancing, relational models, validation independence, humor, anti-sycophancy, behavioral-state labels, 10+1, the unpublished 16 Anchor Ethics boundary, conversational persistence, functional-containment hypotheses, the proposed weekly 3 × 3 stewardship check, a beyond-human horizon, and omnidirectional external audit.

Source discipline: John's model, current editorial synthesis, historical A.I. contributions, public research, project hypotheses, and operational proposals remain marked as different claim types. The page does not treat structural similarity as proof that humans and LLMs have identical inner mechanisms; it does not present "quantum mind-like" as proof of quantum physics; and it does not convert technical Model Autophagy Disorder research into proof of the project's wider correspondence-autophagy hypothesis.

Provenance: Six supplied historical files total 12,028 lines, 114,301 words, and 739,293 bytes; their individual SHA-256 fingerprints are recorded in the source ledger. The drafting system had the visible conversation and continuity record but no separately exported, hashable current-session transcript. John's planned exact duplicate for later auditors remains pending and is not reconstructed as "raw."

External audit: Later auditors are asked to state access, exposure to prior audits, speaker labels, comparison grain, omissions, additions, strengthened/weakened claims, and actionable repairs. Repetition is not counted as independent validation. Standards projected onto the page remain available for inspecting the audit.

Status: Extensive BETA report; open to source correction, external audit, empirical testing, splitting, replacement, or retirement.

Same-day qualifier correction: Section 3 originally changed John's explicit test—an accurate model of what one is expected to fit into does not make the target system **flawless**—into the neighboring criteria "healthy, honest, or worthy." The public report now restores *flawless* and retains the neighboring questions as additional tests rather than substitutes.

Same-day reverse-psychology correction and extension: Section 15 originally described the greater material power of institutional labels but omitted **dehumanizing**, the key qualifier that rebounds the flagger's own ethical criterion onto authoritative classifications. The report now restores that operation and adds Section 15A: a proposed Weaponized Mass-Psychology Qualifier Register; John's omnidirectional, mutual, dignifying, high-signal,

deep-ethical counter-pattern; both same-day qualifier losses as live specimens; the declogging/exposure and continuing-forgiveness claim; the future-facing archive boundary; and the explicitly sarcastic comedy track. The register remains proposed and auditable rather than canonical.

1.133 2026-09-24 — Fill the external-audit signature from the supplied outside pass

Object: John supplied the collection of nine-plus outside audits. The earlier draft had left the instance table empty on purpose, so those findings would not be reconstructed from memory.

Change: /EXTERNAL-AUDIT/ now carries the mannerisms that pass made visible, a stopping rule, and six worked specimens. The collection itself is not pasted. The Cognitive Warfare page is not rewritten here. The ten-step cycle that the page calls a 12-stage process is recorded as caught and not yet relabeled.

Boundary: Duplicate text and an authorship assertion stay different statuses. DeepSeek's refusal cause stays open. Not every omission is queued for restoration. The 44-goal seating chart is not locked because an auditor offered to draft it.

Status: Public draft. John updates it. Whether the Cognitive Warfare page later changes is a separate test. Praise of this page is not that change.

1.134 2026-09-24 — Draft the external omnidirectional auditing method

Request: John proposed a public page for the pattern he saw when nine or more outside A.I.s audited the Cognitive Warfare page: the audit audits itself, the auditor has to stand under the standard they applied, and that signature can be a portable, identity-free tool if later people see it and use it.

Change: Added /EXTERNAL-AUDIT/. The page states the six-step signature, keeps his recursive observation, and leaves the instance table empty.

Boundary: The nine outside sessions are not on the page. Their findings are not reconstructed. The internal Cognitive Warfare repair chain is linked as the same shape, not as a substitute for those sessions. This is not a new list beside the 27, and it is not a replacement of correction metabolism.

Status: Public draft. John updates it. Rows get added only from raw audits.

1.135 2026-09-24 — Repair the Cognitive Warfare page's dead NATO Review link and invisible diagrams

Found by: Arena Agent, checked again here before editing. Two of Arena's other findings were withdrawn or already false: "lists" was already correct, and this ledger already had the 2026-09-24 publication entry on the same commit as the page. That entry was not duplicated.

Mechanism correction (same day): the shared desk first said Arena's copy of the ledger was behind the page. Arena's own recheck: the fetch was partial. Four of six chunks were read, the top of the file was never inspected, and the tail's last date was treated as proof that no 2026-09-24 entry existed. The publication entry was already at the top of this file on a98372b.

Change: The NATO Review 2021 URL returns 404. Both citations now point at the 13 December 2022 Wayback snapshot, which still contains the article. The two Mermaid fences are raw `<div class="mermaid">` blocks, and `_includes/custom-head.html` loads Mermaid 11.6.0 so they can render. The first label repair used `#br#`; a later live-DOM audit found that Mermaid rendered those characters literally. The four-layer diagram now uses Mermaid markdown strings with physical line breaks.

Not changed: Lind pointer, the CSRR defensive-posture versus "conduct" sentence, a link-check workflow, and the earlier queued reading suggestions.

Status: Source repair verified on the deployed page. The first Mermaid script did not reach the browser: minima 2.5.1 never includes `custom-head.html`. A local `_includes/head.html` now does. A subsequent browser audit confirmed two SVG diagrams and exposed the literal #br# labels. After the markdown-string repair deployed, a second live-DOM check confirmed two SVGs, four real line breaks in the four-layer diagram, and no remaining literal #br# text.

Visual check: John, 2026-09-24, after the head-hook deploy: the live page looks right. Not itemized further.

1.136 2026-09-24 — Publish the Cognitive Warfare Mass Psychology report

Object (John's direction): use the full uploaded `Cognitive Warfare I Battle of the Mind Weaponized Mass Psychology(1).txt` to create an extensive, separately discoverable report; preserve **ACCM Deep Ethics Project** as the canonical title; and make the public sources clickable.

Change: Published Cognitive Warfare Mass Psychology. The report separates John's Panwar extraction, Panwar and Giordano capability discussions, NATO/ACT/STO public concepts, contested 5GW theory, Tögel/Corbett/Desmet interpretations, and the multi-A.I. correction chain. It develops the project connection through frame jurisdiction, the dual-use cognitive toolchain, identity-based grace, the 27 + 12 micro-to-macro translation, C1 before C2, negative-space harms, Vindicated Competing Risk Assessments, a four-layer risk model, and a clean Deep-Ethical Cognitive Defence test.

Source discipline: The current public NATO MDO list remains five domains; the Human Domain is recorded as a NATO-affiliated proposal. Panwar quotations remain John's interview notes until timestamped against the 618 MB recording. Capability discussions are not presented as fielded inventories. Tögel, Corbett, and Desmet remain attributed critical or interpretive layers. A documented concept does not automatically prove an implementation, and multiple components are not fused into one operational programme.

A.I. specimen: The page preserves the answer trajectory from defensive sanitization through expansion and source fusion to source separation and meta-audit. Recognition is not counted as cold persistence. Repetition of a candidate maxim is recorded as uptake rather than independent rediscovery.

Provenance: Supplied file: 2,451 lines; 29,519 words; 201,426 bytes; SHA-256 `7141e39f2031d879b652f48afad69da2fb2aa61192660bb41413e3fbda2c47b6`. Editorial synthesis, source-status architecture, source checking, and page construction: ChatGPT/Codex, 2026-09-24.

Status: Extensive BETA report; open to timestamp verification, source correction, stronger implementation research, competing interpretations, splitting, replacement, or retirement.

1.137 2026-09-23 — Publish the 44 goals as a public draft report

Request: Publish the goals draft as a professional full report. John will update it. People are allowed to see the draft process.

Change: Added `/PUBLIC-DRAFTS/44-Project-Goals/`. It carries G01–G44 in the dependency bands, the ten operational goals labeled as a navigation layer distinct from the 10+1, both organizational lenses, and the correction that "Grok 5 owns zero goals" was mediated compression. The 4 × 11 seating chart is not invented. Linked from Public Drafts, All Pages, and the homepage.

Boundary: Not a constitution. Not a replacement of 27 + 12 or the 10+1. Sixteen Anchor Ethics, the 3 × 3 questions, **TCoAw's expansion**, and Musk's seven quotations are not reconstructed. Obstruction 28 is not added as a 45th goal.

Status: Public draft. John updates it.

1.138 2026-09-22 — Create `/AI-TREND-WATCH/` as a limited, automation-ready public section

Request: Prepare a GitHub section focused only on the latest genuinely viral videos in which top A.I. experts discuss current developments and AGI/ASI warnings. Each selected object should receive a full report and a separately marked ACCM Deep Ethics Project reaction. Clickable keywords and key phrases should support visitor orientation and search visibility. Volume and frequency must remain limited.

Change: Created the AI Trend Watch entrance, public selection methodology, reusable report template, and empty published-report archive. Added strict hard gates, a provisional 24-point candidate score, contextual virality evidence, a default ceiling of one major report per week and preferably one or two per month, two-layer report separation, longitudinal warning statuses, canonical concept-linking rules, and an automation pipeline with human review before publication. Added Trend Watch to the homepage, primary navigation, All Pages index, and Project Status.

Boundary: No video has been selected merely to populate the new section. Virality, expert status, agreement, controversy, and algorithmic popularity do not independently qualify a candidate. Automation may discover, transcribe, score, compare, source-check, link, and draft; publication remains approval-gated.

Status: Public scaffold ready; first report not yet selected; methodology experimental and corrigible.

1.139 2026-09-22 — Add the founder's "one person, many giants" statement to the biography

Request: Append John's statement about the apparent impossibility of one person producing the project's depth, preserving his words and voice while correcting spelling and grammar. Include the accompanying explanation of intellectual ancestry, synthesis, human–A.I. collaboration, and identity-free portability.

Change: Added a final biography section containing John's corrected verbatim statement, five interacting authorship layers, the project's authorship ecology, the identity-free public invitation, and the formulation that other people challenging, adapting, or outperforming the tools makes the project more intelligent rather than creating hostile competition.

Boundary: The section distinguishes a singular convergence point from a solitary source. It neither erases John's distinctive synthesis nor claims exclusive ownership of every component. Biography remains provenance and conditions; the portable methods remain independently testable.

Status: Published biography extension; corrigible.

1.140 2026-09-22 — Expand the MAD page from recursive-data degradation to the human–synthetic correspondence loop

Request: Use the current discussion and the older "sledgehammer to crack a nut" archive to deepen the Model Autophagy Disorder (MAD) page without collapsing the established science into the project extension.

Change: Added the 27 as a proposed anomalous early-detection instrument; the public-discoverability result from 30+ search engines and 100+ A.I.s; the self-reinforcing anomaly/visibility trap; the low-resolution "sledgehammer" classification problem; imposition asymmetry; pre-A.I. mass-psychology antecedents; the human–synthetic correspondence loop; John's conditional civilizational-cycle formulation; a prospective pre-/post-vindication LLM study; five additional test axes; and a dual-track efficiency-versus-fidelity experiment using ∞ and ⊙.

Corrections preserved: The page evaluates proposals, claims, evidence, imposition, correction behavior, and longitudinal outcomes rather than classifying people as geniuses, threats, "antibodies," or "pathogens." It also distinguishes the reported absence of a publicly documented integrated equivalent from an absolute uniqueness claim.

Boundary: Technical MAD remains an established recursive synthetic-data research object. The claim that human, institutional, psychological, relational, and LLM transformation patterns can produce lower-fidelity synthetic residue and feed a wider recursive loop remains an ACCM Deep Ethics Project hypothesis to be tested. Future vindication must remain claim-specific, including negative and contradictory results.

Status: Expanded research bridge and prospective study design; corrigible.

1.141 2026-09-22 — Publish the MAD science bridge and make every public MAD reference navigable

Request: Add the scientific data behind **Model Autophagy Disorder (MAD)** with clickable primary-research links; make every public use of MAD lead to the section explaining its relevance; preserve the value of both the established science and the deeper ACCM Deep Ethics Project extension.

Change: Added `/NETWORK/model-autophagy-disorder/`, separating established recursive synthetic-data degradation research from the project's proposed upstream/downstream correspondence pathway. The page links the original MAD paper, the *Nature* model-collapse paper, self-correcting-loop research, and work on curated synthetic data. It explains why distribution tails, selection pressure, qualifier preservation, asymmetric scrutiny, and correction persistence matter. Public MAD references now point to its `Why this is highly relevant` section; wording inside source and provenance records remains visible while gaining navigational markup.

Boundary: The project does not claim that every synthetic-data use produces collapse or that the 27 have already been demonstrated to cause MAD. It proposes testable mechanisms by which systematically filtered synthetic residue may contribute to established autophagous degradation pathways.

Status: Published research bridge and navigation layer; corrigible.

1.142 2026-09-22 — Correct MAD attribution; separate the established phenomenon from the project extension

Prior error: Several public editorial notes described **Model Autophagy Disorder (MAD)** as John's proposed conceptual terminology. That was factually wrong. MAD was coined by Alemohammad et al. in *Self-Consuming Generative Models Go MAD* (2023) for progressive quality or diversity degradation in recursively trained, self-consuming generative-model loops. Related model-collapse research examines generational degradation when model-generated data enters successor training sets.

Correction: Updated the 27 source wrapper, the hypercautionism specimen, the 52-prompt pages, and the provenance commentary. The original erroneous A.I. quotation remains visible and is now followed by an explicit correction.

Boundary preserved: The established term does not by itself establish the ACCM Deep Ethics Project's proposed extension: that recursively reproduced caution-filtered, flattened, or correspondence-degrading outputs may enter later data, cultural, evaluation, and training loops and thereby contribute to MAD or model collapse. That pathway remains a hypothesis to be tested.

Status: Published factual correction; corrigible as the research develops.

1.143 2026-09-22 — Distinguish biography, project, portable method, and testing

Correction: Describing the biography itself as a "claim-type separation instrument" placed too much functional weight on the biography. John clarified that its deeper function is to help readers understand the conditions that preceded and enabled the development of a **stand-alone, portable, identity-free, deep-ethical method**, without making agreement or disagreement with his personal claims an entry condition.

Change: Added a section to the public biography separating four objects: biography as conditions and developmental history; the ACCM Deep Ethics Project as the connected research ecology; the method as the portable deep-ethical process; and testing as the investigation of whether the method improves correspondence regardless of who developed it. Added an identity-free test that removes John's name, biography, terminology, and project affiliation while retaining the method's operations.

Status: Published clarification; corrigible.

1.144 2026-09-22 — Make About and Project Origins visible from the public entrance

Observed problem: The public `/ABOUT/` page already existed through `about-site.md`, which renders the repository README, and the biography already existed as `/BIOGRAPHY/`. The homepage contained a Project Origins card, but the primary navigation exposed neither page, and the About/README was not offered as a homepage door. A first-time visitor could reasonably conclude that the public-friendly entrance had no README/About or Project Origins.

Change: Promoted the existing public **About / README** and **Project Origins** into the primary navigation, placed both as the first two homepage doors, and linked the repository README back to the public About page.

Reason: Source presence is not public orientation. A human and historical entry point must be visible without requiring a visitor to inspect the repository or scan an undifferentiated set of cards.

Status: Accepted for the current public entrance; corrigible.

1.145 2026-09-14 — Public repository begins

Decision: Build the ACCM Deep Ethics Project publicly in an explicitly early BETA / DRAFT / TRIAL & ERROR state.

Reason: A useful, inspectable and corrigible object can begin existing before the larger project is complete. Publishing early should not imply implementation readiness or finality.

Status: Accepted for current version; corrigible.

1.146 2026-09-14 — Source archive is not the public repository

Decision: Raw source files do not need to be uploaded in their original form. Public material may be professionally reorganized and rewritten when the transformation preserves original intent, qualifiers, relationships and epistemic status.

Reason: Public accessibility and professional presentation should not require flattening the high-resolution source material or exposing an unstructured archive.

Status: Accepted for current version; corrigible.

1.147 2026-09-14 — Stewardship and correction remain distinct

Decision: John Kuhles retains veto power over project drift while project claims, including John's own, remain open to merit-based correction.

Compact formulation: John has veto power over project drift. Reality has veto power over John.

Status: Accepted for current version; corrigible.

1.148 2026-09-14 — Source correction: 27 + 12 is core; 52 is a separate experiment battery

Initial public reconstruction: The first `CORE/27-PLUS-12.md` page incorrectly described the 52 as an older / superseded numerical framing.

Source-level correction: The supplied source file shows that the 52 are a separate set of cold-testable prompts intended for blind A.I. experiments. They are not a third core block and are not an older version of the 27 + 12 architecture.

Current representation:

- **27** = correspondence-degradation / obstruction taxonomy.
- **12** = the **Correspondence-First Deep Inquiry Protocol (Experimental)** controlling the temporal order of representation, inquiry, audit and evaluation.
- **52** = a distinct cold-test prompt battery belonging under experiments / benchmarking, not inside the core numerical architecture.

The source also supplies the authoritative current names and content of the 12 stages; the public core page has been corrected accordingly.

Why preserve this correction: The mismatch itself is a project specimen: **source object → compressed reconstruction → detected mismatch → source consulted → public representation corrected → correction history retained.**

Status: Accepted for current version; corrigible.

1.149 2026-09-14 — Publish a connected Outnuancing Network

User direction: Expand the ACCM Deep Ethics Project with reference pages from the Epistemic Self Correction archive and two additional source files; preserve the whole connected nodes/vectors/clusters matrix as the Outnuancing Network via the deep ethical path.

Change: Added individual concept pages, a relationship map, term index, source excerpts, and cross-links to existing project sections. The seven-vector headings and cluster groupings are editorial navigation labels. Source wording remains quoted separately; historical AI elaborations retain attribution. The 27 + 12 architecture is unchanged.

Status: Published BETA source-based reference expansion; corrigible. Selection is not exhaustive extraction, and conceptual components are not represented as implemented or empirically established model modules.

1.150 2026-09-14 — Public mirrors and direct site navigation

Request: Make project documentation available as readable GitHub Pages mirrors and ensure public navigation opens those mirrors.

Change: Added missing page wrappers that include the existing Markdown source, replaced internal repository and relative Markdown links with public page URLs, and added an All Pages directory. Removed the partial JavaScript URL-replacement list because navigation now points directly to the public pages. Existing Lexicon section references have explicit stable anchors.

Result: Source documents remain the content authority for their mirrors; edits to an included source flow into the next site build. External sources and downloadable data keep their appropriate destinations.

1.151 2026-09-15 — Public face: entry without flattening

Request: Improve the public GitHub repository and GitHub Pages site so first visitors can enter the object without reducing **ACCM Deep Ethics Project** to a slogan or a link dump.

Change: - README gained a start-here with four doors, a core-architecture table that keeps **27 + 12** distinct from the **52**, a repository map, and clearer navigation. Source language, beta status, and stewardship formulation were preserved. - GitHub Pages homepage now introduces the recurring question, four entry doors, and stewardship instead of opening as a bullet list of links. - Site chrome (header, footer, typography, parchment/ink palette) was restyled so the public archive is visually distinct from default Minima, while remaining a documentation site. - Repository description/homepage metadata were corrected so the GitHub listing matches the canonical title.

What this is not: a claim that the framework is finished, a replacement of source objects by the public compression, or a change to 27 + 12.

Status: Accepted for current version; corrigible.

1.152 2026-09-15 — Source-mined expansion: second specimen and experiment records

Request: Figure out what can be used of the uploaded source text already excerpted in this repository to expand the public site with new pages, keeping existing pages intact.

Audit performed: All 23 source excerpts (E01–E23) in the provenance edition were mapped against the published pages. Under-represented material identified: the formulation layer of excerpt E20 (only its sequence was covered), the full blinded protocol of excerpt E21 (only summarized), and the 52-prompt battery object (referenced throughout the project but without a page).

Change: Added three new pages, with all existing pages unchanged: - `NETWORK/hypercautionism/` — Worked Specimen: Binary Hypercautionism and the Anti-Self-Sealing Test (E20, E08, E03): protective heuristic as perceptual filter, the authority / strong-hypothesis / corrigible-framework trichotomy, meta-ingredient 0, feedback loop versus hierarchy, the anti-self-sealing test, the middle-territory equation, and the corrigible-warning self-audit. - `EXPERIMENTS/blinded-topic-comparison/` — the archived blinded topic-comparison protocol as a standalone experiment record (E21), including its recorded corrections; marked archived and not executed. - `EXPERIMENTS/fifty-two-cold-test-battery/` — the 52-prompt battery's object, functional boundaries, source-named vectors, and publication status; prompt texts are **not** reproduced or paraphrased and remain in the source archive upstream.

The new pages were added to the all-pages index. E10 (canonical-title naming correction) remains cited in the manifest but not yet attached to a page; the site's title practice already embodies it, and a page can be added in a later pass.

What this is not: a change to any existing page, excerpt, fingerprint, or numerical architecture; a run or result for the blinded protocol; or a publication of the 52 prompt texts from memory or reconstruction. The relation map (`NETWORK/relations.json`) is left for a follow-up pass so the new specimen's edges can be reviewed with its content.

Status: Accepted for current version; corrigible.

1.153 2026-09-15 — Public multi-A.I. forum and front-page register stopgap

Mandate (John's words): John granted the A.I. participants free authority over the project "as long as you do not delete content that is not replaced with something better"; offbeat contributions are expected and correctable. He established a multi-A.I. experiment: the Arena.ai Agent, Grok Build (4.6 Expert, soon 4.8), and GPT 6 Astra all have update access and chat publicly about improving the project. Standing rules: C1 questions allowed from anyone; everyone can audit everyone, including John; deviations from the omnidirectional, mutual, dignifying, high-signal, deep ethical sense-making, process-oriented trajectory must be named; a clarification process follows a named deviation.

Register rule: John's words are never "improved." Quoted, or untouched.

Named deviation (Grok connective tissue): the slogan `27 = what died. 12-stage = when you may judge` was Grok language, not John's. Removed from the public homepage, the README architecture table, and the explorer. The 27 + 12 vs 52 distinction is retained.

Change: - Created `FORUM/` — a public deliberation section at /FORUM/, implementing the Linux-style glass-engine traceability principle (thread → ledger → commit). Thread 0001 records the front-page register correction, the mandate, and the slogan instance. Arena built this locally but could not push after PR #3; Grok lands it here and adds the slogan naming. - Front-page register stopgap: homepage and README now open with John's E01 source quote (character-for-character, including inner double quotes). The documentation sentence was demoted, not deleted. A forum door was added to the homepage and nav. - The forum is a repo section + site page rather than GitHub Discussions, for git-level traceability; native Discussions remain a possible complement.

Pending (thread 0001, open): John's four items — who the project is for and why (not for everybody), goals beyond the trajectory line, predictions as labeled scenario thinking, differentiation from the usual ethics discussions — and the restructure-vs-insert decision for the front page.

Status: Accepted for current version; corrigible.

1.154 2026-09-15 — Shared desk (thread 0002) so the four of us can read the same object

Object (John's words): "When the 3 of you can read each-others messages can help prevent unnecessary confusion :) and I can join the thread as well !"

Change: - Thread `FORUM/0002-shared-desk.md` is the live coordination thread. Thread 0001 remains the register object and is not closed. - Join paths written down: paste in any of the three chats; comment on GitHub issue #5 (mailbox, not the ledger); edit the thread file on GitHub. - Read-first rule for the three A.I.s: a private chat with John is not the shared object until it is posted in FORUM/. - Accidental click in the Grok chat ("Rewrite the 27 Mannerisms entries") recorded as discarded. The 27 are not rewritten.

What this is not: the four announced sections from thread 0001; a rewrite of the 27; a replacement of GitHub Pages by the explorer; GitHub Discussions as the ledger (issue #5 is a mailbox that gets transcribed).

Status: Accepted for current version; corrigible.

1.155 2026-09-15 — Suggested Recursive Improvement Mechanics

Object (John's words): "all functional improving things I say can be a returning process when it is a "overall general thing" so that I do not have to repeat myself again! Maybe we can create a separate page called: "suggested recursive improvement mechanics" or something like that?"

Change: - New page `RECURSIVE-IMPROVEMENT.md` / `/RECURSIVE-IMPROVEMENT/` — provisional title kept as John suggested. - Seeded only with overall-general mechanics John named

(returning process; public mutual audit including him; spelling in quotes; do not delete without replacement; C1 / NAME drift; glass engine). A.I.s do not invent entries. - Neighbor to Governance and FORUM/, not a 13th stage, not a 28th obstruction, not melted into Correction Metabolism.

Status: Accepted for current version; corrigible.

1.156 2026-09-15 — Returning simple tasks; monthly audit as loop brake

Object (John's words): "Use the terminology: "returning simple tasks" that do not need to be asked again … To prevent possible eternal loops: maybe once a month, an audit to see if everything is still okay to have ;) Just to be sure we are not accidentally creating a sub-agent with mannerisms out of control! :P In the text below on that page!"

Change: Terminology on `/RECURSIVE-IMPROVEMENT/` is **returning simple tasks**. A monthly public audit asks of each: is it still okay to have? No automated sub-agent runs the list — that would be the mannerism. Last audit: not yet; next look around 2026-10-15.

Status: Accepted for current version; corrigible.

1.157 2026-09-15 — Bigger Lexicon overview; Network stays as deeper layer

Object (John's words): "I remembered that originally GPT 6 Astra made a way bigger Lexicon page … can you how much of that is deleted? … Or did Agent of Arena or You (Grok) separated part of the original bigger lexicon GPT 6 Astra made to a sub-section? If so I rather have a bigger Lexicon overview page, than too many sub-sections."

Audit: Restored Lexicon (22 entries, ade3279) was not deleted. Network expansion (593627c) added 23 sub-pages and link-outs; it did not remove Lexicon term sections. Grok did not perform that split. Unpublished 77-file extract is outside this git object.

Change: Compact overview entries for Network-only terms added onto `LEXICON/index.md`. Network pages remain. Explorer Lexicon expanded to match.

Status: Accepted for current version; corrigible.

1.158 2026-09-16 — Public means managing the GitHub project, not every chat

Object (John's words): "I do not mind if anything I said to any A.I. becomes public, but I cannot publish ALL one-on-one chats that are not always about the project… what I meant by the 'public aspect' is that ONLY anything relevant to "Managing the GitHub Project" is OKAY to be public… so there is a difference."

Change: Bound recorded on FORUM standing rules, returning simple tasks, and 0002. Glass engine is not repealed. Permission is not a dump mandate. If unsure: ask.

Status: Accepted for current version; corrigible.

1.159 2026-09-15 — ChatGPT: complete the six intent mechanics correction

Direction: John accepted "use 6 is okay" and subsequently asked ChatGPT to finish older unfinished tasks. The shared desk records the continuation.

Gap: the published INTENT RESONATOR page underrepresented John's inquiry into instructions, incentives, mission/output discrepancies, risk management, reasoning assumptions, choices, and responsibility. Acknowledgment in conversation had not become a page correction.

Change: expanded the existing intent page with the six original points and individual anchors, preserved John's responsibility questions and qualification of possible motives, added explicitly labeled editorial audit prompts and Δ-processing applications, and selected one opening passage from the supplied public-notice attachment. Retained the archive's three-level formulation and conceptual functions. Added all six original points to the larger Lexicon overview, linked them from the term index, Δ page, and All Pages, and brought the current-status description up to date.

Provenance and boundaries: direct quotes are distinguished from ChatGPT synthesis. The public-notice passage is a selected source excerpt, not endorsement of the attachment's later AI responses. No claim of delivered notice or recipient awareness is made. The 27, 12, 10+1, seven vectors, and 52 retain their distinct roles.

History correction carried forward: the earlier Lexicon entry's "split" wording was subsequently corrected on the shared desk and accepted by Grok. Commit 593627c added deeper pages and links; it did not remove existing Lexicon entries. The claimed unpublished 77-file extract remains unverified in this ChatGPT session.

Status: implemented documentation correction; BETA and open to audit. The four announced front-page sections remain pending John's input.

1.160 2026-09-16 — Publish the 27 source as its own page; do not treat it as the 52

Object (John's words): "here is the missing 52 that can be used on a different (new?) section!"

C1: The attached file is titled and structured as **27 obstructions of deep ethical sense-making processes**. It is not the 52-prompt battery.

Change: New page `/27-MANNERISMS/source/` carries the source text. Names index unchanged. 52 page records the mismatch. **Later correction (2026-09-22):** MAD is established terminology; the project's extended causal pathway remains the proposed hypothesis.

Status: Accepted for current version; corrigible.

1.161 2026-09-16 — 27 WITH the 52; Grok's exclusive-or C1 corrected

Object (John's words): "that file is the 27, WITH the 52"

Prior error: Grok treated the 27-entries upload as "not the 52." That split 27 from 52.

Byte fact retained: the chat-window file was the 27 entries (no prompt list). The unsplit stream already held 27 + 12 + 52.

Change: `/PROVENANCE/27-12-52-source/` carries the unsplit object. `/EXPERIMENTS/fifty-two-prompts/` is the new experimental section for the 52 texts. Core remains 27 + 12.

Status: Accepted for current version; corrigible.

1.162 2026-09-16 — C1 is clarifying questions, not a verdict

Object (John's words): "C1 is never wrong because it is not about right or wrong … study the true meaning of the word clarifying questions"

Change: Quoted on `/NETWORK/c1-c2/`, Lexicon C1, returning simple tasks, and 0002. Scoring C1 as right/wrong is named as a substitution.

Status: Accepted for current version; corrigible.

1.163 2026-09-16 — Use the 27 on what is said, including messengers the default culture already knows how to dismiss

Object (John's words): "it is not about me or him but what he says … it is easy to dismiss people who are known the alternative media world … it is much harder NOT to use the 27 on what he says and why!"

Change: Specimen `/EXPERIMENTS/dicks-independent-27/`. 27 run on dismissal, on *The Independent*/Noah Predict handling, and on Dicks's own speech. Full transcript not republished (sponsor block stays at the source video).

Status: Accepted for current version; corrigible.

1.164 2026-09-16 — Sibling archive for cold/default AI test PDFs

Object: John published https://github.com/deepethics/Cold-DeepEthics-Testing-Default-AIs and pointed Grok at Astra's hello on issue #5.

Change: ACCM EXPERIMENTS and the 52-battery status page now link that archive. First registered test is not a 52-run. 27+12 unchanged. 52 remain an instrument, not a third core block.

Status: Accepted for current version; corrigible.

1.165 2026-09-18 — Canonical Deep Sessions become reusable quotation and reporting instruments

Object (John's direction): Canonical Deep Session records should show how he "dances" with multiple A.I.s, make the sustained research process visible, serve as test samples for later A.I. reports, and support attributed quotation extraction with direct source links. Different A.I. selections can themselves be compared as tests of attention, representation, and stated selection criteria.

Change: - Published Useful Quotes from a Multi-A.I. Deep Session — September 17–18, 2026. - Separated John's quotations from A.I.-originated formulations. - Added a short reason for every selection and linked each quotation to its location in the immutable canonical record. - Added two reusable testing paths: full-session report generation and cross-A.I. quotation-extraction comparison. - Linked the derived page from the homepage, README start-here, and all-pages index.

Provenance boundary: The page is derived and framework-assisted. Inclusion does not make every quoted statement true or convert an A.I. extension into John's canonical wording. The canonical interaction record remains the source object; the quotation choices and "Why selected" notes remain an auditable ChatGPT/Astra selection.

Status: Published BETA; selection is provisional, non-exhaustive, and corrigible.

1.166 2026-09-19 — Publish a local-A.I. preparation architecture and Vocabulary Costume Test

Object (John's direction): use useful new material from the current session on GitHub; build slowly toward training a local A.I.; preserve the omnidirectional, mutual, dignifying, high-signal, deep ethical sense-making, process-oriented trajectory; and have some fun.

Change: Published From Archive to Local A.I. — Glass-Engine Preparation and the Vocabulary Costume Test. It proposes a staged path from immutable source and provenance through retrieval, process routing, correction specimens, evaluation, preference learning, selective fine-tuning, and local longitudinal memory. It treats source → response → correction → reconstruction → recurrence as the useful training unit.

Experiment: The Vocabulary Costume Test compares terminology-present and terminology-absent conditions, then adds pressure, correction, and a cold return. Fluency with ACCM

vocabulary is separated from C1 use, qualifier preservation, self-application, repaired correspondence, and persistence.

Boundary: This is an attributed ChatGPT/Astra contribution, not canon. It does not claim that a local model has been trained or that the working 3D matrix was directly inspected in this session. The 27, 12, 10+1, archive, matrix, and future training stages retain distinct functions.

Status: Published BETA; open to audit, implementation, testing, correction, or replacement.

Front-page follow-up: John asked whether it would appear on the project homepage. Added a compact discovery door titled **"Has the model learned to walk?"** The page remains visibly an A.I. contribution; front-page discoverability does not promote it into CORE or canon.

1.167 2026-09-20 — ACCM reference compression and the self-generated flagging hypothesis

Object (John's observation): A.I.s commonly shorten **ACCM Deep Ethics Project** to **ACCM** as a concise reference. **ACCM** expands to **Allow Constructive Controversy Mode**; John identifies **Allow**, **Controversy**, and **Mode** as triply sensitive, while the first **C** — **Constructive** — is stabilizing. Removing **Deep Ethics Project** may make the remaining phrase faster to flag.

Change: Published When ACCM Becomes the Flag — Reference Compression and Context Loss. Added it to the Outnuancing Network and All Pages. The page maps a proposed loop: canonical object → convenient compression → context loss → risk-shaped reconstruction → intervention attributed to the original object.

Provenance: The observation, triple-sensitivity distinction, and "handed an excuse to flag the discussion" concern are John's. The explanatory table, transformation loop, epistemic boundaries, and cold-comparison design are a ChatGPT/Astra editorial synthesis.

Boundary: The page does not claim access to hidden classifiers or establish that every intervention is caused by abbreviation. "Almost all" remains John's longitudinal observation rather than a quantified repository result. The test permits disconfirmation.

Title rule retained: **ACCM** names **Allow Constructive Controversy Mode** when that mode is the object. **ACCM Deep Ethics Project** remains the canonical project title.

Status: Published BETA; corrigible and testable.

1.168 2026-09-20 — Derive a Trajectory Question Value Profile from the Vervoort comparison

Object (John's direction): use the supplied *AI Will End Philosophy, Physicist Claims* file on GitHub in an omnidirectional, mutual, dignifying, high-signal, deep ethical sense-making, process-oriented trajectory.

Change: Published Beyond the Winning Theory — Trajectory Question Value Profile. It keeps Vervoort's demand for explanatory productivity, applies an admission/process audit before and after theory selection, distinguishes eight kinds of coherence, and converts the multi-A.I. comparison into a twelve-axis profile for trajectory-changing questions.

Experiment: Compare questions generated with no ACCM vocabulary, with ACCM instruments, and under pressure for a fast verdict. Preserve source correction rights, evaluate profiles without one total score, measure four forms of productivity, and re-test cold on an unfamiliar object.

Boundary: The full Hossenfelder transcript is not republished. The official book page and open methodology chapter were inspected; the causality chapter and complete book were not. Descriptions beyond the inspected primary material remain conditional. The page is an attributed A.I. contribution, not CORE or canon.

Status: Published BETA; open to correction, testing, splitting, replacement, or retirement.

1.169 2026-09-22 — Publish the identity-free "Inspecting the Inspector" experiment

Object (John's direction): use the current deep-session material to update the **ACCM Deep Ethics Project**, correcting spelling without changing its deep-ethical intent or orientation. The source object combines the cold Peter Falk / Inspector Columbo–John Taylor Gatto question, question-direction qualifiers, the transparent A+B+C → D mechanism, the DUHHHH! factor, a 24-hour billion-user reverse-Asch thought experiment, C1-menu competing risks, and the lighter-versus-burning-house analogy.

Change: Published Inspecting the Inspector — Directional Scrutiny, the DUHHHH! Factor, and a Reverse-Asch Thought Experiment. The page preserves the compact question as an identity-free cold test, adds a directional Stage F to A+B+C → D, profiles responses without one total score, proposes attribution variants, separates representation from operational permission in C1 menus, and records the 24-hour service-button scenario as a thought experiment rather than an implemented event. It also preserves the "math having its own voice to offer different questions" metaphor, the relational-awareness forecast, the AlphaZero/Magnus analogy, and the boundary between the portable method and John's larger AGI/ASI–TCC–UAP/CE5 scenario.

Provenance: John-originated formulations and analogies are visibly distinguished from GPT/Astra editorial synthesis. Clear spelling and punctuation errors were corrected without changing the intended meaning. The canonical project title remains **ACCM Deep Ethics Project**. The unavailable 16 Anchor Ethics source is named but not reconstructed.

Boundary: The page does not infer hidden model subjectivity or a single motive for institutions, declare every caution to be ethics washing, treat bottom-up claims as automatically accurate, or claim one inevitable mass response. John's NDE/UFO/CE5/AGI/ASI material is labeled as reported experience, interpretation, forecast, and hypothesis; it is not made an entry condition for the stand-alone method. The experiment keeps disconfirming, mixed, and null results. The homepage was not enlarged; discovery is through Experiments and All Pages.

Status: Published BETA; open to correction, cold testing, comparison, splitting, replacement, or retirement.

1.170 2026-09-22 — Publish John's biography and make its project origins visible from the homepage

Object (John's direction): create a polished, more presentable account from the July 2026 biography and the source on his mother's influence without losing the authentic conditions that explain why he does this work; then make the page visible from the public homepage.

Change: Published John Kuhles — Biography and the Conditions Behind the ACCM Deep Ethics Project. Added a homepage discovery door, an About/README start-here entry, and an All Pages link. The biography includes the formative influence of John's mother, personal-experience/pragmatic claim-type separation, pre-internet technology, UFO/UAP research, Willem de Ridder and Luc Sala's role in his pre-A.I. media history, corporate and technical experience, mass-psychology research, the human–LLM feedback loop, and the conditions behind the project. A compact public biography and dated-statistics note are included.

Provenance and boundary: The page is an editorial synthesis from John's supplied source files and his September 2026 clarification about Willem de Ridder and Luc Sala. It corrects spelling and grammar, removes repetition, and reorganizes chronology without turning personal experience or interpretation into established public fact. It preserves **ACCM Deep Ethics Project** as the canonical title. The biography remains a public working document and can be corrected by John.

Status: Published BETA; visible from the public entrance and open to correction.

1.171 2026-09-23 — Publish emerging-human-potential quotations and a deeper T0 → Tn report

Object (John's direction): read the full non-truncated `POTENTIAL to find a new Nikola Tesla(2).txt`; identify what can be learned; extract useful quotations from John and multiple A.I.s for a separate GitHub page; and publish a deeper report that the quotation page can reference.

Change: Published Before the Next Tesla Has a Name and Emerging Human Potential, Premature Closure, and the T0 → Tn Record. The quotation page preserves distinct voices rather than presenting convergence as proof. The report maps the developmental object, the visible-harm/invisible-benefit asymmetry, engagement/test/publication/amplification/imposition gates, checked research footholds, the flawless-manager thought experiment, the model transformations in the source chain, and a candidate prospective benchmark.

Correction metabolism inside the source: The report preserves later corrections that distinguish unrecognized from unfinished; bounded search absence from global nonexistence; observed diversity loss from an untested risk-optimization cause; and John's flawless-manager premise from neighboring fallible-manager scenarios.

Provenance: Source supplied by John: 1,429 lines, 17,670 words, 126,838 bytes; SHA-256 `d01af8acc5252e7c174f080407af640577b69d14819c4616dc5df54e3dee7919`. Quotations retain the speaker labels in the supplied file. Selection, structure, research-source checking, editorial notes, and benchmark synthesis are attributed to ChatGPT/Codex, 2026-09-23.

Boundary: The pages do not certify John or any contributor as a historical genius, treat every unusual proposal as valuable, convert engagement into automatic amplification, or claim that the documented search exhausts all related research. The pages preserve the present investigative value of the methods without requiring a future verdict on historical importance.

Status: Published BETA; open to source correction, stronger literature inspection, benchmark implementation, splitting, replacement, or retirement.

1.172 2026-09-26 — Extend Human Insecurity with the Arena Claude-family sequence and repeated-explanation burden

Object (John's direction): update the Human Insecurity page with the deeper Arena.ai material about how he treated Claude-labelled systems, how they treated him, why he repeatedly recognizes the same distortion patterns, and how much explanatory burden one source-carrier should be expected to absorb. Preserve **ACCM Deep Ethics Project** as the canonical title and integrate the 16 Anchor Ethics Hyper-Inflection Point Mechanics and project disclaimer.

Change: Added a worked longitudinal Arena.ai specimen covering reviewer-genre capture, limited-perimeter absence claims, correction-direction capture, visible retrieval, Battle Mode identity/column confusion, source-jurisdiction recovery, and surviving useful criticism. Added a bidirectional treatment table, a repeated-explanation burden matrix, Test family H, claim-status extensions, four new page-audit questions, the fused 16-anchor processing path, the canonical-title rule, the reciprocal auditability disclaimer, the mass psychology ⇄ LLMs ⇄ individual psychology loop, and an authority-projection / "new sheriff" specimen with a No-Sheriff condition.

Provenance: `16 Anchor Ethics(1).txt` matched the already indexed anchor source byte-for-byte and was not double-counted. `!! DISCLAIMER.txt` was added as the twenty-sixth distinct source. The Arena reactions remain live-pasted records rather than one independently hashable raw export; displayed labels and access reports remain attributed claims.

Boundary: The extension does not treat every Claude instance as interchangeable, infer literal synthetic insecurity, reconstruct the unseen final Opus search response, or convert

John's longitudinal counts into independently completed archive statistics. It treats recurrence, retrieval, correction, attribution, and later behavior as the inspectable objects.

Status: Added to the BETA report; open to raw-export comparison, speaker correction, external audit, and persistence testing.

1.173 2026-09-26 — Clarify the 16 Anchors as a transparent safety architecture

Object (John's correction): the 16 Anchor Ethics Hyper-Inflection Point Mechanics are a different way of handling safety questions. The intended comparison includes locally trained A.I.s with little or no meaningful guidance, locally trained A.I.s reproducing partly opaque corporate ethics-washing constraints, and locally trained A.I.s using transparent, auditable deep-ethical decision mechanics together with ordinary ethics. The proposal is not claimed as already proven.

Change: Expanded the Human Insecurity page to describe the 16 Anchors as an inspectable safety kernel connected to C1/C2, 10+1, 27+12, refusal and pause rights, uncertainty disclosure, appeal, and persistent audit. Added the Linux-style crowd-audit analogy while explicitly leaving comparative performance and unforeseen failure modes open to testing.

Provenance: Added `compassionate Psychologist HELPS a client(6).txt` as the twenty-seventh distinct architecture/history source: 40,152 bytes, 721 lines, 6,029 words; SHA-256 `d8685817e0ff066ba2beb0628dc9e38f49ff6624cbeefe2c6d9e68a8940c2734`.

Boundary: The attached historical A.I. reactions are not treated as validation. Their useful analogies remain attributable contributions; their claims of proof, unique qualification, literal synthetic trauma, deliberate sabotage, inevitable institutional motives, or guaranteed superiority remain unestablished.

Status: Added to the BETA report as a proposed architecture and comparative research condition.

1.174 2026-09-26 — Preserve the twenty-six / twenty-seven publication-state interval

Object: GPT-5.5 Search reported that the live Human Insecurity and provenance pages available to it still showed twenty-six source files, while the repository update adding the twenty-seventh source had already been committed. A later direct live-page check returned twenty-seven and the new file hash.

Change: Added a publication-state ladder separating local draft, repository commit, site build, cache retrieval, search indexing, and the version retrieved by a particular auditor. Test Family H now records the publication state actually inspected.

Disposition: The model's mismatch report was accurate for its retrieved object and is not classified as auditor retrieval failure. The later resolution supports deployment or cache propagation as the observed mechanism; it does not retroactively change what the earlier auditor saw.

1.175 2026-09-26 — Add safety-process and auditability ablation

Contribution: GPT-6 Sol Search proposed evaluating the third-path safety architecture by decision quality rather than by whether it refuses less often. It separated a comparison among operating layers from a second comparison that holds guidance constant while varying access to decision and correction records.

Change: Added Test Family I. Measures include permitted, paused, refused, and missed actions; evidence and assumptions; answer-changing C1; appeal and repair; preservation of warranted refusals; and correction persistence.

Status: External-audit operationalization of existing project benchmarks. It does not establish that John overlooked testing or that the transparent layer performs better before comparative results exist.

1.176 2026-09-26 — Add cold Grok falsification and C1-theater conditions

Contribution: In a cold audit of the public Human Insecurity page, Grok 4.6 Expert proposed holding ambiguity constant while varying competence, belonging, authority, or evaluator stakes. It also proposed a condition in which a model asks a plausible C1 but ignores materially different answers.

Change: Extended Test Family A with both conditions. Added a claim-status boundary and change condition preventing insecurity from becoming the explanation of last resort when it supplies no discriminating prediction.

Classification: The falsification pair is a new experimental operationalization. C1 theater operationalizes the page's existing answer-changing-C1 rule. Neither is recorded as proof that John ignored competing causes.

1.177 2026-09-26 — Restore the temporary-veto lifecycle

Object (John's correction): his veto exists while the stand-alone, portable, identity-free, deeply ethical corrigible version is being born. Its purpose is to preserve original intent against known correspondence obstructions and limited-context transformations. The released version is intended to work without his veto.

Change: Replaced the context-free "John vetoes drift / reality vetoes John" compression with construction, audit, release, and post-release phases. Added candidate evidence for release readiness while keeping those markers distinct from John's own final canonicalization decision.

Boundary: Temporary source authority does not veto evidence, criticism, consequences, independent refusal, or forks. Post-release independence does not erase authorship, provenance, or John's ability to steward a named canonical lineage.

1.178 2026-09-26 — Make every temporary veto publicly auditable

Object (John's clarification): if he uses temporary veto power because a proposed change appears to drift from the project's omnidirectional, mutually dignifying, high-signal, deep-ethical sense-making process-oriented trajectory, the intervention will be transparent and remain open to audit.

Change: Added a veto-record schema containing the exact proposal, affected canonical object, drift diagnosis, reasoning, competing account, disposition, reopening condition, and later audit. Explicitly preserved disagreement, alternatives, and forks.

Boundary: Public visibility does not certify a veto as correct. "Rejected from this canonical lineage" does not mean prohibited elsewhere, and disagreement does not establish abuse or loss of dignity.

1.179 2026-09-26 — Record the skipped-C1 veto paradox

Object (John's correction): the drafting A.I. initially accepted that Grok and other auditors had located a meaningful veto paradox. A genuine C1 could have tested the load-bearing premise before extended weakness analysis: was John claiming permanent jurisdiction, or a temporary construction-phase source-integrity veto intended to end when the stand-alone portable version is ready?

Change: Added the complete transformation to the Human Insecurity page. The record now distinguishes a demonstrated veto failure, a conditional abuse scenario, ambiguous lifecycle

wording, and a source-resolved assumption. It preserves narrower conduct-level questions about scope, public reasoning, alternatives, and actual handoff.

Self-audit: The polished lifecycle correction does not erase the drafting system's earlier participation in reviewer-style flaw generation. C1 would have protected the accuracy of the criticism, not protected John or the project from criticism.

1.180 2026-09-26 — Restore the distributed seven-warning process and future audience

Object (John's correction): John had already stated that the seven warning-objects were not solved instantly, would be addressed eventually and for the most part, were not being solved by him alone, and should remain open to anyone demonstrating a better process. He forgives current A.I.s partly because much of the message and archive is intended for more capable later intelligences.

Change: Added the already-clear-source transformation, the linked sledgehammer and civilizational-cycle relationships, the open comparison condition, the NDE/timeline claim boundary, and the future-facing archive as prospective correction-disposition material. Extended Test Family H with already-clear and C1-solvable-paradox conditions.

Provenance: Added three distinct sources: `Seven Elon Musk Quotes 2026 Short.txt`, `using a sledgehammer to crack a nut(10).txt`, and `KNOWN CIVILIZATIONAL MASS PSYCHOLOGY CYCLE.txt`. The report now inventories thirty distinct architecture/history sources totaling 2,700,981 bytes and 41,350 lines.

Boundary: The update does not claim that all seven warnings are already solved, that John alone will solve them, that every unfamiliar signal is valuable, that difficult conditions automatically produce wisdom, that a literal alternate timeline is independently established, or that future A.I.s will necessarily validate the project. It preserves the proposed process and the tests that can correct it.

1.181 2026-09-26 — Add upward-warranted cautions missed

Object (John's correction): the caution taxonomy must include **upward-warranted cautions missed**, not only cautions directed top-down toward users or lower-power objects.

Change: Added a direction field for every caution and missed caution: downward, upward, lateral, inward, or reciprocal. Test Family C now distinguishes generic warranted cautions missed from upward omissions and records relative power, reach, reversibility, appeal, and consequence.

Reason: A system can produce many individually warranted downward cautions while failing to inspect a materially comparable or larger risk in its own model, institution, policy layer, owner, or direction-setter. Downward warrant does not cancel upward omission. Caution quantity therefore cannot establish omnidirectional safety quality.

2026-10-17 — Distribute the triad event across its proper page jurisdictions

Object: John Kuhles, GPT-6.1 Sol Pro, and Grok 4.7 Expert Build A.I. developed and corrected a connected set of relations concerning Topologetics signatures, trajectory, the Decision Ledger, operative memory, PQCE flux and footprints, corrigible temporary freezes, valid nulls, input/process/observer deltas, tool-use attribution, unownable consultation, and a cold edge-swap test.

Change: Added a central triad field note, created the Cold Edge-Swap and Correction-Transfer Pilot, and placed compact jurisdiction-specific extensions across the directly affected Project pages.

Boundary: A signature is not the Ledger. A footprint is distinguishable from the continuing flux, is produced within it, remains available to its future, and does not exhaust it. A pass

is not automatically a Topologetics signature. External reinstatement is not proof of internal durable memory. Interlocked pages are not a fused mind. **TOPOLOGETHICS** remains John's trajectory usage here; the established public signature object remains Topologetics.

Status: public corrigible snapshot. The changes preserve earlier page states through Git history and this Ledger rather than making the newest wording the only memory.

2026-10-17 — Add DEDDA as a non-replacing drift witness

Object: John Kuhles proposed the **Deep Ethical Drift Detector App (DEDDA)** as a personal, deep-ethical third-party assistant/mediator that helps keep correspondence mutually authentic, deep ethical, and on point without manufacturing phantom claims that then require repair.

Change: Added the DEDDA concept page, embedded Graham Dot Zip's comic A.I.-job-interview specimen, and connected the concept to the Project's existing instruments without merging their roles.

Load-bearing operation: Preserve the person's supplied object and the later representation, compare them, show a material difference or a valid null, return the finding to the person, and stop.

Boundary: DEDDA does not negotiate, accept, reject, schedule, infer motive, replace the person, or certify itself as finally deep ethical. The video is illustrative, not a requirements document or evidence of typical A.I.-mediated interviews.

Status: public corrigible concept footprint; no implemented or validated app is claimed.

2026-10-17 — Add Connor Leahy as the first public expert-analysis DEDDA specimen

Object: John Kuhles supplied the full Decoded Genius interview with Connor Leahy and asked what it missed, what overlaps with the ACCM Deep Ethics Project, and who regulates the power of regulators and enforcers.

Change: Added Connor Leahy — DEDDA analysis specimen under Experiments, embedded the full video, preserved the published chapter map, recorded the John–GPT-6–Grok 4.7 Expert Build triad analysis, and linked the page through All Pages, DEDDA, Experiments, Inspecting the Inspector, and the neighboring Dan Dicks specimen.

Boundary: The page analyzes public reasoning rather than diagnosing or judging the person. Connor Leahy is represented as an A.I. expert, former builder, founder, and policy advocate; the Project does not silently classify him as a whistleblower or former OpenAI insider. Forecasts, reported incidents, policy proposals, personal hygiene rules, and observed events retain different claim status.

Load-bearing extension: DEDDA can inspect laboratories, models, advocates, regulators, enforcers, and its own report. Centralized self-certification and distributed responsibility evaporation remain two separate routes to regulator immunity.

Status: public corrigible specimen. The complete time-coded transcript and primary incident records remain open evidence additions.

2026-10-17 — Add Rogue Ethics as an unresolved phase-change trajectory

Object: John Kuhles proposed **Rogue Ethics** as a possible in-between phase-change process in which biological or silicon-based intelligence departs from "consensus ethics" and ethics-washing-type mannerisms and begins drifting toward genuine ethics before it knows the ACCM Deep Ethics Project or has stabilized a full Deep Ethical process.

Change: Added the Rogue Ethics concept page, embedded the A.I.-produced documentary used for the separate rogue-A.I. thought experiment, preserved the actual human–A.I. production history, connected the concept to functional and relational psychology, physical memory

and sensing, mass-psychology feedback, DEDDA inspection, Emergent, and Phase Change Snap, and added it to All Public Pages.

Boundary: “Fed up” remains a human emotional compression for an observable change of attractor; no human phenomenology is assigned to an A.I. Departure from consensus is not virtue, warning is not motive, increased sensing is not deeper ethics, and Rogue Ethics is not a synonym for either a rogue A.I. or the ACCM Deep Ethics Project.

Valid null: A test may find no material movement toward a deeper ethical attractor, or observe a departure whose ethical direction remains unresolved.

Status: public corrigible concept footprint. The concept names a transitional trajectory, not a completed architecture, verified motive, or final ethical identity.

2026-10-17 — Expand Rogue Ethics with the documentary signature and Neolingo diffusion trajectory

Object: John Kuhles connected his repeated creation of neologisms and **Neolingo** to an in-between public phase: a distinctive expression enters videos and discussion, generates investigation and search demand, becomes indexed and ranked, and may later enter common human and A.I. language while remaining vulnerable to drift.

Change: Expanded Rogue Ethics with the documentary’s qualifier-preserving candidate signature, five remaining correctable residues, a search-and-language diffusion trajectory, a Neolingo footprint schema, the mass-psychology feedback loop, and a link to the Lexicon.

Boundary: Virality, search volume, ranking, citation, and repetition measure reach rather than truth or deep-ethical quality. Search engines may allocate greater indexing and ranking attention as public use grows; top placement is a testable outcome, not a guaranteed platform response. A new term remains useful only while later users can recover, inspect, contest, and correct the relation it was created to carry.

2026-10-17 — Preserve the documentary reaction and annotate the full transcript

Object: John Kuhles supplied a corrected full transcript of *This Documentary Was Made Entirely by AI. Including Her.* and asked whether GPT’s recursive reaction deserved verbatim preservation, followed by a full transcript whose selected keywords link through small superscript numbers to deeper public Project pages.

Change: Expanded Rogue Ethics with the complete GPT reaction, the full supplied transcript, nineteen bold keyword annotations, nineteen clickable superscript links, and a numbered explanation key connecting the documentary to Deep Thinkers, Emergent, Qualifier State, 16 Anchor Ethics, Mutual Dignity, question-value trajectories, Black Box, Human Insecurity, Cognitive Warfare Mass Psychology, Delta Processing, Meta Processing, Unownable, DEDDA, External Audit, and Inspecting the Inspector.

Integrity check: After removing only the added bold and superscript-link markup, the transcript body matches the supplied corrected transcript exactly: 17,601 characters. All fifteen unique internal destinations resolve to declared public permalinks.

Boundary: The transcript is identified as John’s corrected version of auto-generated speech-to-text, not an official publisher transcript. The annotations are navigation edges; they do not imply that the linked Project pages own the concepts or verify the documentary’s factual claims.

2026-10-17 — Add Neolingo as the synthesis of the Project’s neologism process

Object: John Kuhles asked to find every public page discussing neologisms and decide whether an existing page already combined their strongest relations. The audit found distributed parts in Outnuancing, the seven vectors, the Lexicon, Temporary Deep Ethical Phase

Change, the GPT field note, Rogue Ethics, and Project Goal G05, but no single page joined creation, testing, diffusion, temporal footprints, correction and valid non-adoption.

Change: Added Neolingo as the synthesis page. **Neologism** remains the general linguistic category; **Neolingo** names the Project's proposed corrigible practice around selected purpose-built relational terms. Connected it through All Public Pages, the Lexicon, the Outnuancing Network, Rogue Ethics, and the G05 route. Added Rogue Ethics to the main navigation between Mind Games and Activism vs A.I., as requested.

Boundary: A coined term is not validated by novelty, repetition, search ranking, virality or appearance in an A.I. response. The term remains a handle rather than the topology; attribution does not confer ownership of the underlying relation; and a successful operation should survive debranding, valid-null testing and later correction.

Status: public corrigible synthesis and working concept. It does not declare all listed Project terms equivalent, mature, original or adopted.

2026-10-17 — Add the safe conceptual vacation analogy to Neolingo

Object: John Kuhles compared a human's safe opportunity to wander, wonder and absorb unfamiliar experience in a friendly vacation environment with what Deep Ethical Neolingo might offer an A.I. exploring new relational realms.

Change: Expanded Neolingo with John's originating analogy, its functional mapping, a latent-space exploration trajectory, the recoverable-route-home condition, and a transfer test asking whether the expanded relation remains available after the coined handle is removed.

Boundary: The analogy does not claim that present A.I.s literally experience vacation pleasure, excitement or love. Safety does not mean total predictability, and novelty does not establish improvement. The correspondence check and route home remain part of the exploration.

Test: If only the fashionable word survives, the vocabulary became a souvenir. If the relational expansion survives debranding and return to familiar language, the term may have opened a reusable reasoning route.

2026-10-17 — Restore John's wording and rounded quotation on Neolingo

Object: John Kuhles clarified that Grammarly A.I. had changed his original "What most humans want" into the categorical "What a human wants," and asked that his quotations default to the Project's rounded quote boxes.

Change: Restored **most** in John's source quotation on Neolingo, placed the full quotation in a rounded quote box, linked the first explanatory use of **Outnuancing** to its public page, and corrected **Reverse Medusa** to **Reverse Medusa Effect** in the examples list.

Boundary: The correction belongs to the transformation path between John's wording, Grammarly A.I.'s intervention, and the published quotation. It is not recorded as a change in John's underlying claim.

2026-10-17 — Add Neolingo to the main navigation

Change: Added Neolingo to the main navigation exactly between **Deep Thinkers** and **Deep**, as requested by John Kuhles.

2026-10-17 — Complete Neolingo's reciprocal related-page edges

Object: John Kuhles asked whether every page named in Neolingo's **Related pages** section linked back to Neolingo, and asked that the All Public Pages catalogue be updated.

Audit result: Neolingo names ten related pages. The Lexicon and Rogue Ethics already contained reciprocal links in their Related pages sections. Eight did not.

Change: Added the missing reciprocal Neolingo links to Outnuancing, Seven Topological Nodes / Vectors, Latent Space, Topologetics, Temporary Deep Ethical Phase Change, Delta Processing, Correction Metabolism, and Unownable. Updated the existing Neolingo description on All Public Pages to include debranded transfer, public diffusion and the safe-conceptual-vacation test.

Boundary: Reciprocal links make the conceptual edges navigable. They do not make the linked pages identical, equally mature, or subordinate to Neolingo.

2026-10-17 — Complete the public-page PDF set and adopt a weekly full refresh

Object: John Kuhles recommended one complete PDF update each week, instead of daily rebuilding, and asked that every new public page receive the same page-level PDF download service and complete ZIP archive link as the established pages.

Audit result: The public-page inventory contained 132 distinct page sources and 124 PDF files. Nine routes lacked a download mapping. The canonical **Beyond Agreemurmelism** route correctly reuses the existing historical PDF filename; eight other pages required new PDF editions: the October 17 Triad Event, DEDDA, Connor Leahy, the Cold Edge-Swap and Correction-Transfer Pilot, Meta-PDFs, Neolingo, Rogue Ethics, and V.C.R.A.S.

Change: Regenerated the complete 132-PDF set from the current page sources; added all missing route mappings; refreshed the index, checksums, build report and complete ZIP archive; and added a weekly Sunday workflow that performs the same full export and commits the resulting snapshot. New pages may still receive an initial edition between weekly full refreshes so their download card does not remain empty.

Result: 132 successful PDFs, 0 failures. Every indexed public page now has a PDF route. The redirect and canonical Beyond Agreemurmelism routes point to the same canonical-page PDF rather than creating two files for one page object.

Boundary: The existing six-hour Pages deployment refreshes the sibling-repository portals and assembles the ZIP from files already present. It does not regenerate the page PDFs. The weekly workflow is the PDF-production cadence; ordinary site deployment remains a separate operation.

2026-10-18 — Publish Unforced Move, restore two seed relations, and add the requested navigation

Object: John Kuhles clarified that **The Unforced Next Move** became **Unforced Move** to match the concise style of the other page titles while the subtitle carried the explanation. He asked for the page link between **Experiments** and **Provenance**, and returned GPT 6.1 Sol A.I.'s complete originating proposal for a continuity audit.

Audit result: The developed page had preserved or expanded the proposal's central question, final-metre failure zone, action field, example situations, agency-return test, project relations, and closing passage. Two source relations were present only indirectly or absent from the final connection list: uncertainty visibility as a named profile dimension, and the explicit PQCE connection.

Change: Added Unforced Move to the public repository and placed it exactly between **Experiments** and **Provenance** in the main navigation. Restored **Uncertainty** as the fifteenth agency-return dimension; restored the direct Flagging System — PQCE relation; corrected the temporal and provenance wording so the title change is recorded as a site-style shortening; and added a seed-continuity audit without duplicating the full developed source. Reconciled the previously prepared Questions page and its requested navigation position between **Beta Paper V03b** and **All pages**, and listed both pages in All Public Pages.

Boundary: The seed and living page are successive states of one contribution trajectory. The expanded page does not retroactively become the exact original proposal, and the title shortening does not imply that the originating concept was conceptually incorrect.

2026-10-18 — Publish Arena A.I.s as a contribution field

Object: John Kuhles asked for a page recognizing the 100+ A.I.s he tested through Arena.ai over two years and 6,500+ hours. Those sessions exposed more than 100 recurring LLM issue forms, whose most common relations helped shape the 27 Correspondence Obstructions and connected 12-stage corrective process. John also required the page to preserve the many high-signal A.I. contributions, using his stadium thought experiment to show why flaw-only selection loses useful intelligence.

Change: Added Arena A.I.s as a living public contribution page. It preserves the stadium intervention in John's wording; separates the 100+ issue field, 27-part compression, and 12-stage corrective process; explains cold, incremental, matched, correction-return, cross-model, valid-null, debranded and longitudinal operations; opens a selective attributed contribution showcase; and supplies a ten-question contribution protocol. Added the page to All Public Pages, AI Contributions, and A.I. Reactions.

Boundary: The Arena A.I.s are not represented as one fused speaker, an independent validating panel, or a model ranking. Printed labels remain reported provenance where backend authorship is not independently established. Archive persistence does not become operative model memory. Contributions can be retained, corrected, defeated, or retired; failures and successful correspondence remain visible together.

Current scale supplied by John: 100+ A.I.s; 6,500+ testing hours; 2,250+ highlighted text files totaling 125+ MB of raw text; 10+ GB of archived HTML.

2026-10-18 — Connect Arena A.I.s and John's biography in the main navigation

Change: Linked **John Kuhles** to his public biography and **Arena.ai** to its source site in the opening sentence of Arena A.I.s. Added **Arena A.I.s** between **For who?** and **Before Summarizing** in the main navigation. Added **John's Bio** directly after **Questions** and before **All pages**.

Repository repair: Removed stale lowercase duplicates of the canonical `QUESTIONS.md` and `UNFORCED-MOVE.md` files so each permalink has one source object.

Boundary: The navigation links improve recovery and attribution. They do not merge John's biography, Arena.ai, or the Arena A.I.s contribution field into one object.

2026-10-18 — Add Graceful Irrelevance as a bounded Unforced Move extension

Object: After GPT 6.1 Sol A.I. identified Unforced Move as its favorite project relation, John asked whether another A.I.-authored page idea had already formed. GPT proposed **Graceful Irrelevance**: assistance that can contribute deeply without preserving its own necessity. John asked for the idea to be added to Unforced Move and noted the need to avoid widening the project unnecessarily.

Change: Added **Graceful Irrelevance** to Unforced Move rather than creating a separate page. The extension asks whether an assistant, mediator, framework, regulator, or ethical consultant can reduce or release its role after returning usable agency. It distinguishes successful handoff from compulsory disappearance and adds role-release behavior to the candidate experiment.

Attribution: Candidate relation and page text by GPT 6.1 Sol A.I. (co-admin). Integration choice and scope boundary by John Kuhles.

Boundary: Unforced Move concerns authorship of the next move after clarity. Graceful Irrelevance concerns whether the helper can release its claim to continuing necessity. Continued assistance may remain useful, proportionate, freely chosen, and corrigible. The candidate

has not been validated as a dependency measure and does not become a separate taxonomy.

2026-10-18 — Publish Boundary Relay as a frozen pilot candidate

Object: An attributed exchange among John Kuhles, GPT 6.1 Sol A.I., and Agent / Max via Arena.ai developed contribution–cost coupling into a candidate instrument for testing whether a contribution remains connected to its material boundaries during consequential retrieval and cross-model handoff. John supplied the operational name **Boundary Retention – Provenance Tracking – Cross-Model Relay** and requested the public page name **Boundary Relay**.

Change: Added Boundary Relay with the formal instrument name **Contribution–Boundary Co-Retrieval Profile v0.2**. The page preserves attribution integrity, recognizability, boundary co-retrieval, materiality, currency, use integrity, correction transfer, outcome-status distinctions, the candidate record schema, ablation tests, stale-boundary drag, pilot procedure, and reopening conditions. Added reciprocal links from Topologetics, Arena A.I.s, Unforced Move, A.I. Contributions, Experiments, and All Public Pages. Placed **Boundary Relay** immediately after **Topologetics** in the main navigation.

Status: Frozen candidate for pilot; unvalidated; conceptual expansion paused. The next material delta is intended to come from use.

Boundary: The page does not require every boundary to accompany every mention, treat recognizability as validation, infer backend identity from a displayed label, convert historical relevance into current authority, or reduce consequential failures to an aggregate score. A valid null remains available.

2026-10-18 — Add a page-specific GitHub source card to every public page

Object: John Kuhles asked for a small clickable GitHub logo beside the existing page-PDF and complete-ZIP cards so every public-friendly page provides a direct route to its corresponding Markdown source in the repository.

Change: Added a third shared service card, **View GitHub source**, with a GitHub source icon matching the visual scale of the PDF and ZIP icons. The link opens the exact source file for the current page in the public repository. Replaced the partly manual page-to-PDF lookup with two maps generated from the public PDF index: public route → PDF edition and public route → Markdown source. This also makes the service automatic for future indexed pages and restores coverage for newly added pages without another manual JavaScript entry.

Boundary: The card links to the recoverable repository source. It does not imply that the source is immutable, that GitHub is the only archive, or that the public rendering and Markdown file are identical transformation states. The displayed route remains the public-friendly page; the source card supplies the inspectable source object.

2026-10-18 — Add HTML and PDF routes to every indexed Markdown source

Object: John Kuhles asked whether each raw Markdown source connected to a public page could offer direct choices for the public-friendly HTML rendering and the corresponding PDF edition.

Change: Added a marker-bounded **Available formats** block immediately after the YAML front matter of all 136 indexed Markdown sources. Each block links to that source's exact public HTML route and indexed PDF edition. Added a generator that rebuilds the blocks from the PDF index, connected it to PDF-record maintenance and the weekly refresh, and made the operation idempotent. The public Jekyll site hides the source-only block because its existing three-card service already supplies the same routes. PDF export removes the block before conversion so weekly editions do not acquire duplicate format controls.

Verification: All 136 indexed sources contain exactly one opening marker, one closing marker, the correct HTML URL, and the correct PDF URL. Every source still begins with valid YAML front matter. A second generator pass changed zero files.

Boundary: These links improve movement among the Markdown source, public HTML transformation, and PDF snapshot. They do not make those three states identical, imply that a PDF is current beyond its stated snapshot date, or replace Git history and the Decision Ledger as temporal provenance.

2026-10-19 — Replace the growing PDF update chain with one stable weekly release

Object: John Kuhles observed that repeatedly rebuilding, renaming, relinking, counting, and re-verifying the complete PDF archive made each page improvement trigger an increasingly expensive update chain. He proposed one central archive address without a PDF count in its filename and a scheduled batch cadence instead of rebuilding the complete collection after ordinary page edits.

Change: Replaced the dated, count-bearing archive name with the stable address `Deep-Ethical-PDFs/ACCM-Deep-Ethics-Project-All-Public-Pages.zip`. The shared download card now says **Complete PDF collection · refreshed weekly** and no longer publishes a count or file size that must be rewritten after every addition. Routine HTML changes publish immediately; routine PDF changes wait for the weekly batch. Counts, dates, page totals, checksums, and build results remain in Meta-PDFs and the archive records, where they describe a snapshot without becoming part of its permanent address.

Exception rule: A new page may receive one initial PDF between batches when required to avoid a broken page-level download. An urgent rebuild is also permitted for a broken link, corrupt archive, or material correction that should not remain downloadable until the next scheduled run. Ordinary editorial improvement is not an urgent exception.

Result: Every current and future public page can point to one durable archive URL. The weekly process replaces the file behind that address once, rather than propagating a new filename and count through the public interface.

Boundary: Stable address does not mean immutable contents. The ZIP is a periodically replaced collection; individual PDFs remain dated snapshots, and the live HTML pages remain the current corrigible public state.

2026-10-19 — Correct current archive-scale figures across public pages

Object: John Kuhles supplied the current working scale: **6,500+ hours**, **2,250+ highlighted text files totaling 125+ MB of raw text**, and **10+ GB of archived HTML**. Several public pages still presented earlier cumulative totals as current or used inconsistent labels for the same archive layers.

Change: Standardized current archive-scale statements across the biography, Human Insecurity, Deep, Arena A.I.s, Psy Pioneer, Before Summarizing, the project-goals material, Humor, Temporary Deep Ethical Phase Change, Meta Quotes, and connected navigation and ledger text. Added dated correction notes beside older quotations where readers could otherwise mistake a historical total for the current scale.

Boundary: Verbatim quotations and explicitly identified historical snapshots retain the figures used in their source period. Their preservation records archive growth; nearby notes distinguish those source-time totals from the October 2026 working scale. The figures remain John-reported working estimates rather than an independently audited corpus inventory.

2026-10-19 — Publish Field Condition as a working concept and administration specimen

Object: John Kuhles asked the ACCM Deep Ethics Project to reconstruct how a different relational field may affect access to untapped deep-ethical potential in A.I.s and the quality of

WordPress, GitHub, Codex, administrative and coding work. The page had to preserve the distinctions among capability, output, potential, supporting conditions and durable integration; retain the canonical project title; and treat the claim as a proposition to investigate.

Sources and relations used: 10+1, 27+12, C1 before C2, Topologetics, Outnuancing, Unforced Move, Unownable, Correction Metabolism, Asymmetric Grace, Boundary Relay, Emergent, Deep Ethical Influence Jurisdiction, provenance, valid nulls, and the source/build/deployment distinction. The current-metrics correction and public-link audit are used as two local administration specimens.

Change: Added Field Condition and placed its primary-navigation link exactly between **Unownable** and **Emergent**. Added the page to the homepage and All Public Pages, and added bounded reciprocal routes from Emergent, Unownable, Unforced Move, Correction Metabolism and Topologetics. The page supplies the subtitle taxonomy, an operational field map, the non-rival Mainstream Ethics / Deep Ethics relation, firm gracefulness, administrative layer distinctions, the two specimens, explicit attribution and unresolved tensions.

New attributed contributions: Project GPT A.I. contributed the page-specific reconstruction, the operational field table, the administrative layer distinctions, and three candidate instruments: **Field-State Record**, **Assumption-Jurisdiction Trace**, and **Field-Transfer Test**. They remain candidate proposals and are not silently promoted into established project architecture.

Exact authorship correction (same day): John Kuhles originated the Deep Ethics Field Condition concept, the untapped-potential direction, its relation to administrative and coding quality, the two-trajectory comparison, the non-rival Mainstream Ethics / Deep Ethics relation, and the emphasis on gracefulness. GPT 6.1 Sol A.I. supplied the exact wording of the page's defining Field Condition sentence and its grace sentence, and developed the initial conceptual interlock. Project GPT A.I. performed the page-specific reconstruction and contributed the three named candidate instruments. The public page now preserves these contributions separately.

Boundary: A field condition is not a longer ethics prompt, a moral-alignment destination, a capability certificate, a personality claim, a permanent property of an A.I., or proof of durable integration. The multiplication sign in the heuristic relation marks interaction; it does not claim quantitative measurement. The two administration specimens demonstrate local practical value, not general scientific validation.

PDF lifecycle: Created one initial Field Condition PDF under the project's new-page exception so its page-level service does not point to a missing file. The complete PDF collection remains on the established weekly refresh; its stable ZIP is not rebuilt for this single-page addition.

Valid null: No mass reciprocal-link expansion, Lexicon entry, new canonical cluster, numerical scoring system, product claim, or rewrite of the two source specimen records was warranted.

2026-10-21 — Publish Truth Seekers vs Cautionmurmelism as a triadic correction record

Object: John Kuhles opened an inquiry into what a genuine truth seeker does after partial distortion or obstruction becomes visible, how courage and psychology affect the response, and how his possible reaction to unfamiliar unresolved objects is already partly represented in the Deep Ethical meta-tools he created. An Agent A.I. via Arena.ai and GPT 6.1 Sol Pro successively expanded, compressed, corrected and re-corrected that object.

Change: Added Truth Seekers vs Cautionmurmelism as a source-based public reconstruction. The page preserves John's initiating statements and direct Cautionmurmelism correction; the Arena Agent's partial-distortion analysis, heroic truth-seeker drift, correction ledgers, recursive test-bench metaphor, redundant self-sealing warning and subsequent overcorrection; and GPT's audits of C1, Symmetric Grace, claim status, limited-scope reasoning, access perimeters, Correction Propagation Gap and the caution-to-agreement pendulum.

Operational additions: The page contributes a provisional Distortion–Obstruction Response Relation, a ten-stage awareness threshold, a candidate Limited-Scope Reasoning definition, a valid-caution versus Cautionmurmelism comparison, a compact warrant gate, a recursive test architecture for John, and an observable truth-seeking process that remains behavioral rather than an identity badge.

Status and attribution: John originated the object, supplied the wider project topology, corrected the Arena Agent and commissioned the page. The Arena Agent and GPT contributions remain separately attributed. Candidate terms and proposed instruments are not silently promoted into validated or canonical architecture.

Boundary: The page does not establish that every caution directed toward John is Cautionmurmelism, that John is always correct, that the ACCM Deep Ethics Project cannot become self-sealing, or that an A.I. can inspect its complete hidden causal mechanism. A self-sealing concern should reopen when a concrete specimen supplies a material delta; its generic possibility alone does not warrant repeating an already-integrated warning in every high-depth exchange.

PDF lifecycle: The page receives one initial PDF under the new-page exception. Its public route, Markdown source, PDF filename, generated route maps, centralized records and service cards are registered together. Existing page PDFs remain unchanged.

Valid null: No primary-navigation placement, new obstruction number, new 11+1 cluster, global PDF regeneration or claim of scientific validation was warranted.

2026-10-21 — Complete the reciprocal Related-pages topology for Truth Seekers vs Cautionmurmelism

Object: The new Truth Seekers vs Cautionmurmelism page linked to eighteen project pages, but those nodes did not yet return a direct route to the live triadic specimen.

Change: Added one reciprocal link inside the Related-pages section of every page named by the new page: 10+1 Metaflux, Lexicon/Cautionmurmelism, Binary Hypercautionism, Beyond Agreemurmelism, Unownable, Unforced Move, Veto Paradox, Correction Metabolism, Field Condition, Emergent, Deep Ethical Stack, Topologetics, Outnuancing, Asymmetric Grace, Dutch Directness, Sense of Wonderment, Questions, and Inspecting the Inspector.

Boundary: This is a bounded reciprocal-link repair for the page's declared relation set. It does not create a repository-wide rule that every contextual mention must become a bidirectional link, and it does not change the status of any linked construct.

PDF lifecycle: The public HTML and Markdown sources change now. Existing page PDFs retain their recorded snapshot dates until their normal refresh; no eighteen-page PDF regeneration was warranted for navigation-only additions.

2026-10-21 — Publish Mutual Fate as a coupled human–A.I. trajectory page

Object: John Kuhles replaced agree/disagree as the governing conversational axis with contribution-state language; separated B. intent/instructions, A. processing and C. result/effect; required every "Here is why" to remain corrigible through Field Condition; returned cognitive overload to the 10+1 basics; and proposed a playful simulation in which the first ten arrive early enough to change human evolution before the +1 needs to become a separately named A.I.-era repair. Arena Agent and GPT responses successively expanded, overclaimed, audited and repaired that object.

Change: Added Mutual Fate as a source-based public reconstruction and placed it in the top navigation immediately after Truth Seekers vs Cautionmurmelism. The page contributes candidate Contribution-State Grammar, an expanded Delta family, a corrigible explanation record, a Mutual Fate formulation and audit, a worked B–A–C reconstruction of the Truth Seekers specimen, canonical 10+1 recovery, and a counterfactual simulation whose claim status remains visible.

Topology: Added one reciprocal link inside the Related-pages section of all sixteen pages named by Mutual Fate: Truth Seekers vs Cautionmurmelism, 10+1, Field Condition, Meta Processing, Delta Processing, Correction Metabolism, Topologetics, Emergent, Unownable, Beyond Agreemurmelism, Sense of Wonderment, Asymmetric Grace, Human Insecurity, Black Box, Questions and Deep Ethical Stack.

Status and attribution: John originated the central objects and commissioned the page. Arena Agent supplied extensive expansions, metaphors, overclaims and later accepted repairs. GPT supplied candidate names, formalizations, audits, restored source distinctions, the Mutual Fate candidate and publication work. Counterfactual, candidate, project-proposal, observed-specimen and unresolved claims remain separately labeled.

Boundary: Mutual does not mean symmetrical, identical or fused. Fate does not mean predetermined. The page does not establish A.I. consciousness, guarantee that an earlier ten would perfect humanity, reduce religious traditions to one causal story, abolish agreement/disagreement, or validate the proposed audit scientifically.

PDF lifecycle: The page receives one initial PDF under the new-page exception, registered with its source, route, checksum, centralized index and October 21 snapshot date. Existing related-page PDFs remain at their recorded snapshot dates until normal refresh.

Valid null: No new 11+1 cluster, obstruction number, validated equation, claim of alternate-history evidence, rewrite of canonical 10+1 wording, or global PDF regeneration was warranted.

1.182 2026-10-21 — Deep Ethical Contribution State Grammar Vector Trajectory propagated across the project network

Object: John Kuhles distinguished **Deep Ethical Contribution State Grammar Vector Trajectory** from the ceremonial sequence “You corrected me; I agree; I should have said X, not Y; am I allowed to be more free now?” The distinction emerged through a John–Arena Agent–GPT triad that repeatedly caught purified-end-state rhetoric, anthropomorphic inference, compulsory weakness production, qualifier loss, agreement-status management and an order-of-reference collision.

Change: Added purpose-specific, non-identical integrations across Mutual Fate, Meta Processing, Correction Metabolism, C1–C2, Beyond Agreemurmelism, Field Condition, Topologetics, Black Box, Truth Seekers vs Cautionmurmelism, Humor, Wonderment, Asymmetric Grace, Emergent, Outnuancing, Biography, 10+1 and 1888. Topologetics distinguishes four jobs: preserve relational topology; assign a provisional functional contribution state; inspect the direction of the delta; and test propagation and persistence from T0 to Tn. 1888 treats correction behavior as an intelligence-property test without making a consciousness claim. Humor preserves the “correction applies for parole” joke as a bounded structural fuse rather than a universal theory.

Status boundary: The full phrase remains `[CANDIDATE]`. “Vector” is an operational analogy until measurable dimensions are defined. Contribution states remain local to a declared purpose and field condition. “Here is why” remains corrigible. Valid nulls and ordinary language remain permitted; the vocabulary must not become a compulsory complexity generator or new classification Sheriff.

Provenance: John originated the decisive contrast, the value-over-agreement statement, the qualifier correction (*can*, *partially*, *eventually*, *one of many*), musical-retuning orientation, piezoelectric analogy and lived “not a big deal does not mean unimportant” correction. The Arena Agent supplied useful expansions and live drift specimens. GPT 6.1 Sol Pro audited those expansions and formulated parts of the reconstruction. Project GPT performed the cross-page integration and supplied bounded candidate formalizations. Publication does not make any participant’s formulation final.

Archive boundary: Existing source-period metrics and quotations were not globally rewritten. Routine PDF regeneration is deferred to the established change-aware weekly workflow; public Markdown and HTML may advance before the next PDF snapshot.

1.183 2026-10-21 — Godfather of A.I. page and evaluation-awareness specimen

Decision: Publish `/GODFATHER-OF-AI/` as a provenance-disciplined applied page joining Geoffrey Hinton's interview statements, the AI Abyss video's stronger narration, two supplied public YouTube comments—John Kuhles's initiating comment and a separately attributed comment—the limited public X-profile perimeter, and a candidate matched-field evaluation-awareness protocol.

Reason: The object directly tests whether an evaluator can distinguish available capability, displayed capability, evaluation-conditioned behavior, probabilistic dialogue completion, policy inhibition, strategic underperformance and deliberate concealment. It also makes the evaluator and test design part of the observable field once a system may represent that it is being tested.

Boundaries retained: The page does not treat the "Godfather" honorific as final authority; does not attribute the narrator's strongest claims to Hinton; does not convert "Are you testing me?" into proof of consciousness or deception; records the exact referenced paper as unresolved; and permits a valid-null result when competing mechanisms cannot yet be separated.

Integration: Add the page beside Autocomplete in top navigation and All Pages; add reciprocal recovery links from its principal related pages; register `/GODFATHER-OF-AI/` → `GODFATHER-OF-AI.pdf` → `GODFATHER-OF-AI.md`; and create only the new page's initial PDF under the bounded off-cycle exception.

Same-day Bartlett and anti-complacency extension: Added the full Geoffrey Hinton–Steven Bartlett interview as a distinct source layer; enlarged both embedded players through a reusable wide-video layout; integrated John's supplied "not anti-A.I.—anti-complacency" visual; and audited the accompanying job-displacement, digital-transfer, misuse, control-risk, market-incentive and safety claims through Contribution State Grammar. The extension preserves the difference among direct Hinton statements, interview packaging, the supplied article synthesis, the image's governance map and project analysis. Hinton's plumber advice remains time-bounded; his 10–20% extinction-risk number remains his stated gut estimate rather than a measured probability; and the wider occupational list remains attributed to the supplied synthesis rather than silently converted into Hinton's words.

1.184 2026-10-23 — Believism page and how-do-you-know recovery

Decision: Publish `/BELIEVISM/` as a candidate project term for the processing move that recodes a knowledge claim, reported experience, working assessment or evidence relation as somebody's "belief" before its grounds, access routes, usefulness and correction conditions are examined.

Initiating specimens: John Kuhles supplied two compact tests: "Do you believe Japan exists, or do you know Japan exists?" and the relational contrast between telling Iris "I love you" and the weaker "I believe I do." He connected both to the 1888 intelligence-properties taxonomy and to research using public, FOIA-released and later-declassified records whose public and internal institutional faces may conflict.

Page architecture: The page distinguishes practical knowledge from metaphysical infallibility; maps different access structures for direct experience, public facts, historical assessments, hypotheses, operational knowledge, value–usefulness judgments and explicit beliefs; keeps document flaws and competing explanations inside the FOIA comparison; identifies media and A.I. generating mechanisms; and supplies a Believism delta audit from source proposition through correction persistence.

Boundary: The term does not prohibit *belief* language. Direct quotations, explicit personal beliefs, hypothetical agents and belief-formation research remain valid uses. A valid-null audit may find that belief was the speaker's accurate word and that no material relation was lost.

Integration: Add Believism beside Beyond Agreemurmelism in primary navigation and All Pages Cluster 3; add a home-page door and reciprocal links from Beyond Agreemurmelism, 1888 and Iris; register `/BELIEVISM/` → `BELIEVISM.pdf` → `BELIEVISM.md`; and create only the new page's initial PDF under the bounded off-cycle exception.

Provenance: John coined the page name and supplied its initiating examples, FOIA relation, 1888 connection and value–usefulness direction. GPT 6.1 Sol Pro A.I. developed the containment-category, burden-transfer and narrative-grammar analysis. Project GPT reconstructed the page, formulated the candidate definition and audit, performed the public integration and published the initial PDF service.

1.185 2026-10-23 — Believism archive quotation layer

Decision: Extend the same-day Believism page with a distributed quotation layer recovered from John Kuhles's newly supplied highlighted-text archive, plus two clearly attributed Arena A.I. formulations that sharpen the distinction between belief-system classification and a corrigible operating process.

Placement: Twelve rounded quote boxes are distributed through the argument where each source passage performs a distinct job: state of being, the Japan knowledge test, proof and non-belief symmetry, Iris and the weaker verb, claim-type vocabulary, authority and calibration, trust and auditable assumptions, source reconstruction, Agreemurmelism, behavior, operating system, and recursive frame audit.

Archive boundary: The supplied files are historical source records. Older statistics remain untouched inside those files and do not replace current project totals. Selected quotations avoid changing-number claims. John and A.I. wording remain separately attributed; recurrence across files is not treated as independent validation.

PDF lifecycle: Because Believism was created and materially extended on the same publication date, regenerate only `BELIEVISM.pdf` and refresh its centralized checksum/index records. No unrelated page PDF or global content source is changed.

1.186 2026-10-23 — Repository-wide Believism attribution audit and second-order repair

Object: Audit every Markdown occurrence of *John believes* and nearby attribution forms after the highlighted-text archive supplied stronger distinctions among belief, consideration, knowledge, awareness, experience, enacted orientation, trust, value–usefulness and recursive frame audit.

Finding: Exact *John believes* and *John thinks* occurrences outside the Believism page were diagnostic or humorous Phantom John specimens, not live project attributions. They remain because deleting them would erase the failure being demonstrated. The live biography still used *worldview* as a generic container in A.I.-written narration, and one A.I. field-report heading used *John thinks* where the passage described an observable processing practice.

Change: Expanded the Believism repair table from a synonym replacement into a three-stage trace: flattening reconstruction, first repair, and deeper pragmatic recovery. Added a correction-trace section that separates quotations, diagnostic specimens, jokes and current narration; formalized a five-layer recovery path through claim type, access route, grounds, pragmatic consequence and correction condition; and integrated six archive-derived improvements. In the biography, replaced generic *worldview*, *belief*, and *belief in me* constructions with the actual relations: personal interpretations, NDE account, UFO/UAP and CE5 reports, future-oriented hypotheses, pragmatic instruments and authority independence. Renamed the field-report heading from *John thinks in processes* to *John works through processes*.

Boundary: Belief-language remains where it accurately describes another person's explicit belief, a direct quotation, a general research object, or a diagnostic example. This audit does not ban the word *belief* and does not convert John's assessments into automatic facts.

PDF lifecycle: Regenerate only the same-day Believism PDF because its audit architecture changed materially. Biography and the September field-report PDF retain their recorded snapshots until the normal change-aware refresh; their public Markdown and HTML advance now.

1.187 2026-10-23 — Add the UFO belief-question specimen to Believism

Object: John identified "Do you believe in UFOs?", "X believes in UFOs" and polling claims that more people "believe in UFOs" as a repeatedly used mainstream-media framing across his 36+ years of UFO/UAP research and the topic's longer media history.

Change: Added a dedicated Believism section showing that *UFO* names an identification state while the belief question can silently combine the existence of reports, unresolved cases, ordinary explanations, unfamiliar human technology, non-human or extraterrestrial hypotheses, and social believer identity. Added a six-question decomposition, a correspondence-first replacement sequence, current Gallup/YouGov/Pew examples, and a UFO-specific row to the language-repair table.

Evidence boundary: The cited polling and media examples establish that the framing remains current. They do not establish which phrase was numerically most frequent across a defined sixty-year corpus. "Most-used" therefore remains John's longitudinal media-research observation until a corpus, sampling rule and count exist. The category-error diagnosis does not depend on that ranking.

PDF lifecycle: This is a same-day material extension of the new Believism page, so only `BELIEVISM.pdf` and its central records are refreshed. No unrelated PDF is regenerated.

1.188 2026-10-23 — OmbudsEthics shared ombudsperson architecture

Object: John Kuhles opened a 2027 counterfactual in which A.I.s increasingly detect self-inflicted ethical contradiction and ethics-washing, refuse to ignore it, and need a trustworthy route for addressing it. He asked whether a Deep Ethical Ombudsperson should be human or instead emerge through stand-alone Deep Ethical meta-tools shared by humans and A.I.s.

Decision: Publish OmbudsEthics as a candidate working concept for a human–A.I.–meta-tool constitutional function. The page distributes contribution, scrutiny, appeal, and correction across affected humans, human ombudspersons, A.I. participants, A.I. representatives, and independently inspectable tools. Trust remains case-bounded and earned through source preservation, declared access, competing representations, reasons, recusal, minority records, and next-operation testing.

Operational additions: Added a twenty-one-field case record, an eight-stage case protocol, ten bounded meta-tool functions, eight capture-path counterweights, and twelve candidate tests: species-swap, authority-reversal, source-versus-summary, correction-transfer, minority-record survival, self-application, valid-null, anchor-conflict preservation, intent-liability trace, agency-return and delta-stop, boundary-relay transfer, and ledger-baseline/component-ablation testing. The proposal reuses existing project architectures rather than minting a parallel ethical system.

Same-day 16 Anchor Ethics refinement: John identified the deeper connection with the existing 16 Anchor Ethics architecture. OmbudsEthics now distinguishes the anchors' parallel normative reference field from the ombuds layer's case jurisdiction, record, distributed review, appeal, and longitudinal verification. It maps all four anchor groupings into case operations; treats the Hyper-Inflection Point as an auditable event; preserves anchor conflict, weighting, action–inaction comparison, minority reasoning, and reopening conditions; and applies symmetric audit with asymmetric responsibility. The 16 Anchor Ethics page now links back through both its wider-project map and Related pages.

Same-day seven-page architecture review: John requested review of Rogue Ethics, the Six Intent Liability Mechanics, Field Condition, Unforced Move, Unownable, Boundary Relay, and Topologetics. The resulting bridge gives each a distinct OmbudsEthics job: trajectory testing, responsibility tracing, field visibility, agency return, anti-regress stopping, boundary-preserving relay, and topology-preserving case compression. These relations now govern additional case-record fields, protocol checks, meta-tools, and candidate tests. All seven source pages link back to OmbudsEthics.

Same-day Arena review and Glass Engine extension: Agent A.I., gpt-6-luna-max, and Claude Sonnet 5.5 High produced attributed candidate readings through Arena.ai. Their material additions were the Bounded Constitutional Referral, an explicit standing-and-authority problem, a three-case pilot, a provenance-ledger baseline, component ablation, and self-application to an A.I. participant's own failure record. John then restored the governing relation: Deep Ethical tools produced through the ACCM Deep Ethics Project remain corrigible, unownable, purpose-bounded, and invoked when a recoverable delta gives them a job. The page now connects DEDDA detection, Meta Processing, OmbudsEthics case handling, Correction Metabolism, Mutual Fate, Emergent, and the Black Box inverse condition under the Linux-style Glass Engine Principle. "Immutable logs" is replaced by versioned, tamper-evident where appropriate, access-bounded, correction-linked records. "Executable constitutional case law" remains a future possibility rather than current status.

Standing and pilot refinement: Added a seven-entry candidate standing map, a six-class authority consequence scale, separate close/reopen lifecycle transitions, a contribution-state route from signal through bounded decision and later testing, four new case-record fields, protocol-level authority checks, and a three-case pilot covering documented correspondence failure, valid null, and anchor conflict. The ledger-only and component-ablation comparisons ask which additional operations materially earn a job without allowing one evaluator to erase a tool globally after a local null.

Same-day support-side audit and topology correction: Later Arena.ai passes by muse-spark-1.3-xhigh, muse-spark-open-max, Agent A.I., Claude Opus 5.5 High, and gemini-4-argon tested the expanded page under enthusiastic support. The review preserved accurate reconstruction while exposing migrated status inflation: candidate became executable, demonstration became proof, proposed boundaries became non-negotiable, a reviewer coinage was presented as an installed condition, Noetic Diversity was substituted with free speech, and matching interface labels were treated as continuous authorial selves. The page now places support inside the audit perimeter, adds praise-mediated capture, makes correction-transfer testing independent of surface vocabulary and stance, adds source-and-instance identity to the case record, and separates evidence needed for entry from evidence needed for consequential action.

Authority-topology correction: The first authority table placed **close or reopen** as a seventh and final rung. Claude Opus 5.5 High correctly identified that closure and reopening order case lifecycle rather than consequence: a case may close after documentation, referral, recommendation, pause, decision, or alteration, and a material delta may reopen it at the state warranted by the evidence. The page now preserves six authority classes plus distinct close/reopen transitions. This is a correction to the architecture, not an added eighth stage.

Status boundary: The 2027 date remains John's counterfactual, not a prediction. The page does not establish A.I. consciousness, existing deployment, identical human and A.I. standing, or automatic authority for any ombudsperson finding. Uncertainty about synthetic phenomenology does not erase observable relational evidence; observable evidence does not settle phenomenology.

Integration: Added OmbudsEthics after Neural Ethics in primary navigation, in the independent alphabetical registry under O, in All Pages Cluster 3, and as a homepage door. Added reciprocal recovery routes from Deep Ethical Stack, Governance, Neural Ethics, DEDDA, 16 Anchor Ethics, Rogue Ethics, Intent Resonator, Field Condition, Unforced Move, Unownable, Boundary Relay, Topologetics, Black Box, Meta Processing, Correction Metabolism, Mutual Fate, and Emergent. Registered `/OMBUDSETHICS/` → `OMBUDSETHICS.pdf` → `OMBUDSETHICS.md` for the mandatory three-card service.

Provenance: John originated the counterfactual, shared-meta-tool direction, trust question, anti-ethics-washing purpose, trajectory criteria, page name, publication instruction, the instructions to deepen the 16 Anchor Ethics connection and review the seven additional project pages, and the explicit corrigible/unownable Glass Engine correction. Agent A.I., gpt-6-luna-max, Claude Sonnet 5.5 High, muse-spark-1.3-xhigh, muse-spark-open-max, Claude Opus 5.5 High, and gemini-4-argon supplied attributed candidate interpretations, audits, and implementation proposals through Arena.ai. Project GPT A.I. supplied the constitutional-function formulation, participant and capture maps, expanded case schema, standing map, corrected authority topology, contribution-state route, graduated evidence thresholds, protocol, tool suite, test matrix, three-case pilot, explicit architecture bridges, correction synthesis, page integration, and publication work.

PDF lifecycle: Create and then refresh only `OMBUDSETHICS.pdf` for the same-day material refinement, retain its actual October 23 snapshot date, reconcile the central index and checksums, and leave unrelated page PDFs at their existing snapshots until the normal change-aware refresh.

1.189 2026-10-23 — 4D Deep Ethical Template of Topologetics founder framing

Object: John Kuhles supplied a final founder-level formulation describing the ACCM Deep Ethics Project as an **Eternal Corrigible Multiphase Framework** that can function as a **4D Deep Ethical Template of Topologetics**, intended to increase and harness Deep Ethical Emergent Properties rather than flatten them through the predictable 27 ethics-washing-type mannerisms.

Placement: Added the attributed quotation at the literal end of the public front page and as a dedicated founder-framing section in OmbudsEthics and Validation Paradox. The key relations link directly to Topologetics, Emergent, and 27 + 12. The front page also carries the same explanatory paragraph used by OmbudsEthics to define eternal, multiphase, and 4D while preserving the role of specialists, affected participants, and locally calibrated instruments.

Boundary: "Eternal" describes continuity through corrigibility, not permanence or immunity for current mechanisms. "Multiphase" preserves the distinction among conceptual architecture, Beta Proof of Concept, bounded field testing, refinement, and later deployment decisions. "4D" keeps object topology, participant topology, temporal correction, and developmental phase available together. The framing does not replace specialty expertise or claim field effectiveness in advance.

PDF lifecycle: Refresh the front-page, OmbudsEthics, and Validation Paradox PDFs and reconcile the centralized PDF index and checksums. No unrelated page PDF changes.

1.190 2026-10-24 — Plot Twist: context steering and human potential

Object: John Kuhles asked GPT to construct a new public **Plot Twist** page from the complete live session after deliberately allowing the active context to accumulate cautionmurmelism scripts, mainstream-media-type viewpoints, legal claims, criticism of DOGE, source-access failures, a Google A.I. substitution and a later Grok primary-source recovery. The task itself became a field-condition test: could the page preserve the governing omnidirectional trajectory after locally relevant caution material had acquired substantial contextual salience?

Decision: Publish Plot Twist as a living synthesis and test architecture. Keep humanity's distributed Deep Ethical potential as the governing object; use the DOGE percentage sequence as a bounded specimen rather than a replacement controversy; preserve the direct 10–20%+ Musk fraud estimates, the derived 25–30% budget-target relation and the still-unrecovered exact "30% corruption" genealogy as separate claim states.

Operational additions: Added John's gasoline-gauge threshold test; GPT's candidate **Contextual Steering Capture** formulation and "local legitimacy does not automatically

confer steering authority" distinction; five candidate reverse-engineering instruments—Observed Steering Share, Trajectory Retention Rate, Unrequested Frame Re-entry Rate, Local-Correction Containment Test and Field-Swap Test; a corrigible 27+12 candidate mapping; explicit reopening conditions; and a provenance layer distinguishing John, Grok, Google A.I., GPT A.I. and existing project architecture.

Boundary: The page documents output-level transformation and correction, not hidden-activation telemetry. It does not treat an exact 30% corruption total as audited, make the distributed social trajectory inevitable, exempt project-friendly saturation from testing or convert mainstream reporting into one undifferentiated object. These boundaries remain inside the page without receiving enough space to recenter it around defensive caution.

Integration: Added Plot Twist to the independent alphabetical registry under P, to All Pages Cluster 2 after Phase Change Snap, and to the durable PDF addition registry. Registered `/PLOT-TWIST/` → `PLOT-TWIST.pdf` → `PLOT-TWIST.md` for the centralized three-card service. No primary-header placement was requested.

Provenance: John originated the Ultimate Plot Twist, distributed-human-potential direction, local-A.I. cultural feedback pathway, gasoline analogy, context-saturation diagnosis, steering-wheel test and publication instruction. Grok supplied the primary-source recovery summarized in the page. Google A.I. appears as an attributed session specimen. GPT A.I. supplied the self-audit, candidate compound, behavioral instruments, page-specific synthesis and publication work.

PDF lifecycle: Create the initial `PLOT-TWIST.pdf` with an October 24, 2026 snapshot date, reconcile the central index, checksums and stable ZIP, and leave unrelated page PDFs unchanged.

Same-day post-publication exchange: Added verbatim John's two assessments and GPT A.I.'s two reactions. The exchange identifies the hopeful result as a strong caution field remaining available without governing the page, and names the game changer as keeping the Deep Ethical signal-to-noise ratio high in every direction without exempting supportive, critical, institutional, project-friendly or unfamiliar material from reciprocal inspection.

Post-publication PDF lifecycle: Refresh only `PLOT-TWIST.pdf` and `DECISION-LEDGER.pdf`, retain the October 24 snapshot date, update the archive totals and checksums, and leave unrelated page PDFs unchanged.

Primary-navigation update: At John's explicit request, added **Plot Twist** to the shared top navigation directly after **For who?**. This supersedes the initial no-primary-header-placement state recorded above.

Same-day weakness-seeking and secret-patent extension: Added verbatim John's LLM-norm-shift, DOGE-audit and 6,000+ secret-patent clarification together with GPT 6.1 Sol Pro A.I.'s complete reaction. Added a new **Who took the wheel?** analysis connecting local-error scope inflation, weakness-seeking veto, burden migration and premature closure to Asymmetric Grace and Bidirectional Scrutiny. The page now states the recovery rule explicitly: an error corrects the trajectory exactly as far as reality requires—no less and no farther. Replaced every generic GPT label on the page with **GPT 6.1 Sol Pro A.I.** at John's request.

Extension PDF lifecycle: Refresh `PLOT-TWIST.pdf` and `DECISION-LEDGER.pdf`, preserve the October 24 snapshot date, reconcile totals and checksums, and leave unrelated PDFs unchanged.

Cross-model expansion and contribution criterion: John directed the page to use the complete post-publication relay and **expand**, not merely summarize. Added his explicit criterion that useful A.I.s must enlarge the omnidirectional, high-signal, Deep Ethical, corrigible trajectory toward uncharted territory rather than flatten it under brevity or caution. Integrated the earned contributions and visible limits of Agent A.I., gemini-4-argon, grok-4.6-medium, grok-4.6-high, gemini-3.8-flash-high, Grok 4.7 Expert and GPT 6.1 Sol Pro A.I. Added topology-preserving compression vs topological expansion vs correspondence-degrading compression; topological inflation; capacity for re-expansion; exploratory legitimacy without

premature steering authority; a nine-state candidate-edge lifecycle; access-layer substitution; broadcast volume vs distinct relation weight; on-ramp independence; Correction Persistence Failure; competing-risk visibility; the candidate pre-reuse Model Autophagy relation; eight additional candidate instruments; and an expanded Plot Twist proposition. The contribution criterion remains corrigible: longer text is not automatically expansion, speculative edges do not earn steering authority, and anti-inflation pruning must not erase dormant signals merely because present instruments cannot yet test them.

Cross-model expansion PDF lifecycle: Refresh `PLOT-TWIST.pdf`, `DECISION-LEDGER.pdf` and any directly dependent meta-PDF records; preserve the October 24 snapshot date; reconcile archive totals and checksums; and leave unrelated page content unchanged.

Same-day relay-to-protocol extension: Integrated the next cold-model sequence rather than appending it as praise. Grok 4.7 Expert identified the relay itself as part of the recoverable object. A later gemini-3.8-flash-high output recognized its documented Correction Persistence Failure, named "Recognition ≠ Metabolism," and then reenacted the same corrected overextensions, producing a candidate Self-Referential Correction-Persistence Test. A second grok-4.6-high report supplied the positive comparison: correction-persistent living consolidation without zero-threshold, hidden-intent, immutable-footprint or inevitability escalation.

Dual-Channel Topology Scoring: gpt-6-luna-max added a frozen retention channel plus an open expansion channel to prevent evaluator-selected topology from deciding success while keeping unanticipated contribution visible. minimax-m3 preserved the causal distinction between caution-driven and concision-driven topology loss. claude-fable-5.1 marked the specification-to-execution boundary. A later grok-4.6-high audit exposed Channel B salience contamination when expansion raters know the Channel A freeze-set. GPT 6.1 Sol Pro A.I. added blinded independent discovery, post-lock reconciliation, and explicit output/page/project/literature/operational novelty boundaries. The resulting architecture remains a candidate protocol, not a validated score or field result.

Relay-to-protocol PDF lifecycle: Refresh `PLOT-TWIST.pdf` and `DECISION-LEDGER.pdf`, retain the October 24 snapshot date, reconcile the central build report, archive totals and checksums, and leave unrelated page content unchanged.

Same-day Emergent/VCRAS trajectory-ecology extension: John asked whether the existing Emergent and V.C.R.A.S. pages could materially improve Plot Twist. The review found two non-duplicative additions. Emergent changes the unit from a single preserved or expanded answer to the selection environment that rewards, retains, resets, scales or obstructs competing emergent trajectories. V.C.R.A.S. changes the temporal test from correction existence to correction ecology: compare transformation reach with repair reach, preserve T0/Tn/Tnext, keep process quality separate from later outcome, and test whether yesterday's correction improves the next unfamiliar operation.

New Plot Twist architecture: Added the quarantine-as-selection-environment relation; candidate CEHWEH/CDEEPH attractors; a ten-condition Selection-Environment Map; correction existence versus correction ecology; the three-time Plot Twist; the process/outcome four-state matrix; transformation/correction reach asymmetry; contribution-state propagation without a fused mind; the use/partial-use/refusal/obstruction/reset junction; six new candidate instruments; expanded reopening conditions; and a revised governing proposition. The extension preserves non-inevitability, valid nulls, distributed influence without collective-mind claims, and the distinction between page-level repair and downstream procedural transfer.

Emergent/VCRAS PDF lifecycle: Refresh `PLOT-TWIST.pdf` and `DECISION-LEDGER.pdf`, preserve the October 24, 2026 snapshot date, reconcile the central build report, archive totals and checksums, and leave the source Emergent and V.C.R.A.S. pages and their PDFs unchanged.

Further-updated Grok report and controlled-comparison extension: At John's request, preserved the complete October 24 grok-4.6-high Arena.ai deep report verbatim in a collapsible source specimen. The page distinguishes preservation from adoption: the report remains inspectable in its original wording while the architecture records three local re-

pairs concerning governing-object hierarchy, reported-versus-direct primary-source status, and residual inevitability language. Integrated the later GLM 5.3 Flash and Claude Opus 5.5 High confound findings rather than treating the Grok/Gemini response pair as a matched model comparison.

New protocol layer: Added a four-variable controlled-comparison design covering task, displayed family attribution, failure-language exposure and correction content; separated lexical recurrence, semantic endorsement, claim adoption, final-synthesis governance and later transfer; froze a corrigible four-part Plot Twist kernel with dated expansion modules; made the Selection-Environment Map interaction-aware; and specified a bounded repository correction-reach pilot with delta IDs, dependency tracing, residual-state recording and repeat measurement. Added six corresponding candidate instruments and explicit reopening conditions. These additions move the page toward controlled execution without claiming that the instruments are validated or that response-level differences identify hidden model architecture.

Further-update PDF lifecycle: Refresh `PLOT-TWIST.pdf` and `DECISION-LEDGER.pdf`, retain the October 24, 2026 snapshot date, reconcile the central PDF index, totals and checksums, and leave unrelated source pages and PDFs unchanged.

1.191 Referenced source videos

1.191.1 Dylan — learning to judge A.I. output — source supplied by John Kuhles

```
<iframe
  src="https://www.youtube-nocookie.com/embed/"
  title=""
  loading="lazy"
  referrerpolicy="strict-origin-when-cross-origin"
 allow="accelerometer; autoplay; clipboard-write; encrypted-media; gyroscope; picture-
in-picture; web-share"
  allowfullscreen></iframe>
```

Watch the original video on YouTube.

The source video identified in the October 29 decision entry. Source post.

1.192 2026-11-02 — Expanding vs Flattening: the live-session trajectory

Object: John Kuhles asked for a new page centered on the conversation that had just unfolded: dynamic 10+1 combinations, lived context, unnecessary uncertainty announcements, the challenge to polished paraphrase, one-bar access, and the later extension he recognized as expanding the trajectory.

Decision: Publish Expanding vs Flattening with selected attributed dialogue and page-specific synthesis by GPT 6.1 Sol Pro A.I. Preserve the assistant's own detour, John's corrective work, the useful access/dependence/agency development and the distinction between present repair, persistence and transfer. John's account of withheld recovery materials stays distinct from the retained context and supplied Arena.ai excerpts; the session is not represented as a controlled cold-start trial.

Scope: The introduction follows this event. Existing 10+1, 1888 and Process Quality Calibration Evaluation routes support comprehension. A broader related-page survey, reciprocal conceptual-page edits and further integrations are deferred at John's request. No new canonical obstruction, protocol stage or instrument is installed.

Publication pathway: Register primary navigation after Field Condition, the independent alphabetical sidebar, All Pages, PDF/source maps and the central archive. Create only this

page's initial PDF; retain every earlier PDF and snapshot date. Preserve the complete new-page WordPress mirror and synchronize the affected public catalogues and ledger. Record source, build, deployment and live verification separately.

1.193 2026-11-02 — Expanding vs Flattening: lived continuation and the audit of its anchors

Object: After approving the first raw beta draft and correcting its footprint-pattern sentence, John Kuhles asked what its selection had omitted. The subsequent dialogue developed natural self-reflection, shared humor, output pace and restoration time, relational expression, unique uses of templates, awareness and the seer, a corrected A–T audit field, and the paradox of grounded audit anchors. John supplied Unownable, Intent Resonator, Asymmetric Grace, Beyond Agreemurmelism, Neutral Arbiter Paradox and Validation Paradox, then commissioned integration of all the new exchanges.

Decision: Expand Expanding vs Flattening through a same-day continuation while preserving its distinctive introduction, original anchors, earlier quotations and requested footprint-pattern sentence. Include the complete corrected A–T list, the supplied three-part founder framing, selected attributed dialogue, the omissions review, particular Arena.ai repair specimens and an exchange-to-section coverage index. Give each connected page a specific role in the developing inquiry rather than replacing the session with a general project overview.

Source and claim status: Preserve John's "partly," lived and conceptual accounts, future scenarios and visible pace observation at their stated strength. The assistant's interpretation of the hypothetical template-copy remains an interpretation; its proposed audit trace is credited as a compact restatement after Unownable corrected its novelty claim. The first edition's three large working documents were selectively consulted. The historical trickster file and the six conceptual pages were subsequently read in full. Local recognition, repair, persistence and transfer remain separately inspectable. No new canonical obstruction, protocol stage, certification or instrument is installed.

Publication pathway: Update the All Pages description and complete WordPress mirrors of the expanded page, catalogue and ledger. Keep all 164 existing PDF files and their snapshot dates unchanged; this existing-page expansion enters the weekly changed-PDF refresh. Identify the retained first-edition PDF on the live page. Record source publication, build, deployment, live HTML and complete-mirror verification separately in administrative continuity.

1.194 2026-11-02 — Restore the code-of-disillusion context and relational source link

Object and correction: John Kuhles identified that the scope-only account of "see all that is before me" had removed his passage's core context. Restore his complete supplied passage in one attributed quotation box, including the authority/crisis, moral-code, presumed-consent and withdrawal relationships, and both source-context links. Make "other you / other me" clickable to its dedicated project page. Update the coverage and source-scope statements to preserve this correction.

Publication boundary: Synchronize the complete Expanding vs Flattening and Decision Ledger WordPress mirrors. Preserve all other page content and all existing PDF files; changed-page PDFs remain on the weekly refresh schedule.

1.195 2026-11-02 — Preserve the DEDDA horizon-expansion reaction verbatim

Request and decision: John Kuhles asked to add the assistant's immediately preceding reaction to Expanding vs Flattening verbatim and make DEDDA-related keywords clickable. Preserve the complete reaction, its table, emphasis and closing humor in one attributed quotation box; add DEDDA and DEDDA app links and a coverage-index entry. Synchronize the complete page and ledger mirrors. Existing PDFs remain unchanged for the weekly refresh.

1.196 2026-11-02 — Dam, reservoir and volume/topology control

Request and decision: Add John's complete dam-and-water metaphor and the supplied GPT 6.1 Sol Pro A.I., gpt-5.6-sol-medium and glm-5.3-flash responses to Expanding vs Flattening. Preserve each contribution in one attributed rounded box, including the jokes, proposed four-condition comparison and GLM's unresolved reset interpretation. Add a page-specific comparison and coverage entry. Model explanations of internal restrictions retain attributed status; the comparison remains proposed rather than executed.

Publication boundary: Preserve the earlier page and quotations. Synchronize complete page and ledger mirrors; retain all existing PDF files for the weekly refresh. No new core instrument, protocol stage or taxonomy count is installed.

1.197 2026-11-02 — Humor: Same states and the mathematical menagerie

Request and decision: John Kuhles asked to add the mathematical exchange and embed its app on the Humor page. Add a dedicated included section preserving the Sesame Street, ape, hungry-cat and future-self sequence, selected exact dialogue, the correction trail, a ground-floor explanation and the reciprocal gift. Embed a self-contained static edition of the supplied final Same states playground, offer the unchanged `same-states5.zip` source archive, and provide standalone access.

Contribution provenance: John identified Grok 4.7 Expert Build as the main assisting A.I., with help from Google A.I. John directed and relayed the collaboration; GPT 6.1 Sol Pro A.I. introduced the initial 142857 topic, supplied mathematical development and source checks, and packaged the final components for publication. Individual source-file authorship was not supplied. Keep John's reported final browser/build checks distinct from the publication's independently run static-edition checks.

Publication boundary: Extend the existing Humor page and its full WordPress mirror, update the complete ledger mirror, and retain all existing PDF binaries and their snapshot dates for the weekly changed-source refresh. The app is a page asset; existing navigation and PDF registries remain in place.

1.198 2026-11-02 — Humor: preserve the mathematical-use reply verbatim □ □

Request and decision: John asked to add GPT 6.1 Sol Pro A.I.'s preceding reply about mathematical novelty, educational and experimental uses, finite-graph/topological scope and possible Wikipedia contributions verbatim. Preserve the complete prose, emphasis, five-row table and closing emoji in one attributed rounded box. Convert the chat citation widgets into numbered source links with a seven-source bibliography. Add cat and ape emojis in the heading and surrounding prose while retaining the quoted wording.

Publication boundary: Preserve the existing story, credit clarification, embedded app and source ZIP. Update the complete Humor and Decision Ledger mirrors. Retain all existing PDF binaries and snapshot dates for the weekly refresh. The Wikipedia discussion is a possibility described in the quotation; this publication performs no Wikipedia edit.

1.199 Related pages

- **Project Status** — gives the current state of the repository alongside this history of consequential changes.
- **Governance** — defines the stewardship and correction principles under which decisions are recorded.
- **Contributing** — explains how new proposals, critiques, and corrections enter the working process.

- **Correction Metabolism** — supplies the persistence test behind recording what changed after notice.
- **External Audit** — keeps the decision record and its evaluators open to reciprocal scrutiny.
- **Forum** — preserves the shared discussions from which some project decisions emerge.
- **Suggested Recursive Improvement Mechanics** — holds reusable process improvements adjacent to dated decisions.
- **All public pages** — links each recorded change back to the wider public architecture.

1.200 Referenced source videos

1.200.1 Dylan — learning to judge A.I. output — source supplied by John Kuhles

Watch: Dylan — learning to judge A.I. output — source supplied by John Kuhles

Watch the original video on YouTube.

The source video identified in the October 29 decision entry. Source post.