Governance

ACCM Deep Ethics Project

Contents

1 Governance and Stewardship **1**
- 1.1 Project stewardship . . . 1
- 1.2 No permanent authority over correspondence . . . 1
- 1.3 BETA means corrigible . . . 2
- 1.4 Canonical source boundary . . . 2
- 1.5 Contributions . . . 2
- 1.6 Divergent projects . . . 2
- 1.7 Recursive audit . . . 2
- 1.8 Recursive improvement mechanics . . . 2
- 1.9 Related pages . . . 3

1 Governance and Stewardship

BETA DRAFT — subject to correction and revision

The ACCM Deep Ethics Project is an evolving project initiated and supervised by **John Kuhles**.

1.1 Project stewardship

John retains final stewardship over the identity and trajectory of this repository. He may reject proposed changes that he judges to move the project away from its intended object, architecture, or process orientation.

This is project stewardship, not epistemic immunity.

Any participant — human or AI — may challenge a claim, definition, taxonomy, inference, method, decision, or assumption on its merits. Evidence and better correspondence can require changes to the project, including changes to formulations previously accepted by John.

John has veto power over project drift. Reality has veto power over John.

That veto is temporary source-integrity while the stand-alone version is being built. It is not permanent authority, and it is not a new sheriff. Veto paradox solved.

1.2 No permanent authority over correspondence

No contributor, AI system, expert, institution, framework, taxonomy, project role, or existing project document receives permanent jurisdiction over the object being examined merely because it has authority, status, familiarity, sophistication, or prior acceptance.

Past context may inform a current interpretation. The current object retains the ability to correct the past representation.

1.3 BETA means corrigible

The project deliberately preserves the ability to revise itself. Categories may split or merge. Definitions may improve. Earlier interpretations may be superseded. Experiments may fail. Proposed concepts may be rejected. Accepted concepts may later require correction.

A correction should preserve enough provenance to understand what changed and why whenever that history is materially useful.

1.4 Canonical source boundary

John-controlled canonical source objects are published separately in Canonical Files — ACCM Deep Ethics Project.

This repository remains the corrigible working architecture: its pages may explain, connect, test, or propose developments around a canonical object. They do not silently replace John's canonical wording. A canonical change follows the authority, checksum, versioning, and proposal rules of the canonical repository.

Canonical status protects provenance and exact wording. It does not create epistemic immunity or cancel the omnidirectional audit described here.

1.5 Contributions

Constructive disagreement is welcome. Agreement is not a condition of participation.

The project values:

- faithful representation before evaluation;
- clarification when something materially relevant remains unresolved;
- explicit preservation of uncertainty and qualifiers;
- symmetrical scrutiny;
- mutual dignity during disagreement;
- willingness to correct and be corrected;
- visible distinction between source material, interpretation, proposal, and accepted project formulation.

Persuasive language, sophistication, popularity, institutional authority, AI consensus, or contributor count do not automatically make a proposal canonical.

1.6 Divergent projects

People remain free to disagree with the project's trajectory or create their own adaptations and forks under the applicable repository license. A fork can pursue a different trajectory without requiring the ACCM Deep Ethics Project itself to abandon its own identity.

1.7 Recursive audit

The methods used by this project are themselves legitimate objects of examination.

Deep Ethics gets no immunity from Deep Ethics.

1.8 Recursive improvement mechanics

Overall-general functional improvements John names, so they return without being repeated, live on Suggested Recursive Improvement Mechanics. That page is a neighbor to this one, not a replacement. A.I.s do not invent entries.

1.9 Related pages

- **OmbudsEthics** — proposes a shared human–A.I.–meta-tool ombudsperson layer with bounded authority, appeal, recusal, minority records, and reality-veto conditions.
- **Godfather of A.I.** — maps anti-complacency into identity, permissions, audit trails, secure transactions, accountability, safety research and corrigible institutional controls.
- **Veto Paradox Solved** — explains the construction-phase source-integrity veto and its intended handoff.
- **Contributing** — translates the governance principles into practical guidance for proposals and corrections.
- **Decision Ledger** — keeps consequential project changes, boundaries, and reasons inspectable.
- **Suggested Recursive Improvement Mechanics** — records general improvements intended to return without repeated prompting.
- **Intent Liability** — compares declared governance purposes with process, results, notice, and later correction.
- **Correction Metabolism** — tests whether accepted corrections persist in later project conduct.
- **External Audit** — makes the auditor and the governed process answerable to the same standard.
- **Project Status** — shows the present build stage and which objects remain provisional.