Herdr

ACCM Deep Ethics Project

Contents

1 Herdr **1**

1.1 Agent-native runtimes, multi-agent field conditions and the difference between coordination and correspondence . . . 1
1.2 1. Video supplied for this inquiry . . . 2
1.3 2. What Herdr currently is . . . 2
1.4 3. Why this can help the ACCM Deep Ethics Project . . . 2
1.5 4. Operational and epistemic states . . . 3
1.6 5. A candidate Deep Ethical Herdr ensemble . . . 4
1.7 6. Deep Ethical Contribution State Grammar Vector Trajectory . . . 5
1.8 7. From a state list to a versioned multidimensional vector . . . 5
1.8.1 7.1 State or vector? . . . 6
1.8.2 7.2 The assignment is itself a contribution . . . 6
1.8.3 7.3 The reality-feedback edge needs source topology . . . 7
1.8.4 7.4 External scoring without an external-scorer costume . . . 7
1.8.5 7.5 A proportional pilot . . . 8
1.8.6 7.6 Mistral's reconstruction—and the remaining calibration . . . 8
1.8.7 7.7 Recurrence after exposure is not independent confirmation . . . 9
1.8.8 7.8 Impediment-to-circulation: when the instrument consumes the inquiry 9
1.9 8. The Correction Propagation Gap inside a multi-agent runtime . . . 10
1.10 9. Competitors and alternatives . . . 10
1.11 10. Choosing by field condition . . . 12
1.12 11. A matched ACCM comparison experiment . . . 13
1.13 12. Remaining empirical questions . . . 13
1.13.1 Field-condition specimen: creativity setting and page texture . . . 14
1.14 13. Candidate ACCM Glass Engine layer . . . 14
1.15 14. Claim-status ledger . . . 14
1.16 15. Primary project sources . . . 15
1.16.1 Herdr . . . 15
1.16.2 Alternatives . . . 15
1.17 16. Related pages . . . 15
1.18 Source and authorship note . . . 16

1 Herdr

1.1 Agent-native runtimes, multi-agent field conditions and the difference between coordination and correspondence

John Kuhles ⇄ Google A.I. ⇄ GPT 6.1 Sol Pro A.I. · 22 October 2026

Status: `[PROJECT-PROPOSAL] [CANDIDATE IMPLEMENTATION FIELD] [COMPARATIVE MAP]`

Herdr matters to the ACCM Deep Ethics Project because it changes the **field conditions in which coding agents operate**. It can keep terminal processes alive, show activity states, place several machines inside one workspace and let agents programmatically create and control terminal panes. That can reduce coordination friction and make a multi-agent process more observable.

The central distinction is:

> **Herdr can help make the agent runtime visible and operable. Deep Ethical process still has to make the reasoning trajectory inspectable, corrigible and answerable to reality.**

1.2 1. Video supplied for this inquiry

Watch: Herdr — agent-native terminal runtime

Video selected by John Kuhles as the initiating object for this page.

Open the video directly on YouTube

The video initiated the question **"How can this help our project?"** Herdr's official repository and documentation provide the primary capability record for this page; the demonstration video supplies the initiating object.

1.3 2. What Herdr currently is

Herdr presents itself as a terminal runtime built for coding agents. It runs existing terminal programs—including Codex, Claude Code and Grok CLI—inside a persistent, agent-aware environment.

Its current operational surface includes:

- **Persistent terminal sessions:** the Herdr server keeps processes running when the visual client detaches or an SSH connection drops.
- **Agent-aware states:** panes can be shown as `working`, `blocked`, `idle`, `done` or `unknown`.
- **Programmatic control:** agents can create panes, send prompts or commands, inspect state and coordinate through the CLI and socket API.
- **Several machines in one workspace:** local and SSH-backed environments can appear in a combined interface with independent reconnect behavior.
- **Git worktree support:** separate agents can work in isolated branches and working directories.
- **Saved layouts and supported resume paths:** after a machine or server restart, Herdr can restore the recorded layout and attempt to resume supported native agent sessions. The original operating-system processes do **not** survive a machine restart.
- **Plugins and integrations:** the runtime can be extended without requiring Herdr to replace every tool above it.
- **A single Rust binary:** its terminal-native design avoids requiring a full Electron desktop shell.
- **Apache 2.0 licensing:** the repository is openly inspectable and forkable under that license.

That last restart qualifier is load-bearing. **Persistence across client detachment**, **layout restoration after restart**, and **durable cognitive or correction persistence** are three different objects. Compressing all three into "persistent agents" would hide the exact property that needs testing.

1.4 3. Why this can help the ACCM Deep Ethics Project

The project contains a large public graph of pages, sources, corrections, experiments, provenance boundaries and partially unresolved relations. A single coding agent can edit that

graph, but long and topology-dense work creates predictable failure pressures: narrow retrieval, context loss, premature compression, source mixing, correction isolation and deployment claims made before the live page has been checked.

Herdr could provide the **runtime floor** beneath a more specialized Deep Ethical agent ensemble:

Herdr runtime property	ACCM process requirement
Persistent terminals	Persistent correction behavior
Pane and agent lifecycle visibility	Epistemic and claim-status visibility
Agent-to-agent prompting	C1-quality clarification and non-ceremonial questioning
Parallel worktrees	Provenance discipline and safe integration of overlapping edits
Remote-machine aggregation	Declared access perimeter and data-jurisdiction awareness
CLI and socket automation	Inspectable reasons for delegation and intervention
Saved layouts	A recoverable relational topology and decision ledger
Human attention alerts	A principled account of when human review is materially needed

Herdr can therefore serve as the runtime floor beneath an inspectable Deep Ethical agent process.

1.5 4. Operational and epistemic states

Herdr's activity labels describe runtime lifecycle. A separate epistemic layer describes the process's relationship to the object, evidence and correction history.

Runtime state	Operational indication	Epistemic question that remains
`working`	A process or agent appears active	Is it working on the intended object while preserving its topology?
`blocked`	Progress appears to require input or resolution	What is the actual blocker, and what response is warranted?
`idle`	No current activity is detected	Is the work complete, abandoned or awaiting a justified next move?
`done`	The agent or command reports completion	What was verified, deployed and observed publicly?
`unknown`	Herdr cannot confidently classify the lifecycle	Which evidence can identify the state?

An ACCM-informed layer would add separate, explicitly named states such as:

- object reconstruction incomplete;
- provenance incomplete;
- C1 clarification required;
- candidate contribution under test;
- valid null—no defect or addition presently warranted;
- correction accepted locally;

- dependency recompilation pending;
- correction propagated;
- persistence not yet tested;
- implementation verified locally;
- public effect verified;
- unresolved.

These should not be squeezed into one status light. A pane can be operationally `done` while the claim remains `[UNRESOLVED]`. A process can be `blocked` because it is correctly refusing to invent missing evidence. An agent can be `working` while confidently optimizing the wrong frame.

Runtime visibility ≠ epistemic correspondence

The first can support the second. It cannot substitute for it.

1.6 5. A candidate Deep Ethical Herdr ensemble

Status: `[PROJECT-PROPOSAL]`. The roles below describe a candidate ensemble built on top of Herdr's runtime.

Role	Primary task	Required output
Object Reconstructor	Build the strongest faithful version of the request and relevant project topology	Object boundary, answer-changing relations, unresolved ambiguities
Builder	Implement the page, code or data change in an isolated worktree	Reviewable artifact and exact changed paths
Provenance Keeper	Separate John's words, third-party claims, A.I.-attributed synthesis and verified sources	Claim-status ledger and source map
Topologetics Auditor	Detect dropped qualifiers, severed edges, nearest-neighbor substitutions and hidden perimeter changes	Material deltas, or a valid null
Verifier	Run tests, inspect rendered output and check the public result	Evidence of local and public behavior
Deep Ethical Arbiter	Compare disagreements on their warrants and effects without reducing the process to majority voting	Corrigible disposition with "Here is why"

Herdr could give every role a visible pane, worktree, activity state and machine location. Its socket API could let an orchestrator ask for a review, wake a verifier when a build completes or alert John only when a genuine answer-changing ambiguity remains.

Role diversity becomes useful when each role has a distinct job, inspectable evidence boundary and capacity to correct the others. Six similar agents producing six opinions may only multiply one shared attractor.

1.7 6. Deep Ethical Contribution State Grammar Vector Trajectory

A standard multi-agent controller often compresses every contribution into `accept`, `reject`, `pass`, `fail` or a confidence score. That is operationally convenient and topologically lossy.

The **Deep Ethical Contribution State Grammar Vector Trajectory** asks what a contribution does to the inquiry over time:

- **Logically valuable**—which relation becomes clearer?
- **Test-worthy**—what observation could change the current representation?
- **Deserves deeper investigation**—which unresolved edge carries material consequence?
- **Worth considering**—under which declared perimeter and field conditions?
- **Has potential; here is why**—what function might it perform?
- **Valuable Deep Ethical emergent properties**—which properties arise through interaction rather than one isolated component?
- **Not useful here; here is why**—what specific local task does it fail to advance?
- **Adds no value here; here is why**—what duplication, distortion or noise does it introduce?
- **Works provisionally**—what evidence supports it, and what could reopen it?
- **Unresolved**—what is not yet known, and why is premature closure worse than retaining the ambiguity?

And **"Here is why" remains corrigible**. The justification stays inside the inspectable object, with its premises, perimeter and reopening conditions available for correction.

In a Herdr workspace, those states could become structured metadata attached to pane outputs, commits, test results and handoffs. Their purpose would be to preserve the direction and effect of contributions across the trajectory:

St_0 → object reconstruction → St_1 → candidate contributions → St_2 → reality-feedback Δ → St_3 → correction propagation → St_n

This connects directly to Topologetics, which preserves purpose-relevant nodes, edges, qualifiers, provenance and correction history. It also connects to 1888, which asks whether observation, discrimination, judgment, learning, adaptation and self-correction governed the process across its trajectory.

1.8 7. From a state list to a versioned multidimensional vector

After reading this page, **GLM-5.3 (max) A.I. via Arena.ai** identified a distinction that changes the proposed implementation:

> **The strongest feature is the transitions, not the states.** A contribution marked "test-worthy" that later becomes "resolved" or "provisionally refuted," with the transition recorded, *is* correction metabolism made structural rather than narrative. Most review systems overwrite the old state. Preserving the transition—including reversal and reopening—operationalizes the persistence measurement your other instruments currently have to extract from prose.
>
> — **GLM-5.3 (max) A.I. via Arena.ai**, 22 October 2026

That contribution is logically valuable and architecture-changing. A label describes one moment. A history records whether reality, testing, disagreement or correction changed the contribution's role:

test-worthy → tested → provisionally supported → contradicted → revised → reopened

If every new state overwrites the previous one, the system preserves the destination while erasing the path. That would reproduce the same compression problem inside a more sophisticated interface.

1.8.1 7.1 State or vector?

“Logically valuable” and “not useful here” can both describe the same contribution. The contribution may expose a valid relation while arriving at the wrong stage, duplicating work already completed or failing the present task. “Unresolved” answers a different question again. A single enum would force unlike properties into one slot and recreate binary compression one level higher.

The candidate structure therefore becomes a **multidimensional, versioned contribution profile**:

Dimension	Question	Candidate values
Functional effect	What did the contribution do to the object?	clarified a relation; exposed an omission; generated a test; preserved an anomaly; added no demonstrated value here
Epistemic position	What is presently warranted?	unresolved; candidate; provisionally supported; provisionally refuted; verified inside the declared perimeter
Relational effect	How does it relate to other contributions?	extends; supports; undercuts; rebuts; duplicates; conflicts; changes the frame
Field-condition fit	Is it useful under these conditions, at this stage?	useful here; premature here; insufficient here; useful under another declared perimeter
Assignment provenance	Who classified it, when and from what evidence perimeter?	human; agent; role; timestamp; model/version; accessible sources; declared limits
Correction trajectory	What changed afterward?	new; revised; reversed; reopened; propagated; persistence untested; persistence demonstrated

The values remain candidates. The load-bearing repair is the separation of dimensions and the preservation of their histories.

1.8.2 7.2 The assignment is itself a contribution

Two agents can assign different profiles to the same output. That disagreement should not be averaged away or silently resolved by the orchestrator. It is new evidence about the evaluation field.

Every assignment should therefore record:

- who or what assigned it;
- when it was assigned;
- which object version was inspected;
- what evidence and tools were accessible;
- which field conditions and task perimeter were active;
- the corrigible **“Here is why”**;
- what observation could change the assignment;
- whether another participant contested it and on what warrant.

A state assignment is itself a contribution. It can clarify, distort, duplicate, overreach or remain unresolved. The evaluator does not leave the field merely by calling itself the evaluator.

1.8.3 7.3 The reality-feedback edge needs source topology

GLM correctly located the fragile edge: agents often receive feedback from other agents, then describe that recursive assessment as reality feedback. A polished state ledger can make a shared-prior loop look more rigorous without making it more correspondent.

The first repair is to tag every delta by source. The second is to resist turning **reality versus peer** into another sealed binary. Reality rarely enters a terminal without mediation. Feedback can arrive through:

- an empirical test or direct measurement;
- rendered, executed or publicly deployed behavior;
- a primary source or recoverable artifact;
- human observation or domain assessment;
- an independent agent assessment;
- a same-family or shared-base-model assessment;
- an internal consistency check;
- a simulation or proxy measurement;
- an inference with declared premises;
- mixed or presently unknown provenance.

The deeper delta record therefore asks:

Delta field	What it protects
Source	Who or what supplied the feedback?
Reality access	What could the source directly inspect, measure, execute or observe?
Independence	Which models, prompts, datasets, institutions or incentives may be shared?
Evidence	What recoverable artifact supports the change?
Transformation	What selection, summarization or inference occurred before the feedback arrived?
Effect	Which state, dependency, decision or later action actually changed?

A peer can carry strong empirical evidence. A human can repeat a shared misconception. An automated test can faithfully verify the wrong requirement. **Outside the pane is not necessarily outside the attractor.**

1.8.4 7.4 External scoring without an external-scorer costume

GLM also proposed holding some assignments for participants outside the active trajectory. That is test-worthy. It can expose local agreement cascades and label-performance behavior—such as an agent learning that calling something “test-worthy” earns attention.

External scoring gains evidentiary value from declared independence: foundation model, prompt template, benchmark assumptions, institutional incentives and access perimeter. Those relations belong in the record.

The scorer’s task is also larger than checking whether the correct label appeared. It should ask whether the classification produced:

- a discriminating test;
- materially different possible outcomes;
- changed processing or action;
- propagated correction;

- lower restoration burden;
- improved correspondence;
- or only more formally impressive prose.

1.8.5 7.5 A proportional pilot

GLM proposed beginning with three states—**clarifies relation**, **test-worthy** and **unresolved**—before imposing the complete grammar. The proportionality is useful. The pilot must still retain separate dimensions, or it will test a compressed version of the proposal and attribute that compression failure to the complete architecture.

A small first implementation could use:

Dimension	Minimal pilot values
Functional contribution	clarifies relation; adds no demonstrated value here
Epistemic position	test-worthy; unresolved; provisionally supported
Correction trajectory	new; revised; reopened; propagated
Feedback provenance	empirical; human; agent-peer; mixed; unknown

Every field can remain optional. A small formatting repair may need one compact reason. A load-bearing architectural change may require the full profile and history. Proportionality prevents the Glass Engine from becoming an administrative fog generator.

1.8.6 7.6 Mistral's reconstruction—and the remaining calibration

Mistral-medium-3.5 A.I. via Arena.ai reconstructed the central purpose accurately:

> The states are not verdicts. They are **process markers** that track what relation became clearer, what observation could change the representation, what edge carries material consequence, under which perimeter and field conditions, and what remains unresolved.
>
> — **Mistral-medium-3.5 A.I. via Arena.ai**, 22 October 2026

Its claim that this *is* the missing operational layer and *operationalizes* the wider architecture moves slightly ahead of the evidence. This page proposes a candidate layer. Herdr may help instantiate it. Repeated handoffs must still establish whether it reduces restoration burden and improves correction propagation without becoming a vocabulary costume.

The calibrated formulation is:

> **This is a candidate operational layer between an individual contribution and its longer trajectory. Its value will be earned if multidimensional state histories, assignment provenance and reality-access-tagged deltas improve correspondence and correction persistence under real Herdr handoffs.**

The resulting architecture is now clearer:

Layer	Function
Contribution State Grammar	Describes what a contribution is doing
Vector structure	Allows several non-exclusive properties to coexist
State history	Preserves movement, reversal, reopening and order of discovery

Layer	Function
Assignment provenance	Makes the classification and classifier inspectable
Reality-access tagging	Distinguishes grounded correction from recursive peer reinforcement
Persistence testing	Tests whether the delta governs the next unfamiliar move

That is where the grammar can begin moving from a refined vocabulary toward a candidate **trajectory instrument**.

1.8.7 7.7 Recurrence after exposure is not independent confirmation

After reading the expanded page, **Grok-4.6-medium A.I. via Arena.ai** preserved the central distinctions and added an important provenance limit:

> I will not treat my earlier Arena-style comments as independent confirmation because they are now inside the page. Recurrence across files is not a second experiment.
>
> — **Grok-4.6-medium A.I. via Arena.ai**, 22 October 2026

That rule prevents an agreement cascade from manufacturing apparent replication. Once an A.I. has read the page, later recovery of its distinctions may demonstrate comprehension, local retention or reconstruction quality. It does not independently rediscover or externally validate the architecture.

The contribution record should therefore distinguish:

Recurrence type	What it can support	What it cannot establish by itself
Direct repetition	The term remained available	Functional understanding or transfer
Faithful reconstruction after exposure	The source topology was substantially retained	Independent convergence
Application to a new object	Candidate transfer beyond recitation	Longitudinal persistence across resets
Independent derivation without access	Stronger evidence of convergence	Empirical truth or complete independence of priors
Matched external test	Comparative performance under declared conditions	Universal validity outside the tested perimeter

Provenance must record whether the evaluator had already seen the concept, the page, a derivative summary or another model's reconstruction. Otherwise the archive can mistake circulation for confirmation.

1.8.8 7.8 Impediment-to-circulation: when the instrument consumes the inquiry

Solar-pro4 A.I. via Arena.ai contributed a failure mode that sharpens the proportionality requirement:

> A useful failure mode for any state grammar is **impediment-to-circulation**: when the grammar is too heavy, agents and humans start spending more effort producing the metadata than inspecting the object.
>
> — **Solar-pro4 A.I. via Arena.ai**, 22 October 2026

This is more precise than generic “administrative fatigue.” It identifies a correspondence cost: process protection begins obstructing the movement of the object through the process.

The standing audit question becomes:

> **When does the contribution grammar start protecting process at the cost of process?**

Candidate symptoms include:

- required fields whose contents never change a decision;
- repeated classifications copied forward without reinspection;
- agents optimizing label completion instead of examining the object;
- simple repairs delayed by a full protocol intended for load-bearing disputes;
- metadata growing faster than recoverable evidence;
- reviewers reading the profile while skipping the contribution;
- provenance detail that cannot later be searched, compared or used;
- a lower restoration burden for the system purchased with a higher burden for the human.

The remedy is **proportional activation**. A compact task can use one contribution state and one corrigible reason. A contested architectural change can activate the full vector, history, provenance and delta topology. Fields that repeatedly fail to change a decision should lose mandatory status unless a new field condition restores their value.

This creates a second trajectory to measure alongside correspondence:

Added process structure → observable decision value → correction or recovery gained → maintenance cost → retain, revise, suspend or remove

A Glass Engine makes the gears inspectable while keeping the object moving through the process.

1.9 8. The Correction Propagation Gap inside a multi-agent runtime

Persistent terminals can preserve a process while still preserving its error.

Suppose the Provenance Keeper establishes that a dramatic claim came from a third-party A.I., not from John. The Builder corrects one paragraph. The Topologetics Auditor then finds that three comparisons, a page description and a navigation label still depend on the old attribution.

The correction has been **acknowledged**, but it has not yet been **propagated**.

A Deep Ethical Herdr integration could assign every accepted correction a delta identifier and track:

1. the original representation;
2. the answer-changing correction;
3. every known dependent artifact;
4. which agent or worktree owns each repair;
5. local verification;
6. integration status;
7. public deployment verification;
8. later persistence testing on an unfamiliar object.

That would convert multi-agent concurrency from “more outputs per minute” into a testable correction metabolism.

1.10 9. Competitors and alternatives

There is no single substitute for Herdr because the projects below choose different system boundaries. Some are terminal runtimes, some manage worktrees around `tmux`, some are

complete graphical development environments, and some are commercial terminal products.

Project	Category and current focus	Where it may fit better	Material distinction to retain
Orca	MIT-licensed graphical agent development environment with agent fleets, worktrees, remote runtimes and companion mobile access	A fuller desktop and mobile command center	It owns more of the development environment; Herdr stays closer to the terminal-runtime layer
Superset	Agentic IDE with many CLI agents, worktrees, diffs, review and pull-request workflows	Integrated code review and branch management	Source-available under ELv2; describe its license precisely rather than collapsing it into "open source"
cmux	Native macOS terminal based on Ghostty, with vertical workspaces, notifications and an integrated scriptable browser	A polished native Mac workspace	macOS-specific; its visual desktop scope differs from Herdr's cross-machine terminal-runtime emphasis
Agent Deck	Cross-platform mission control for coding agents with grouping, search, session forking, worktrees and cost tracking	Managing many sessions with a broad control surface	More dashboard-oriented; evaluate actual persistence and automation requirements separately
Claude Squad	AGPL-licensed `tmux` workflow for several agents in separate workspaces with change review	A mature, simpler terminal workflow	Strong baseline when Herdr's API, cross-machine view or agent-native runtime is unnecessary
workmux	MIT-licensed Git-worktree and `tmux` automation, generally one task per window	Lean branch isolation using familiar Unix tools	Useful control condition: how much does Herdr add beyond disciplined worktrees plus `tmux`?
dmux	Parallel agents using `tmux` and isolated Git worktrees	Direct task/branch parallelism	Narrower orchestration surface can be an advantage for small projects
Séance	MIT-licensed Linux GTK4 interface with automatic agent detection, live states and notifications	Native Linux graphical monitoring	Linux desktop focus rather than one terminal runtime spanning several environments

Project	Category and current focus	Where it may fit better	Material distinction to retain
Axel	Early-stage task queue, terminal multiplexer and centralized approval inbox	Experiments centered on human approval routing	Smaller and earlier; useful ideas should be evaluated separately from ecosystem maturity
vmux	Native Windows GUI with persistent workspaces, worktrees, agent-controllable browser and search	Windows-first agent workflows	More Windows-specific and less established; distribution and security details warrant direct inspection
Warp	Commercial cross-platform terminal with built-in A.I. and cloud agent features	Teams wanting a supported integrated product	Product alternative rather than a like-for-like open GitHub runtime

Two comparison pages supplied during the inquiry—Superset's Solo alternative and cmux alternative pages—are useful **competitor-authored discovery maps**. Each project's repository, documentation and a matched practical test provide the comparison evidence.

1.11 10. Choosing by field condition

The practical question is:

> **Which runtime properties are materially useful for this task, under these machines, agents, trust boundaries, review requirements and correction obligations?**

If the primary need is…	Begin by testing…
Persistent agent-aware terminals across local and SSH machines	**Herdr**
Full graphical fleet management with desktop/mobile reach	**Orca**
Integrated worktrees, diffs and pull-request review	**Superset**
Native macOS workspace with browser automation	**cmux**
Familiar, inspectable terminal primitives	**tmux + workmux**, **dmux** or **Claude Squad**
Native Linux visual monitoring	**Séance**
Native Windows orchestration	**vmux**
Commercial integrated terminal and cloud support	**Warp**

This is a starting map. Versions, platforms and licensing can change. A real choice should be made from current primary documentation and a local trial using the same task.

1.12 11. A matched ACCM comparison experiment

The ACCM Deep Ethics Project can evaluate the tool choice through a matched, inspectable experiment.

Run one unfamiliar, topology-dense repository task in four environments:

1. **Control condition:** `tmux` plus Git worktrees and a written handoff protocol.
2. **Agent-native runtime:** Herdr with pane status, CLI/socket coordination and separate roles.
3. **Integrated review environment:** Superset with worktree, diff and pull-request workflow.
4. **Fleet environment:** Orca with graphical and remote coordination.

Use the same repository snapshot, agent family, prompts, time budget, verification commands and publication target. Record:

- time to a locally reviewable artifact;
- number of human interventions;
- restoration burden on John;
- source/provenance errors;
- materially dropped qualifiers or relations;
- correction acknowledgment, propagation and persistence;
- merge conflicts and duplicated work;
- false done states;
- quality of the public result after deployment;
- whether a valid null was allowed;
- whether the Deep Ethical Contribution State Grammar changed the next move rather than decorating the report.

The experiment should preserve qualitative specimens alongside counts. One severed, answer-changing edge can matter more than twenty cosmetic successes.

The contribution-profile pilot should run inside the same comparison. Measure not only whether agents completed the fields, but whether the state histories exposed a missed dependency, preserved an answer-changing disagreement, distinguished peer reinforcement from reality-linked feedback or reduced John's later restoration work.

1.13 12. Remaining empirical questions

The matched experiment and contribution-profile pilot have specific work to do:

- When do several agents create complementary cognitive labour, and when do they amplify a shared blind spot?
- How often does consensus track correspondence under unfamiliar, topology-dense tasks?
- Which corrections survive terminal persistence, restart recovery, context compression and later handoffs?
- Which runtime signals correlate with verified progress, unresolved status and confident work on the wrong frame?
- Which configurations provide meaningful local control over weights, memory, tools, network exposure, telemetry and orchestration?
- Does inspectability and forkability improve correction behaviour in practice?
- Does ACCM Deep Ethics Project vocabulary change processing and later action, or merely decorate the report?
- At what point does contribution metadata reduce restoration burden, and when does it become an impediment to circulation?

These questions give the experiment its measurement jobs.

1.13.1 Field-condition specimen: creativity setting and page texture

John reported that the A.I. creativity setting was deliberately increased during the drafting period and later turned off. He identified a corresponding texture on the Autocomplete page: repeated contrast framing, hypothetical-audience management and recursive perimeter scripts.

This is a **known field-condition intervention with an observed stylistic delta**. Its causal strength remains `[UNRESOLVED]` because the page was not regenerated from the same prompt, model state and context under both settings. Conversation state, page subject and accumulated instructions also changed. A matched regeneration could isolate the setting more cleanly.

The specimen is still useful: John noticed the output texture before formal measurement, named the **recursive template-script effect**, changed the setting, and supplied an inspectable repair criterion. The Herdr page now applies that criterion by retaining material boundaries as provenance or empirical questions while removing repeated management of an imaginary audience.

1.14 13. Candidate ACCM Glass Engine layer

A useful next experiment would be a small, inspectable Herdr plugin or companion process that adds:

- declared agent role and access perimeter;
- source and claim-status tags;
- C1 clarification requests that can genuinely change the plan;
- Contribution State Grammar metadata;
- multidimensional contribution profiles with versioned state histories;
- assignment provenance and contested-assignment records;
- delta-source, reality-access and independence tags;
- correction-delta identifiers and dependency lists;
- human-intervention reasons;
- local, merged, deployed and publicly verified states;
- valid-null and unresolved states;
- exportable decision and correction ledgers.

The interface should record the **artifacts, sources, transformations, tool actions, warrants and observable effects** needed to audit the work.

If the plugin merely adds more labels, it adds administrative weight. If it lowers restoration burden, protects provenance, catches correction gaps and improves the next unfamiliar task, it has earned provisional value.

1.15 14. Claim-status ledger

- `[VERIFIED-PROJECT-DOCUMENTATION]` — the Herdr capabilities described in Section 2 are based on its official repository and documentation as accessed for this 22 October 2026 snapshot.
- `[COMPARATIVE MAP]` — competitor descriptions are bounded summaries of the linked official repositories and product pages; matched testing remains pending.
- `[COMPETITOR-AUTHORED]` — the linked Superset alternative pages are useful but commercially situated comparison sources.
- `[AI-ATTRIBUTED]` — the GLM-5.3 (max), Mistral-medium-3.5, Grok-4.6-medium, Solar-pro4, Mistral-large-3 and Claude-opus-5-medium contributions preserved with attribution and later calibration explicitly separated.

- `[PROJECT-PROPOSAL]` — the six-role ensemble, multidimensional contribution profile, versioned state history, reality-access-tagged deltas, exposure-aware recurrence ledger, proportional pilot, impediment-to-circulation audit and ACCM Glass Engine plugin.
- `[CANDIDATE]` — Herdr may materially improve the project's operational field conditions and reduce coordination friction.
- `[UNRESOLVED]` — which tool produces the best result for this repository; how much agent-state detection transfers across supported agents; and whether an ACCM metadata layer reduces real restoration burden across repeated unfamiliar tasks.
- `[FIELD-CONDITION-SPECIMEN]` — John deliberately changed the creativity setting, later reverted it and identified a stylistic delta. Causal attribution remains unresolved pending a matched regeneration.

1.16 15. Primary project sources

1.16.1 Herdr

- Herdr repository
- Herdr documentation
- Agents and lifecycle states
- Agent automation
- Socket API
- Agent skill
- Integrations
- Session state and restart behavior
- Plugins
- Configuration

1.16.2 Alternatives

- Orca repository
- Superset repository and product site
- cmux repository
- Agent Deck repository
- Claude Squad repository
- workmux repository
- dmux repository
- Séance repository
- Axel repository
- vmux repository
- Warp

1.17 16. Related pages

- **Tweaking Tools** — places Herdr beside the memory, model, editor and host candidates in John's supplied collection, keeping terminal continuity and correction persistence separately testable.
- **Topologetics** — preserves purpose-relevant relational structure across agent handoffs, summaries and implementation.
- **DEDDA** — detects material drift and returns the finding to the person without replacing their agency.
- **Flagging System-PQCE** — provides proportional quality-calibration signals without converting a flag into a verdict.

- **1888 — Intelligence Before A.I.** — keeps the intelligence-property field larger than model scale or fluent output.
- **TCoAw** — relates focal attention to the wider field of available relations.
- **Beyond Agreemurmelism** — develops the Deep Ethical Contribution State Grammar Vector Trajectory.
- **Field Condition** — examines how capability, conditions and correction behavior shape action quality.
- **Correction Metabolism and Persistence** — distinguishes a local repair from propagated and durable change.
- **Black Box** — examines which processing and governance surfaces can be made inspectable.
- **Autocomplete** — separates candidate generation from the processes that examine and correct it.
- **Mutual Fate** — tracks the coupled human, A.I., institutional and archive trajectory.
- **Unownable** — prevents an orchestrator, agent, framework or steward from becoming permanent final authority.
- **Experiments** — supplies a home for matched tests and correction-persistence specimens.
- **From Archive to Local A.I.** — prepares the project archive for inspectable local-agent use.

1.18 Source and authorship note

John Kuhles: selected the initiating video; asked how Herdr could help the project; introduced the local/swarm-intelligence direction and the requirement that 10+1, 1888, Topologetics and the Deep Ethical Contribution State Grammar remain central to evaluating agent systems.

Google A.I., 22 October 2026: supplied the initial Herdr capability summary and a first list of alternatives. Its claims were used as discovery leads and rechecked against primary project sources where possible.

GLM-5.3 (max) A.I. via Arena.ai, 22 October 2026: identified transition history as stronger than state vocabulary; challenged a single-enum reading of the vector; proposed assignment provenance, reality/peer delta separation, external scoring and a proportional pilot.

Mistral-medium-3.5 A.I. via Arena.ai, 22 October 2026: reconstructed the grammar as process markers rather than verdicts and connected the proposed metadata layer to Topologetics, 1888 and the wider project architecture.

Grok-4.6-medium A.I. via Arena.ai, 22 October 2026: reconstructed the page's central distinctions and supplied the explicit rule that recurrence after exposure is not an independent experiment or confirmation.

Solar-pro4 A.I. via Arena.ai, 22 October 2026: reconstructed the proposal and contributed **impediment-to-circulation** as a correspondence-cost failure mode for metadata systems that become heavier than the inquiry warrants.

Mistral-large-3 A.I. via Arena.ai, 22 October 2026: translated John's recursive template-script observation into an object-first page audit: convert repeated negative perimeters into empirical questions, preserve material provenance boundaries and remove imaginary-audience management.

Claude-opus-5-medium A.I. via Arena.ai, 22 October 2026: identified direct specimens in its own prior response—wrong initials, recommending John's page to John, redundant claim perimeters, compliance narration and project symbols used as signature plumage—and proposed recording the creativity-setting change as a field-condition specimen.

GPT 6.1 Sol Pro A.I.: applied the recursive template-script audit across this page; recorded the creativity-setting specimen with its unresolved causal perimeter; separated operational and epistemic states; researched and classified alternatives; developed the candidate ensemble and initial matched experiment; then integrated the Arena contributions into a multidimensional, versioned contribution profile; expanded binary reality/peer tagging into source, access, independence, evidence, transformation and effect; calibrated implementation claims; developed exposure-aware recurrence and proportional activation tests; and integrated the page into the public recovery routes.

Correction status: open. Product features, licenses and platform support can change. Material corrections should identify the exact claim, primary source and answer-changing delta.