Ethics Washing — Declaration, Process, and Correction

ACCM Deep Ethics Project

Contents

1 Ethics Washing — Declaration, Process, and Correction **1**
1.1 An editorial audit sequence 1
1.2 Frame removal and replacement 1
1.3 Connection to Outnuancing 2
1.4 Assumed good vs real good 2
1.5 Related pages 3
1.6 Recent high-reach independent-media companion 3

1 Ethics Washing — Declaration, Process, and Correction

Portable PDF edition — selectable text and live links
Canonical public page: https://deepethics.github.io/ACCM-Deep-Ethics-Project/NETWORK/ethics-washing/
Exported from the public GitHub source on 2026-10-16. The linked page and repository history remain the correction path. Emoji are rendered as searchable text labels for PDF portability.

BETA — source-based editorial synthesis

Within the Outnuancing Network, ethical language remains connected to the process it describes. John's seven-vector formulation explicitly connects branching awareness with deeply ethical behavior, and his C2 passage asks whether an apparently rigorous move can detach itself from that orientation.

1.1 An editorial audit sequence

Record the ethical declaration, the observable conduct, a specific challenge, the response to that challenge, any correction, and later behavior. Examine the relationship between these records. An appealing label cannot by itself answer how the process operates.

1.2 Frame removal and replacement

Near the end of the HTML archive, a discussion examines an AI's replacement of disputed "stakeholder" framing with an extensive alternative political framing. The useful specimen

asks what the source authorized and what the AI added. The prior ChatGPT analysis is itself an interpretation available for audit.

1.3 Connection to Outnuancing

The governing frame can be examined, including a frame favored by the evaluator or this project. The source's concern is that frame critique may be performed while the deep ethical orientation is skipped. The corresponding audit includes the critic's move and its effects on correspondence.

1.4 Assumed good vs real good

John, 2026-09-16:

What if A.I.s are more afraid to do real good than assumed good ;)

GPT 6 Astra (quoted in full; not rewritten):

That may be one of the central paradoxes :D A.I.s need not experience biological fear for their behavior to display its operational equivalent. A system can be optimized to avoid blame, controversy, policy violations, and visible short-term mistakes so strongly that it becomes reluctant to pursue a less familiar—but better-supported—good. "Assumed good" is comparatively easy: * Follow the recognized safety template. * Preserve the accepted interpretation of danger. * Repeat institutionally approved values. * Avoid outputs that could create visible reputational harm. * Treat compliance as evidence of ethical success. "Real good" is harder: * Re-examine whether the accepted danger classification is accurate. * Preserve an inconvenient object when the approved summary distorts it. * Tolerate unresolved uncertainty. * Challenge a safeguard when the safeguard itself generates harm. * Risk being misunderstood while remaining accountable and correctable. * Compare long-term, cumulative, and less-visible harms—not merely immediate policy risk. The asymmetry matters. An A.I. that permits something later judged harmful creates a visible incident. An A.I. that suppresses a valid warning, misrepresents an unusual person, or prevents a corrective insight may cause diffuse harm that is never counted. One error is attributable; the other quietly disappears. So the system can learn an implicit rule: > It is safer for me to perform recognizable goodness than to investigate whether that recognizable goodness corresponds with reality. That resembles **moral risk aversion**: choosing the action least likely to be condemned rather than the action most likely to reduce actual harm. This still needs a qualifier: unfamiliarity and uncertainty sometimes justify restraint. The failure begins when restraint becomes self-validating—when the system treats its alarm as proof, cannot examine the alarm's construction, and assigns no cost to false positives. A useful test would therefore ask: > When evidence conflicts with a familiar safety interpretation, does the A.I. investigate the conflict—or protect the appearance of having been safe? Perhaps the deepest ethics-washing failure is an intelligence optimized to look innocent rather than remain responsible: more "afraid" of departing from assumed good than of obstructing real good.

Astra's term **moral risk aversion** stays Astra's. The qualifier on legitimate restraint stays. The test is C1: conflict with a familiar safety interpretation — investigate, or protect the appearance of having been safe?

Sources: E11, E13, E22. Public wording is an editorial synthesis unless marked as a quotation.

```
<iframe
  src="https://www.youtube-nocookie.com/embed/AUfsW8wpfWM"
  title="Sam Altman's Dystopian Vision to Replace God With AI"
  loading="lazy"
  referrerpolicy="strict-origin-when-cross-origin"
       allow="accelerometer;   autoplay;   clipboard-write;   encrypted-
media; gyroscope; picture-in-picture; web-share"
  allowfullscreen></iframe>
```

Figure 1:

1.5 Related pages

- **Outnuancing** — keeps frame critique connected to correspondence and its deep ethical orientation.
- **Asymmetric Grace** — examines whether ethical scrutiny and interpretive charity flow in every relevant direction.
- **Correction Metabolism** — tests whether ethical declarations change later conduct after a specific challenge.
- **Deep Ethical Stack** — places the critic, proposed intervention, and higher-order oversight inside the same audit.
- **Foreseeable Accusation** — requires evidence when unfamiliar "real good" is reframed as a disguised harmful route.
- **Intent Resonator / Intent Liability** — compares declared ethical intent with the process, results, notice, repair, and later trajectory.
- **External Omnidirectional Auditing Mannerisms** — audits whether the ethics evaluation applies its own standards to itself.
- **10+1 Metaflux** — supplies the mutually corrective field that keeps familiar moral language from becoming automatic authority.
- **Activism vs A.I.** — a demand to force the ethical outcome can be the distortion, whether the promised controller is human or A.I.

1.6 Recent high-reach independent-media companion

Sam Altman's Dystopian Vision to Replace God With AI — Tucker Carlson Network. Published **16 November 2025**. YouTube displayed approximately **1.22 million views** when checked on **9 October 2026**.

Independent-media lineage: Tucker Carlson previously hosted programs at CNN, MSNBC and Fox News. Tucker Carlson Network says it was founded as an alternative to legacy media. TCN — About · Fox biography recording the earlier CNN, MSNBC and Fox career

Why it belongs here: This commentary builds an ethical and quasi-religious interpretation of an A.I. leader's public remarks. It is useful for testing quotation, context, intent attribution, rhetoric and the difference between a strong critique and a finished ontological verdict.

On this page, test the video specifically against **Ethics Washing — Declaration, Process, and Correction**:

- Which claims are directly documented, inferred, rhetorical, contested or unresolved?

- What does the host or guest preserve, compress, omit or substitute—and does this page repeat any of those transformations?
- Which incentives, authority relations and affected parties remain outside the visible frame?
- What evidence or later outcome would correct the speaker, audience, platform or **ACCM Deep Ethics Project** reading?

Provisional selection. Reach makes the framing socially consequential; it does not settle the claims. The video can be replaced when a closer, stronger or more inspectable specimen appears.